

Great Ideas in Computer Architecture

Dependability: Parity, ECC, RAID

Instructor: Sean Farhat



This lecture is sponsored by:



(for legal purposes, it's actually not)

Systems & Fault Tolerance

- So far in this course:
 - How software works
 - How hardware works
 - How computers combine software/hardware
- Now:
 - What happens when the real world gets involved!
 - Things break in mysterious ways
 - We can “count on” things breaking!

How can we ensure the systems we’ve developed function reliably (or predictably) in the face of failure?

Measurements of Fault Tolerance, Mitigations

- Dependability
 - Avoiding failure! Build “strong” systems :)
- Redundancy
 - Replication!
 - Make copies of your important state/info so that if it breaks, you can use a replica
- Error Correcting
 - If we do incur some damage, how can we:
 - Find the problem? (Error detection)
 - Fix the problem? (Error correction)

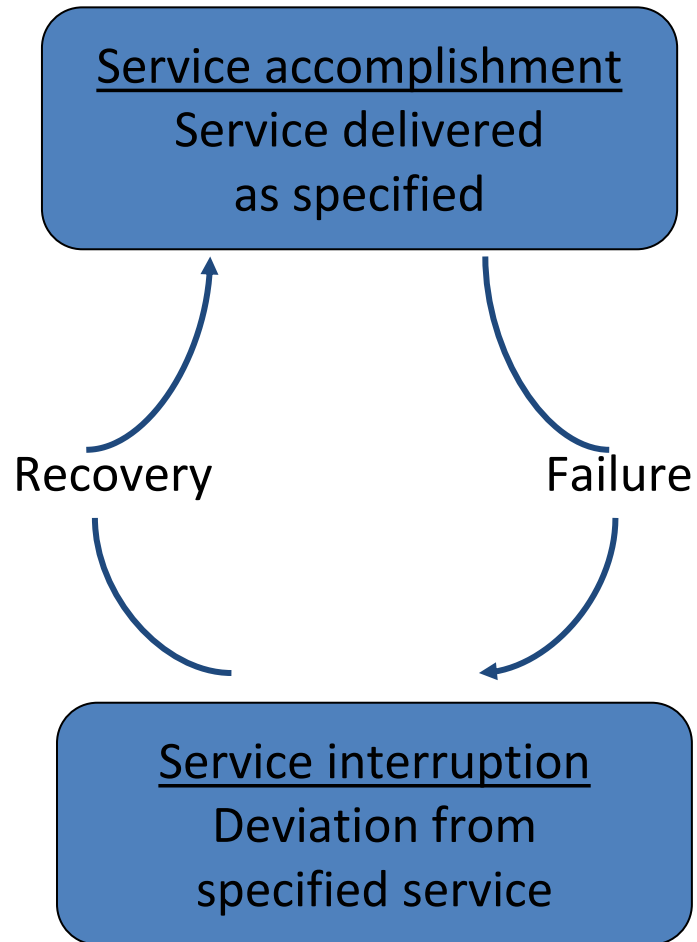
Agenda

- Dependability
- Error Correcting Codes (ECC)
- RAID

What is Dependability?

- A component, or system, is considered highly dependable if it is highly available, or has a low probability of service outages.
 - Unlikely to encounter failures
 - Recovers quickly when failures occur
- The dependability of a system is determined by the overall dependability of its components

Dependability



- **Fault:** failure of a component
 - May or may not lead to system failure
 - Applies to *any* part of the system
- **Repair:** the act of restoring normal service after a fault

Dependability Measures

- *Reliability*: Mean Time To Failure (MTTF)
- *Service interruption*: Mean Time To Repair (MTTR)
- Mean Time Between Failures (MTBF)
 - $MTBF = MTTR + MTTF$
- $Availability = \frac{MTTF}{MTTF + MTTR} = \frac{MTTF}{MTBF}$
- Improving Availability
 - Increase MTTF: more reliable HW/SW + fault tolerance
 - Reduce MTTR: improved tools and processes for diagnosis and repair

Reliability Measures

- 1) MTTF, MTBF measured in hours/failure
 - e.g. “average HDD MTTF is 100,000 hr/failure”
- 2) Annualized Failure Rate (AFR)
 - Average rate of failures per year (%)

$$\text{AFR} = \underbrace{\left(\frac{\text{Disks}}{\text{MTTF}} \times 8760 \frac{\text{hr}}{\text{yr}} \right)}_{\text{Total disk failures/yr}} \times \frac{1}{\text{Disks}} = \frac{8760 \text{ hr/yr}}{\text{MTTF}}$$


Rate that disks
fail per hour

Availability Measures

- Availability = $MTTF / (MTTF + MTTR)$ usually written as a percentage (%)
- Want high availability, so categorize by “number of 9s of availability per year”
 - 1 nine: 90% => 36 days of repair/year
 - 2 nines: 99% => 3.6 days of repair/year
 - 3 nines: 99.9% => 526 min of repair/year
 - 4 nines: 99.99% => 53 min of repair/year
 - 5 nines: 99.999% => 5 min of repair/year


Dependability Example

- 1000 disks with MTTF = 100,000 hr and MTTR = 100 hr
 - MTBF = MTTR + MTTF = 100,100 hr



Dependability Example

- 1000 disks with MTTF = 100,000 hr and MTTR = 100 hr
 - MTBF = MTTR + MTTF = 100,100 hr
 - Availability = $\frac{\text{MTTF}}{\text{MTBF}} = 0.9990 = 99.9\%$



Calculating MTTR Example

- Faster repair to get 4 nines of availability?
- GOAL: Closer to $x/x == 1$, smaller MTTR!
 - $0.9999 = \text{MTTF} / (\text{MTTF} + \text{MTTR})$
 - $0.9999 \times (\text{MTTF} + \text{MTTR}) = \text{MTTF}$
 - $0.9999 \times \text{MTTF} + 0.9999 \times \text{MTTR} = \text{MTTF}$
 - $0.9999 \times \text{MTTR} = 0.0001 \times \text{MTTF}$
 - Plug in $\text{MTTF} = 100,000 \text{ hr}$
 - $\text{MTTR} = 10.001 \text{ hr}$

Dependability Design Principle

- No single points of failure
 - “Chain is only as strong as its weakest link”
- Dependability Corollary of Amdahl’s Law
 - Doesn’t matter how dependable you make *one* portion of system because dependability is limited by the part you do not improve
- In 2013, Google had 99.978% availability for their cloud clusters!

Question

- There's a hardware glitch in our system that makes the Mean Time To Failure (MTTF) *decrease*. Are the following statements TRUE or FALSE?
- 1) Our system's Availability will *increase*.
 - 2) Our system's Annualized Failure Rate (AFR) will *increase*.

	12
(A)	FF
(B)	FT
(C)	TF
(D)	TT

Question

- There's a hardware glitch in our system that makes the Mean Time To Failure (MTTF) *decrease*. Are the following statements TRUE or FALSE?
- Our system's Availability will *increase*.
 - Our system's Annualized Failure Rate (AFR) will *increase*.

	12
(A)	FF
(B)	FT
(C)	TF
(D)	TT

Question

- There's a hardware glitch in our system that makes the Mean Time To Failure (MTTF) *decrease*. Are the following statements TRUE or FALSE?
- Our system's Availability will *increase*.
 - Our system's Annualized Failure Rate (AFR) will *increase*.

	12
(A)	FF
(B)	FT
(C)	TF
(D)	TT

$$\text{Availability} = \text{MTTF} / (\text{MTTF} + \text{MTTR})$$

As MTTF shrinks, our fraction is outweighed by MTTR:

$$100/(100 + 50) \rightarrow 66\%$$

$$50/(50 + 50) \rightarrow 50\%$$

$$10/(10 + 50) \rightarrow 16\%$$

Availability decreases (less 9's, lower percentage)

Question

- There's a hardware glitch in our system that makes the Mean Time To Failure (MTTF) *decrease*. Are the following statements TRUE or FALSE?

- 1) Our system's Availability will *increase*.
- 2) Our system's Annualized Failure Rate (AFR) will *increase*.

	12
(A)	FF
(B)	FT
(C)	TF
(D)	TT

MTTF = "Mean time to Failure" = average time until something goes wrong!

If this number /decreases/, this means failures happen more often!

100 years until failure, 50 years until failure, 10 years until failure, etc.

Failures happening more often == an increased rate of failures per year


As MTTF decreases, AFR increases

Agenda

- Dependability
- Error Correcting Codes (ECC)
- RAID

Great Idea: Dependability via Redundancy

- Redundancy so that a failing piece doesn't make the whole system fail



Great Idea: Dependability via Redundancy

- Applies to everything from datacenters to memory
 - Redundant datacenters so that can lose 1 datacenter but Internet service stays online
 - Redundant routes so can lose nodes but Internet doesn't fail
 - Redundant disks so that can lose 1 disk but not lose data (Redundant Arrays of Independent Disks/RAID)
 - Redundant memory bits of so that can lose 1 bit but no data (Error Correcting Code/ECC Memory)



Error Detection/Correction Codes

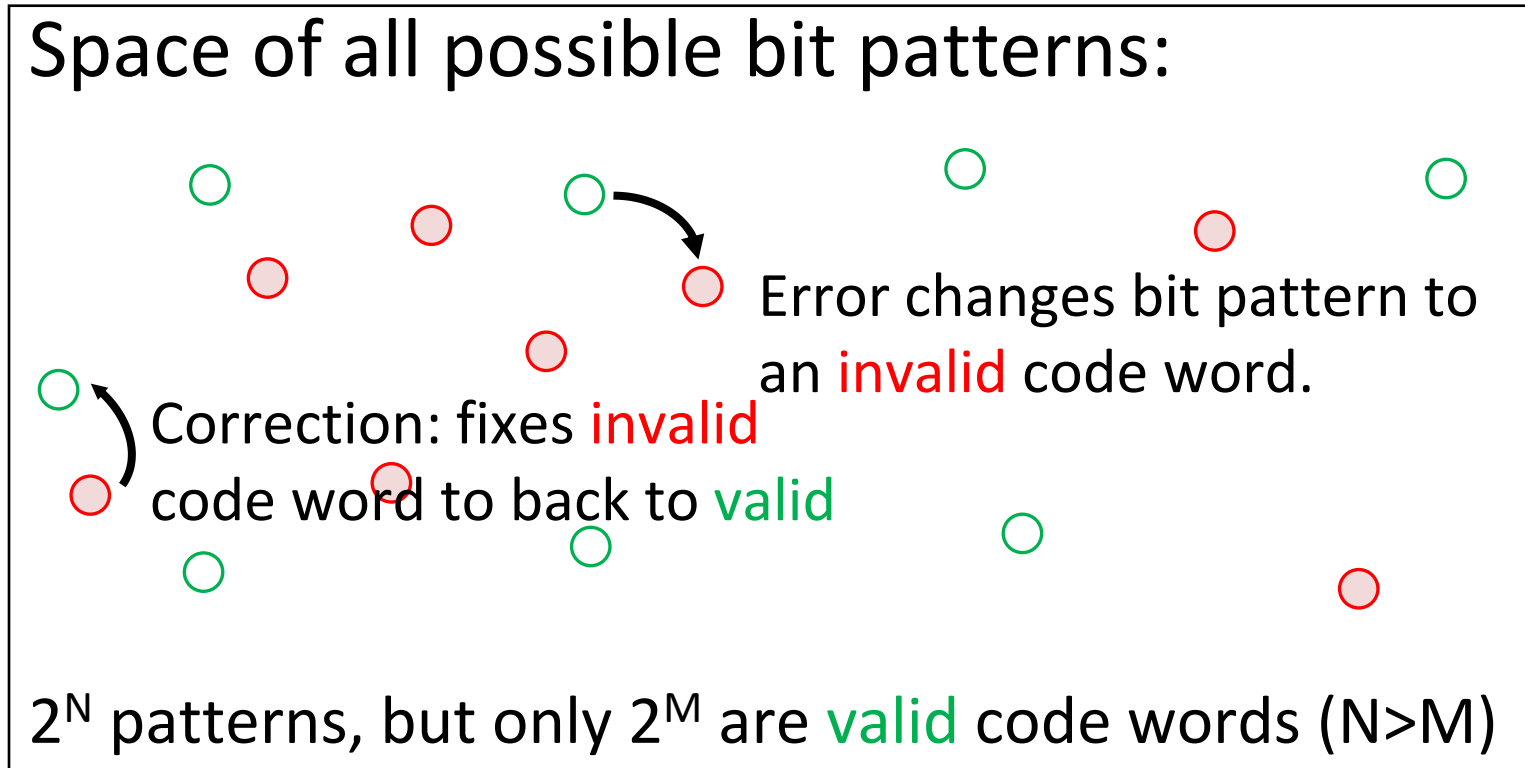
- Memory systems generate errors (accidentally flipped-bits)
 - DRAMs store very little charge per bit
 - “Soft” errors occur occasionally when cells are struck by alpha particles or other environmental upsets
 - “Hard” errors occur when chips permanently fail
 - Problem gets worse as memory systems get denser and larger

Error Detection/Correction Codes

- Protect against errors with **EDC/ECC**
- Extra bits are added to each M-bit data chunk to produce an N-bit “**code word**” ($N > M$)
 - Extra bits are a function of the data
 - Each data word value is mapped to a valid code word
 - Certain errors change *valid* code words to *invalid* ones (i.e. you can tell something is wrong)

Detecting/Correcting Code Concept

- *Detection*: fails code word validity check
- *Correction*: can map to nearest valid code word



Hamming Distance

- Hamming distance = # of bit changes to get from one code word to another
 - $p = 0\underline{1}1\underline{0}11$,
 $q = 0\underline{0}1\underline{1}11$, $Hdist(p,q) = 2$
 - $p = 011011$,
 $q = 110001$, $Hdist(p,q) = 3$
 - If all code words are valid, then
min Hdist between valid code words is 1
– Change one bit, at another valid code word



Richard Hamming (1915-98)


Turing Award Winner

Why does this matter?


- If the entire “space” of possible codewords is valid, we can’t tell if our word has errors (maybe it morphed into another valid word!)
- If some words are valid, and others are invalid, we can detect errors
- If there are a small subset of valid words and a large subset of invalid words, we can trace/track our errors and revert them!
 - Wait really???

3-Bit Visualization Aid

- Want to be able to *see* Hamming distances
 - Show code words as *nodes*, Hdist of 1 as *edges*
- For 3 bits, show each bit in a different dimension:




Minimum Hamming Distance 2



- If 1-bit error, is code word still valid?
 - No! So can *detect*
- If 1-bit error, know which code word we came from?
 - No! Equidistant, so cannot *correct*

Minimum Hamming Distance 3





- How many bit errors can we detect?
 - Two! Takes 3 errors to reach another valid code word
- If 1-bit error, know which code word we came from?
 - Yes!

Parity Bit

- Describes whether a group of bits contains an even or odd number of 1's
 - Define 1 = odd and 0 = even
 - Can use XOR to compute parity bit!
- Adding the parity bit to a group will always result in an even number of 1's (“even parity”)
 - 0b100 Parity: 1, 0b101 Parity: 0
- If we know number of 1's must be even, can we figure out what a single missing bit should be?
 - 10?11 → missing bit is 1

Parity: Simple Error Detection Coding

- Add parity bit when writing block of data:
- Check parity on block read:
 - Error if odd number of 1s
 - Valid otherwise



- Minimum Hamming distance of parity code is 2
- Parity of code word = 1 indicates an error occurred:
 - 2-bit errors not detected (nor any even # of errors)
 - Detects an odd # of errors

Parity Examples

1) Data 0101 0101

- 4 ones, even parity now
- Write to memory
0101 0101 **0**
to *keep* parity even

2) Data 0101 0111

- 5 ones, odd parity now
- Write to memory:
0101 0111 **1**
to *make* parity even

3) Read from memory

0101 0101 0

- 4 ones → even parity, so no error

4) Read from memory

1101 0101 0

- 5 ones → odd parity, so error

- What if error in parity bit?
 - Can detect!


How to Correct 1-bit Error?

- **Recall:** Minimum distance for correction?
 - Three
- Richard Hamming came up with a mapping to allow Error Correction at min distance of 3
 - Called Hamming ECC for Error Correction Code

Hamming ECC (1/2)

- Use *extra parity bits* to allow the position identification of a single error
 - Interleave parity bits *within* bits of data to form code word
 - **Note:** Number bits starting at 1 from the left
- 1) Use *all* bit positions in the code word that are **powers of 2** for parity bits (1, 2, 4, 8, 16, ...)
- 2) **All other bit positions** are for the data bits (3, 5, 6, 7, 9, 10, ...)

Hamming ECC

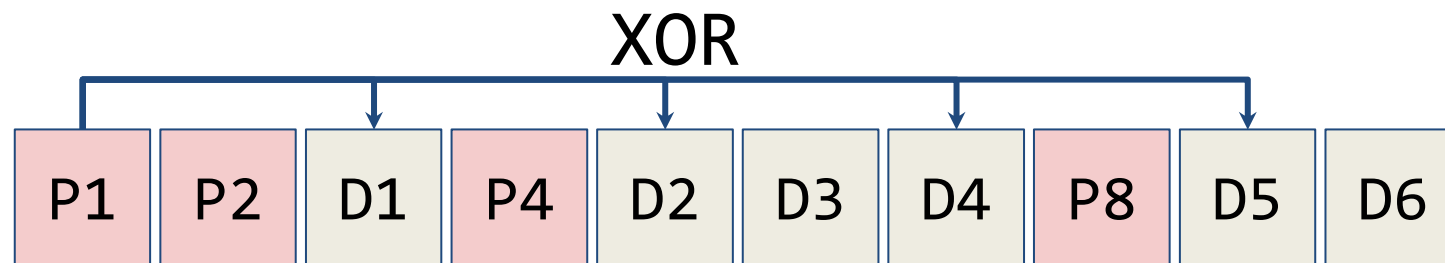



Bit position		1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
Encoded data bits		p1	p2	d1	p4	d2	d3	d4	p8	d5	d6	d7	d8	d9	d10	d11	p16	d12	d13	d14	d15
Parity bit coverage	p1	X		X		X		X		X		X		X		X		X		X	
	p2		X	X			X	X			X	X			X	X			X	X	
	p4				X	X	X	X					X	X	X	X					X
	p8								X	X	X	X	X	X	X	X					
	p16																X	X	X	X	X

Hamming ECC (2/2)


- 3) Set each parity bit to create even parity for *a group* of the bits in the code word
- The **position** of each parity bit determines the group of bits that it checks
 - Parity bit p checks every bit whose position number in binary has a 1 in the bit position corresponding to p
 - Bit 1 (0001_2) checks bits 1,3,5,7, ... ($XXX1_2$)
 - Bit 2 (0010_2) checks bits 2,3,6,7, ... ($XX1X_2$)
 - Bit 4 (0100_2) checks bits 4-7, 12-15, ... ($X1XX_2$)
 - Bit 8 (1000_2) checks bits 8-15, 24-31, ... ($1XXX_2$)

Hamming ECC




Bit position		1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
Encoded data bits		p1	p2	d1	p4	d2	d3	d4	p8	d5	d6	d7	d8	d9	d10	d11	p16	d12	d13	d14	d15
Parity bit coverage	p1	X		X		X		X		X		X		X		X		X		X	
	p2		X	X			X	X			X	X			X	X			X	X	
	p4				X	X	X	X					X	X	X	X					X
	p8								X	X	X	X	X	X	X	X					
	p16																X	X	X	X	X

Hamming ECC



XOR



Bit position		1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
Encoded data bits		p1	p2	d1	p4	d2	d3	d4	p8	d5	d6	d7	d8	d9	d10	d11	p16	d12	d13	d14	d15
Parity bit coverage	p1	X		X		X		X		X		X		X		X		X		X	
	p2		X	X			X	X			X	X			X	X			X	X	
	p4				X	X	X	X					X	X	X	X					X
	p8								X	X	X	X	X	X	X	X					
	p16																X	X	X	X	X

Hamming ECC Example

- A byte of data: 10011010
- Create the code word, leaving spaces for the parity bits:

$_1 _2 1 _3 0 _4 0 _5 1 _6 1 _7 0 _8 1 _9 0 _10 1 _11 0 _12$

Hamming ECC Example

- Calculate the parity bits:
 - Parity bit 1 group (1, 3, 5, 7, 9, 11):
? _ 1 _ 0 0 1 _ 1 0 1 0 → **0**
 - Parity bit 2 group (2, 3, 6, 7, 10, 11):
? ? 1 _ 0 0 1 _ 1 0 1 0 → **1**
 - Parity bit 4 group (4, 5, 6, 7, 12):
? ? 1 ? 0 0 1 _ 1 0 1 0 → **1**
 - Parity bit 8 group (8, 9, 10, 11, 12):
? ? 1 ? 0 0 1 ? 1 0 1 0 → **0**

Hamming ECC Example

- Valid code word: 011100101010
- Original data: 1 001 1010

Suppose we read $0_1 1_2 1_3 1_4 0_5 0_6 1_7 0_8 1_9 \textcolor{red}{1}_{10} 1_{11} 0_{12}$ instead – fix the error!

But how would we figure out where the error is if we *just* see the code word? Hmm....

- Let's examine the parity bits that are dependent on bit 10
 - Maybe we can figure out a pattern?

Hamming ECC Example

- Calculate the parity bits for $0_1 1_2 1_3 1_4 0_5 0_6 1_7 0_8 1_9 \mathbf{1}_{10} 1_{11} 0_{12}$
 - Parity bit 1 group (1, 3, 5, 7, 9, 11):
 $0 + 1 + 0 + 1 + 1 + 1 = 4$ (EVEN)
 - Parity bit 2 group (2, 3, 6, 7, 10, 11):
 $1 + 1 + 0 + 1 + 1 + 1 = 5$ (ODD)
 - Parity bit 4 group (4, 5, 6, 7, 12):
 $1 + 0 + 0 + 1 + 0 = 2$ (EVEN)
 - Parity bit 8 group (8, 9, 10, 11, 12):
 $0 + 1 + 1 + 1 + 0 = 3$ (ODD)

Graphic of Hamming Code

0 1 1 1 0 0 1 0 1 **1** 1 0

Bit position	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Encoded data bits	p1	p2	d1	p4	d2	d3	d4	p8	d5	d6	d7	d8	d9	d10	d11
Parity bit coverage	p1	X		X		X		X		X		X		X	
	p2		X	X			X	X		X	X			X	X
	p4				X	X	X	X				X	X	X	X
	p8								X	X	X	X	X	X	X

- Looks like p2 and p8 were responsible
 - p2 & p8 will be the only incorrect parity bits IFF bit 10 is incorrect
 - Notice that $2 + 8 = 10$
 - Seems like incorrect parity bits tell us where the error is... o:

Hamming ECC Example

- Valid code word: 011100101010
- Recover original data: 1 001 1010

Suppose we see 0₁1₂1₃1₄0₅0₆1₇0₈1₉**1**₁₀1₁₁0₁₂ instead – fix the error!

How to figure out where the error is? Hmm....

- Check the parity bits responsible for bit 10
 - Parity bits 2 and 8 are incorrect
 - As $2+8=10$, bit position 10 is the bad bit, so flip it!
- Corrected value: 011100101010

Hold on...

Replace all the X's with 1's
Everywhere else, put a 0

Bit position	1	2	3	4	5	6	7	8	9
Encoded data bits	p1	p2	d1	p4	d2	d3	d4	p8	d5
Parity bit coverage	p1	X		X		X	X		X
	p2		X	X			X	X	
	p4				X	X	X	X	
	p8							X	X

- This tells you which bit of the code word is incorrect!

Encoded data bits	Parity bit coverage				Bit position
	p1	p2	p4	p8	
p1	1	0	0	0	1
p2	0	1	0	0	2
d1	1	1	0	0	3
p4	0	0	1	0	4
d2	1	0	1	0	5
d3	0	1	1	0	6
d4	1	1	1	0	7
p8	0	0	0	1	8
d5	1	0	0	1	9
d6	0	1	0	1	10
d7	1	1	0	1	11
d8	0	0	1	1	12
d9	1	0	1	1	13
d10	0	1	1	1	14

Hamming ECC Example

Going in reverse:

We see $0_1 1_2 1_3 1_4 0_5 0_6 1_7 0_8 1_9 1_{10} 1_{11} 0_{12}$

Figure out which bit is wrong:

$p_1: 0_1 1_3 1_5 1_7 1_9 1_{11} = \text{even number of 1's} : 0$

$p_2: 1_2 1_3 0_6 1_7 1_{10} 1_{11} = \text{odd number of 1's} : 1$

$p_4: 1_4 0_5 0_6 1_7 0_{12} = \text{even number of 1's} : 0$

$p_8: 0_8 1_9 1_{10} 1_{11} 0_{12} = \text{odd number of 1's} : 1$

Incorrect code bit:

$0b\ p_8 p_4 p_2 p_1 = 0b\ 1010 = 10!$ So flip bit 10

Is This Enough?

- For 1-bit error detection: parity bit
- For 1-bit error correction: Hamming ECC
- What about more bits?
 - For 2-bit error detection: modified Hamming ECC! (see bonus slides)
 - For 2-bit error correction: more advanced schemes such as Reed-Solomon codes (CS 70)

Agenda

- Dependability
- Error Correcting Codes (ECC)
- **RAID**



RAID: Redundant Array of Inexpensive/Independent Disks

- Files are “shared” across multiple disks
 - Concurrent disk accesses improve throughput
- Redundancy yields high data availability
 - Service still provided to user, even if some components (disks) fail
- Contents reconstructed from data redundantly stored in the array
 - Can detect when data is corrupted
 - Can fix data/restore correct version
 - Like before, but “bit” is now “disk”




RAID 0: Data Striping

logical record




striped
physical
records →




- “Stripe” data across all disks
 - Generally faster accesses (access disks in parallel)
 - No redundancy, cannot tolerate any failures
 - Bit-striping shown here, can do in larger chunks

RAID 0




RAID 1: Disk Mirroring




- Each disk is fully duplicated onto its “mirror”
 - Very high availability can be achieved
- Bandwidth sacrifice on write:
 - Logical write = two physical writes
 - Logical read = one physical read
- Most expensive solution: 100% capacity overhead

RAID 1




RAID 2-4: Data Striping + Parity

- Logical data is striped across disks
 - 2: bit, 3: byte, 4: block
- Parity disk P contains parity of other disks
- If any one disk fails, can use other disks to recover data!
 - We have to *know* which disk failed
- Must update Parity data on EVERY write
 - Logical write = min 2 to max N physical reads and writes
 - $\text{parity}_{\text{new}} = \text{data}_{\text{old}} \oplus \text{data}_{\text{new}} \oplus \text{parity}_{\text{old}}$



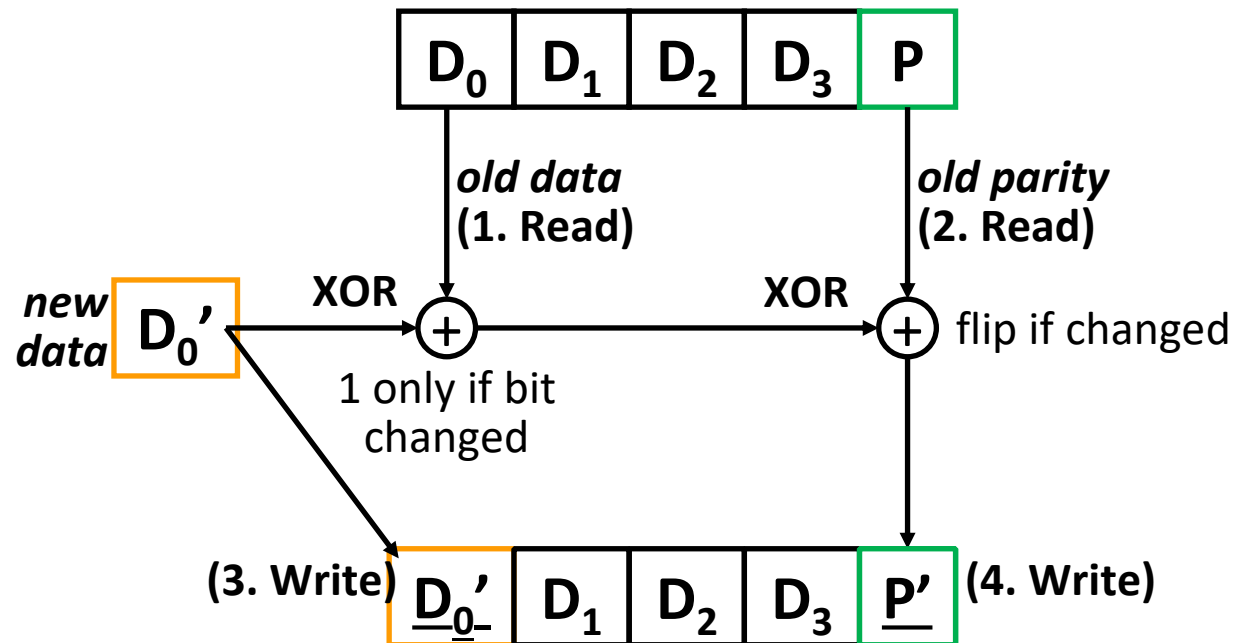
RAID 4



Updating the Parity Data

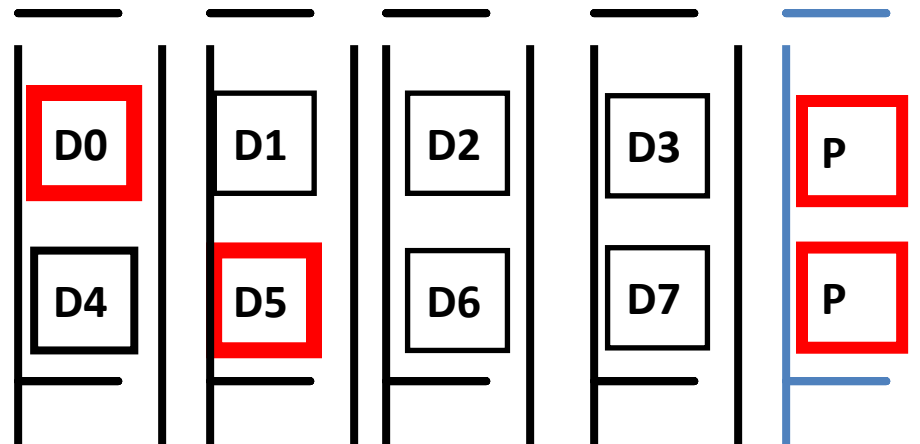
- Examine small write in RAID 3 (1 byte)
 - 1 logical write = 2 physical reads + 2 physical writes
 - Same concept applies for later RAIDs, too

D_0'	D_0	P	P'
0	0	0	0
0	0	1	1
0	1	0	1
0	1	1	0
1	0	0	1
1	0	1	0
1	1	0	0
1	1	1	1



Inspiration for RAID 5

- When writing to a disk, need to update Parity
- Small writes are bottlenecked by Parity Disk: Write to D0, D5 but also write to P disk twice!

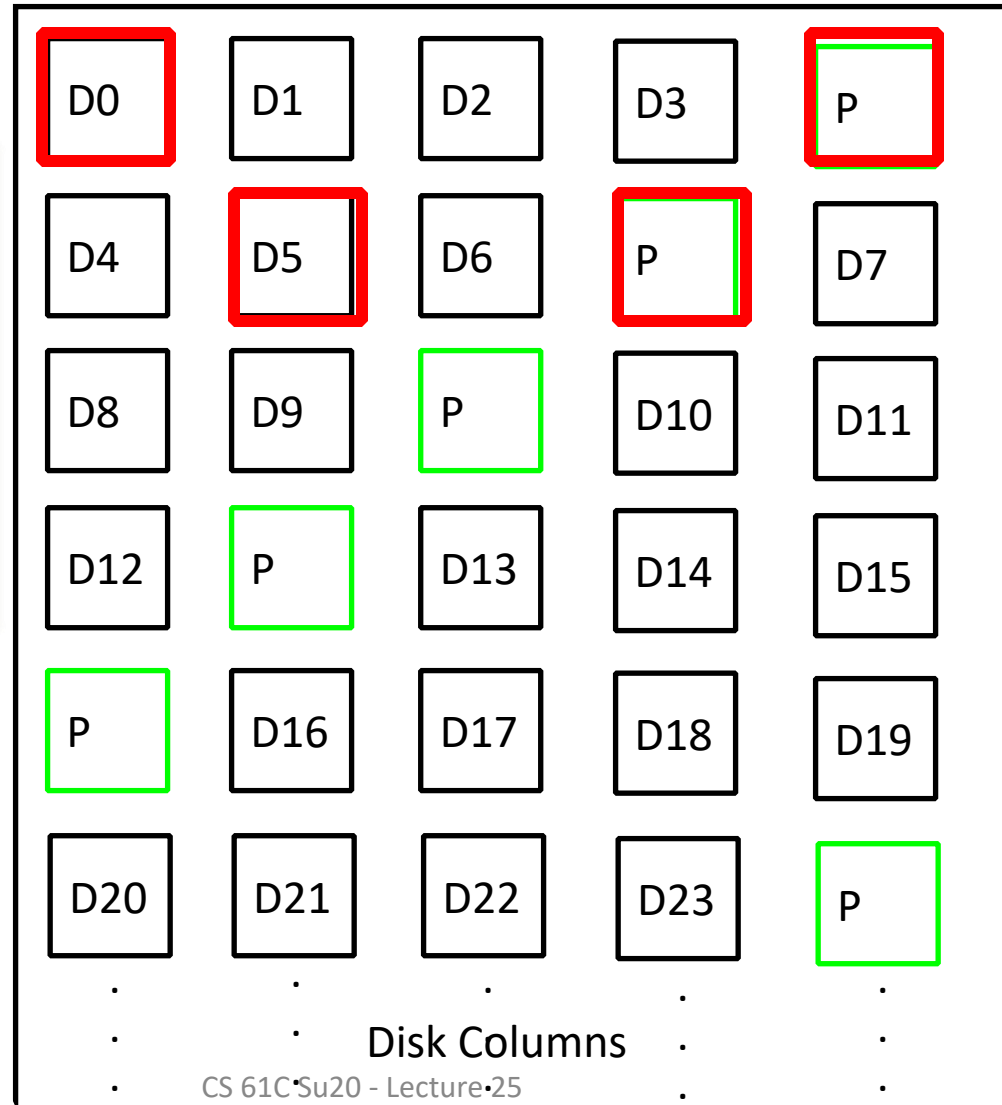


RAID 5: Interleaved Parity




Independent writes possible because of interleaved parity

Example:
write to D0,
D5 uses disks
1, 2, 4, 5





RAID 5



Newer RAID

- RAID 6 – RAID 5 with more parity blocks, so can tolerate 2 failed disks
- RAID 10 – really just RAID 1 + 0,
- You can setup a RAID array on your computer (if you have enough disks) and choose which version you want!



Modern Use of RAID and ECC (1/2)

- RAID 0 has no redundancy
- RAID 1 is too expensive
- RAID 2 is obsolete due to on-disk ECC
- RAID 3 is not commonly used (bad I/O rates)
- Typical modern code words in DRAM memory systems:
 - 64-bit data blocks (8 B) with 72-bit codes (9 B)
 - $d = 64 \rightarrow p = 7, +1$ for DED

Modern Use of RAID and ECC (2/2)

- Common failure mode is bursts of bit errors, not just 1 or 2
 - Network transmissions, disks, distributed storage
 - Contiguous sequence of bits in which first, last, or any number of intermediate bits are in error
 - Caused by impulse noise or by fading signal strength; effect is greater at higher data rates
- **Other tools:** cyclic redundancy check, Reed-Solomon, Fountain Codes, LaGrange

Summary

- Great Idea: Dependability via Redundancy
 - Reliability: MTTF & Annual Failure Rate
 - Availability: % uptime = $MTTF/MTBF$
- Memory Errors:
 - Hamming distance 2: Parity for Single Error Detect
 - Hamming distance 3: Single Error Correction Code + encode bit position of error
 - Hamming distance 4: SEC/Double Error Detection
- RAID:
 - Many different flavors, all with their pros/cons
 - Generally: disks are cheap enough that many of them + some smart organization yields safe storage

Bonus Slides

How does 2-bit detection work for Hamming ECC?

Hamming ECC “Cost”

- Space overhead in single error correction code
 - Form $p + d$ bit code word, where p = # parity bits and d = # data bits
- Want the p parity bits to indicate either “no error” or 1-bit error in one of the $p + d$ places
 - Need $2^p \geq p + d + 1$, thus $p \geq \log_2(p + d + 1)$
 - For large d , p approaches $\log_2(d)$

Hamming Single Error Correction, Double Error Detection (SEC/DED)

- Adding extra parity bit covering the entire SEC code word provides *double error detection* as well!

1 2 3 4 5 6 7 8
 p_1 p_2 d_3 p_4 d_5 d_6 d_7 p_8

- Let H be the position of the incorrect bit we would find from checking p_1 , p_2 , and p_4 (0 means no error) and let P be parity of complete code word (p's & d's)
 - $H=0$ $P=0$, no error
 - $H \neq 0$ $P=1$, correctable single error ($P=1 \rightarrow$ odd # errors)
 - $H \neq 0$ $P=0$, double error detected ($P=0 \rightarrow$ even # errors)
 - $H=0$ $P=1$, an error occurred in p_8 bit, not in rest of word

SEC/DED: Hamming Distance 4

