

# **CS 61C:**

## **Great Ideas in Computer Architecture**

### ***RISC-V Instruction Formats***

# Administrivia

- Project 2: Will be released soon
  - Do Lab 4 first, it is designed to help you get started
- Sign up for a check-in slot!
  - This is your chance to talk to course staff about anything you want

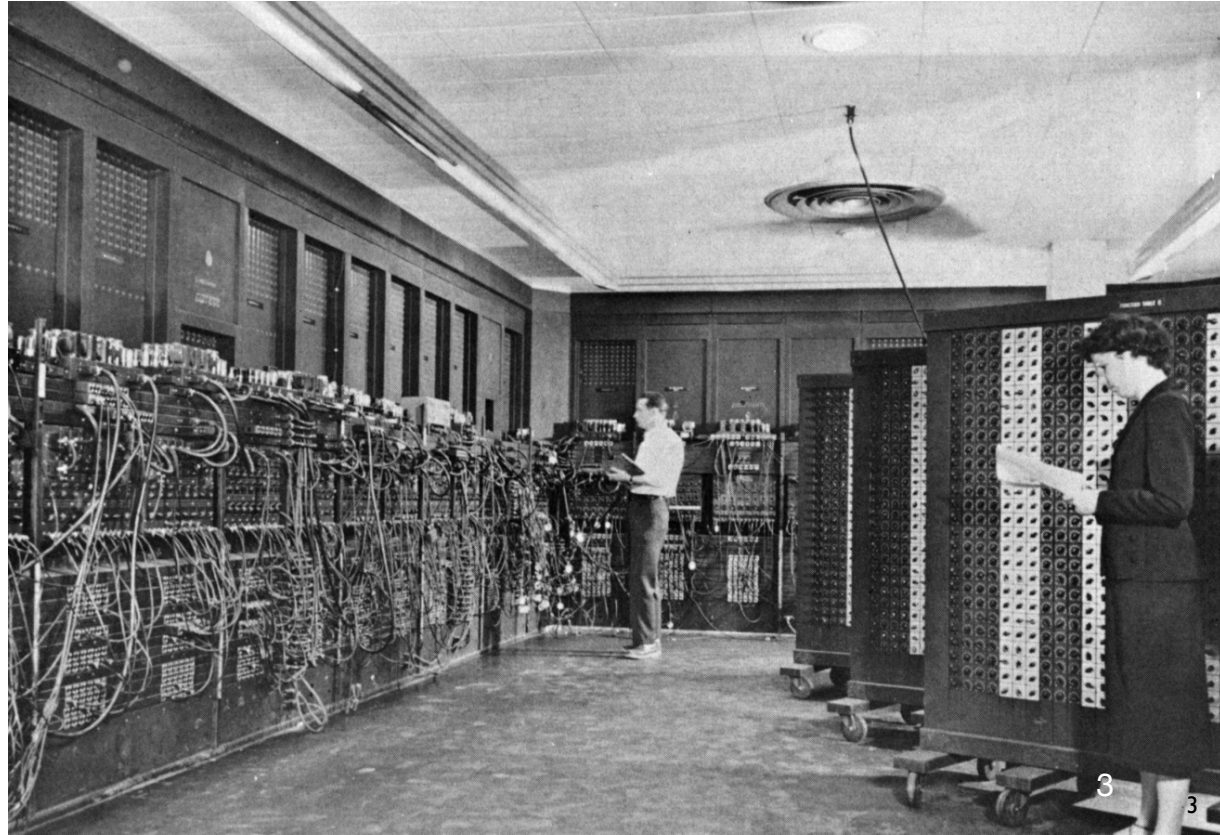
# ENIAC (U.Penn., 1946)

## First Electronic General-Purpose Computer

Computer Science 61C Spring 2021

Kolb and Weaver

- Blazingly fast  
(multiply in 2.8ms!)
  - 10 decimal digits x 10 decimal digits
- But needed 2-3 days to setup new program, as programmed with patch cords and switches



# Big Idea:

## Stored-Program Computer

- Instructions are represented as bit patterns - can think of these as numbers
- Therefore, entire programs can be stored in memory to be read or written just like data
- Can reprogram quickly (seconds), don't have to rewire computer (days)
- Known as the “von Neumann” computers after widely distributed tech report on EDVAC project
  - Wrote-up discussions of Eckert and Mauchly
  - Anticipated earlier by Turing and Zuse

First Draft of a Report on the EDVAC

by

John von Neumann

Contract No. W-670-ORD-4926

Between the

United States Army Ordnance Department

and the

University of Pennsylvania

Moore School of Electrical Engineering

University of Pennsylvania

June 30, 1945

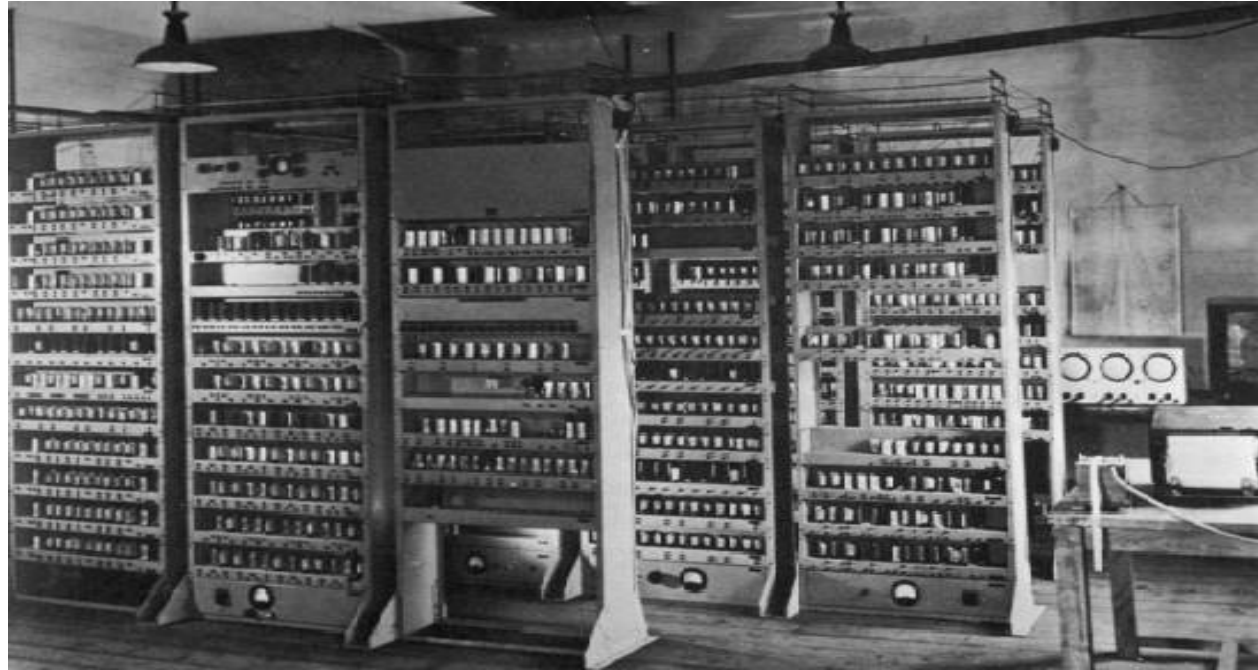
# EDSAC (Cambridge, 1949)

## First General Stored-Program Computer

Computer Science 61C Spring 2021

Kolb and Weaver

- Programs held as “numbers” in memory
- 35-bit binary 2’s complement words



# Consequence #1: Everything Has a Memory Address

- Since all instructions and data are stored in memory, everything has a memory address: instructions, data words
  - Both branches and jumps use these
- C pointers are just memory addresses: they can point to anything in memory
- One register keeps address of instruction being executed: “Program Counter” (**PC**)
  - Basically a pointer to memory
  - Intel calls it Instruction Pointer (a better name)

# Consequence #2: Binary Compatibility

- Programs are distributed in binary form
  - Programs bound to specific instruction set
  - Different version for phones and PCs, etc.
- New machines in the same family want to run old programs (“binaries”) as well as programs compiled to new instructions
- Leads to “backward-compatible” instruction set evolving over time
  - Selection of Intel 8088 in 1981 for 1st IBM PC is major reason latest PCs still use 80x86 instruction set; could still run program from 1981 PC today

# Instructions as Numbers (1/2)

- Most data we work with is in words (32-bit chunks):
  - Each register is a word
  - **lw** and **sw** both access memory one word at a time
- So how do we represent instructions?
  - Remember: Computer only represents 1s and 0s, so assembler string “**add x10, x11, x0**” is meaningless to hardware
  - RISC-V seeks simplicity: since data is in words, make instructions be fixed-size 32-bit words also
    - Same 32-bit instruction definitions used for RV32, RV64, RV128



# Instructions as Numbers (2/2)

- Divide 32-bit instruction word into “**fields**”
- Each field tells processor something about instruction
- We could define different set of fields for each instruction, but RISC-V seeks simplicity, so group possible instructions into six basic types of instruction formats:
  - **R-format** for register-register arithmetic/logical operations
  - **I-format** for register-immediate ALU operations and loads
  - **S-format** for stores
  - **B-format** for branches
  - **U-format** for 20-bit upper immediate instructions
  - **J-format** for jumps

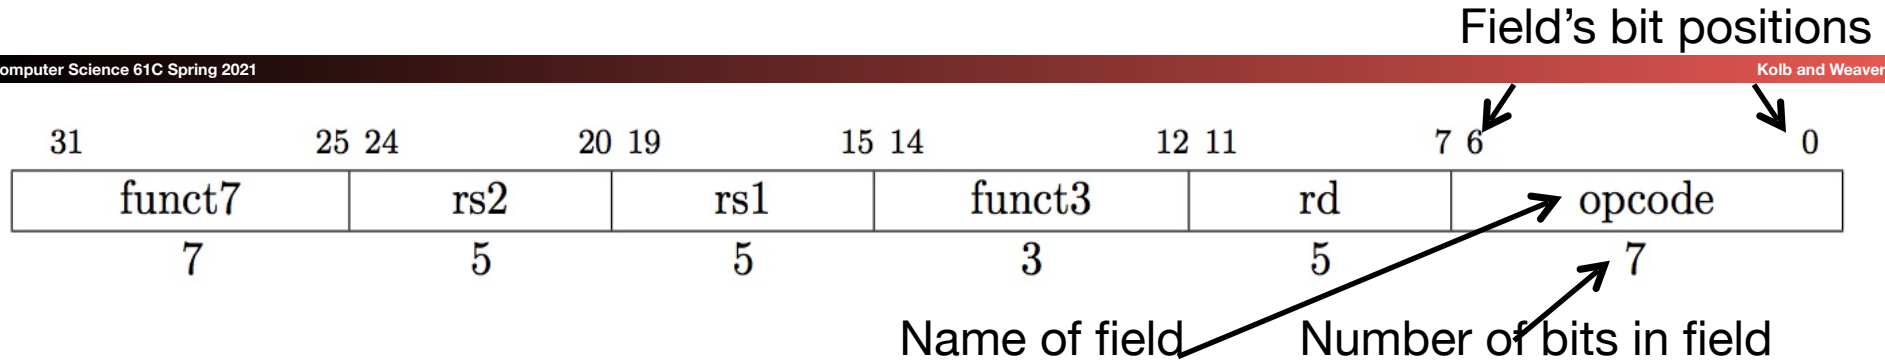
# Summary of RISC-V Instruction Formats

31	30	25	24	21	20	19	15	14	12	11	8	7	6	0				
funct7				rs2			rs1		funct3		rd			opcode		R-type		
imm[11:0]						rs1		funct3		rd			opcode		I-type			
imm[11:5]				rs2			rs1		funct3		imm[4:0]			opcode		S-type		
imm[12]		imm[10:5]			rs2			rs1		funct3		imm[4:1]		imm[11]		opcode		B-type
imm[31:12]										rd			opcode			U-type		
imm[20]		imm[10:1]				imm[11]		imm[19:12]				rd			opcode		J-type	

# R-Format Instruction Layout

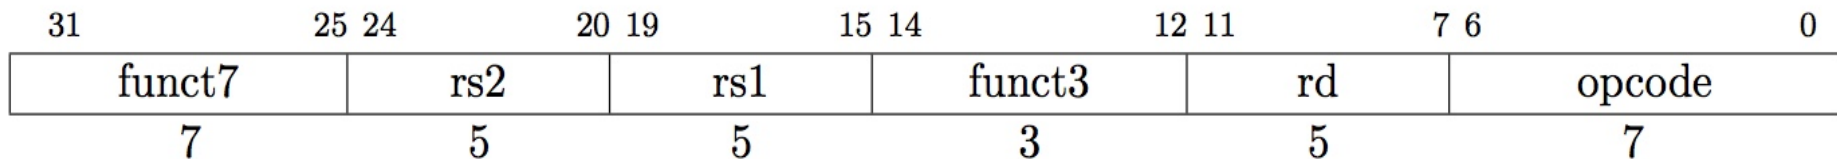
Computer Science 61C Spring 2021

Kolb and Weaver



- This example: 32-bit instruction word divided into six fields of varying numbers of bits each:  $7+5+5+3+5+7 = 32$
- In this case:
  - **opcode** is a 7-bit field that lives in bits 6-0 of the instruction
  - **rs2** is a 5-bit field that lives in bits 24-20 of the instruction
  - etc.

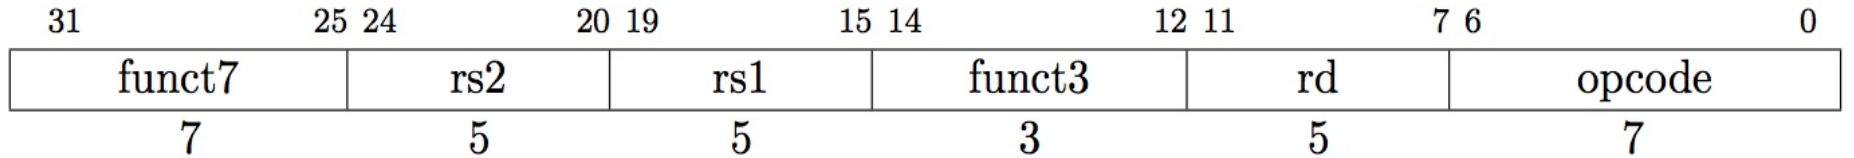
# R-Format Instructions opcode/funct fields



- **opcode**: partially specifies which instruction it is
  - Note: This field is equal to **0110011**<sub>two</sub> for all R-Format register-register arithmetic/logical instructions
- **funct7+funct3**: combined with **opcode**, these two fields describe what operation to perform
- Question: Why aren't **opcode** and **funct7** and **funct3** a single 17-bit field?

– We'll answer this later

# R-Format Instructions register specifiers

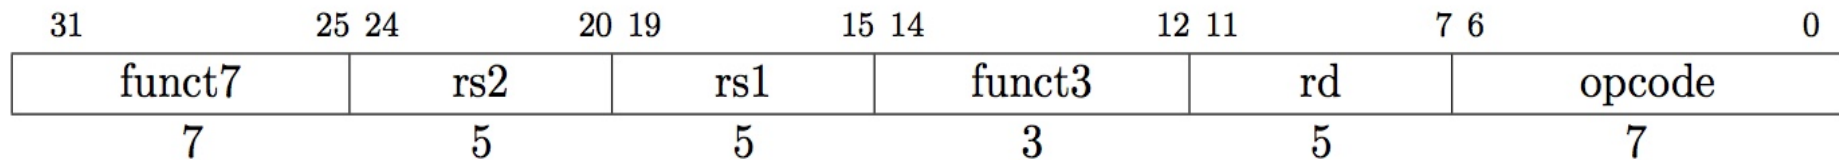


- Each register field (rs1, rs2, rd) holds a 5-bit unsigned integer (0-31) corresponding to a register number (**x0-x31**)
- rs1 (**S**ource **R**egister #1): specifies register containing first operand
- rs2 : specifies second register operand
- rd (**D**estination **R**egister): specifies register which will receive result of computation

# R-Format Example

- RISC-V Assembly Instruction:

**add x18,x19,x10**



ADD

rs2=10

rs1=19

ADD

rd=18

Reg-Reg OP

# All RV32 R-format instructions

funct7

funct3

opcode

0000000	rs2	rs1	000	rd	0110011	ADD
0100000	rs2	rs1	000	rd	0110011	SUB
0000000	rs2	rs1	001	rd	0110011	SLL
0000000	rs2	rs1	010	rd	0110011	SLT
0000000	rs2	rs1	011	rd	0110011	SLTU
0000000	rs2	rs1	100	rd	0110011	XOR
0000000	rs2	rs1	101	rd	0110011	SRL
0100000	rs2	rs1	101	rd	0110011	SRA
0000000	rs2	rs1	110	rd	0110011	OR
0000000	rs2	rs1	111	rd	0110011	AND



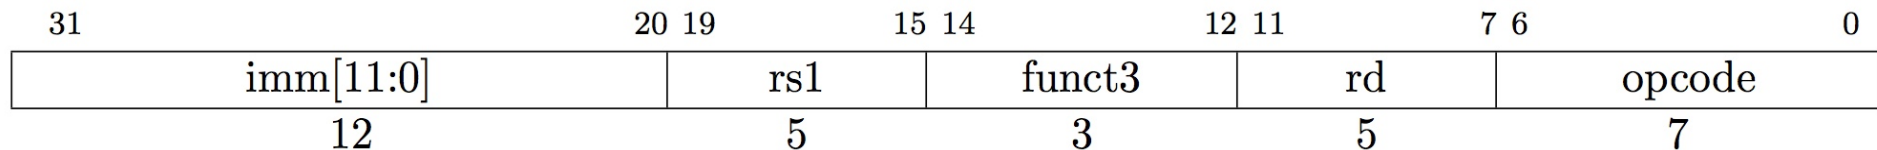
Encoding in funct7 + funct3 selects particular operation

# I-Format Instructions

- What about instructions with immediates?
  - Ideally, RISC-V would have only one instruction format (for simplicity): unfortunately, we need to compromise
  - 5-bit field only represents numbers up to the value 31: would like immediates to be much larger
- Define another instruction format that is mostly consistent with R-format
  - Note: if instruction has immediate, then uses at most 2 registers (one source, one destination)



# I-Format Instruction Layout

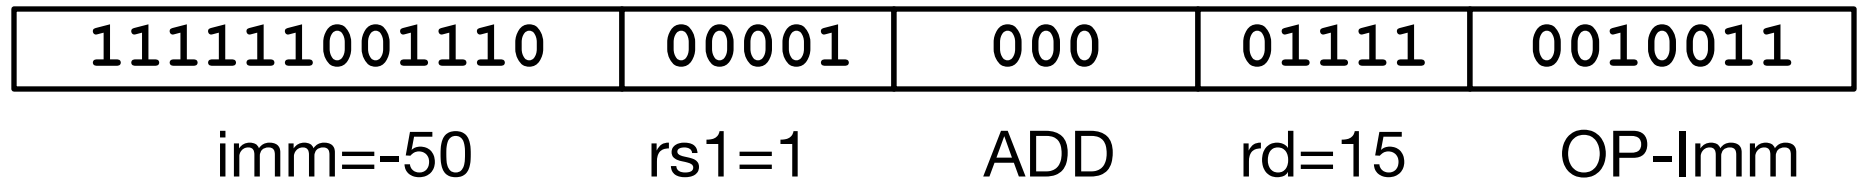
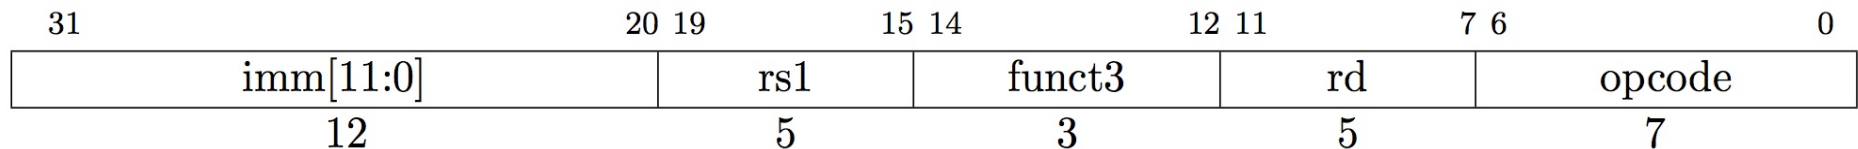


- Only one field is different from R-format, rs2 and funct7 replaced by 12-bit signed immediate, **imm[11:0]**
- Remaining field format (**rs1**, **funct3**, **rd**, **opcode**) same as before
- **imm[11:0]** can hold values in range  $[-2048_{\text{ten}}, +2047_{\text{ten}}]$
- Immediate is ***always*** sign-extended to 32-bits before use in an arithmetic operation
- We'll later see how to handle immediates  $> 12$  bits

# I-Format Example

- RISC-V Assembly Instruction:

**addi x15, x1, -50**



# All RV32 I-format Arithmetic/Logical Instructions

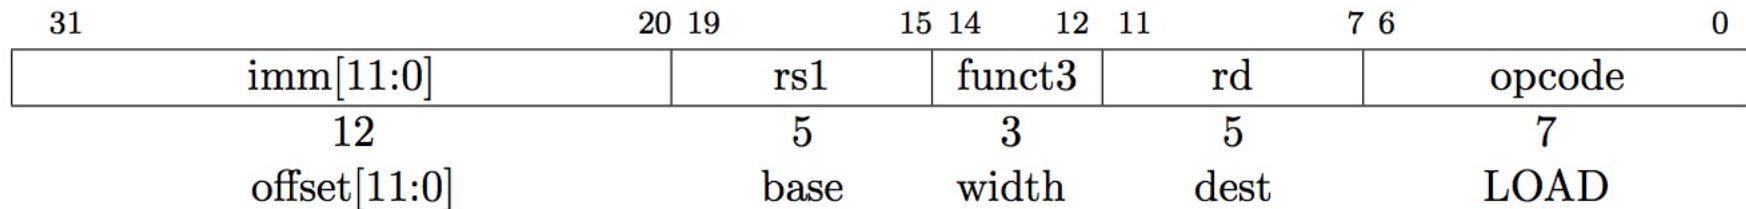
imm		funct3		opcode	
imm[11:0]		rs1	000	rd	0010011
imm[11:0]		rs1	010	rd	0010011
imm[11:0]		rs1	011	rd	0010011
imm[11:0]		rs1	100	rd	0010011
imm[11:0]		rs1	110	rd	0010011
imm[11:0]		rs1	111	rd	0010011
0000000	shamt	rs1	001	rd	0010011
0000000	shamt	rs1	101	rd	0010011
0100000	shamt	rs1	101	rd	0010011

ADDI  
SLTI  
SLTIU  
XORI  
ORI  
ANDI  
SLLI  
SRLI  
SRAI

One of the higher-order immediate bits is used to distinguish “shift right logical” (SRLI) from “shift right arithmetic” (SRAI)

“Shift-by-immediate” instructions only use lower 5 bits of the immediate value for shift amount (can only shift by 0-31 bit positions)

# Load Instructions are also I-Type

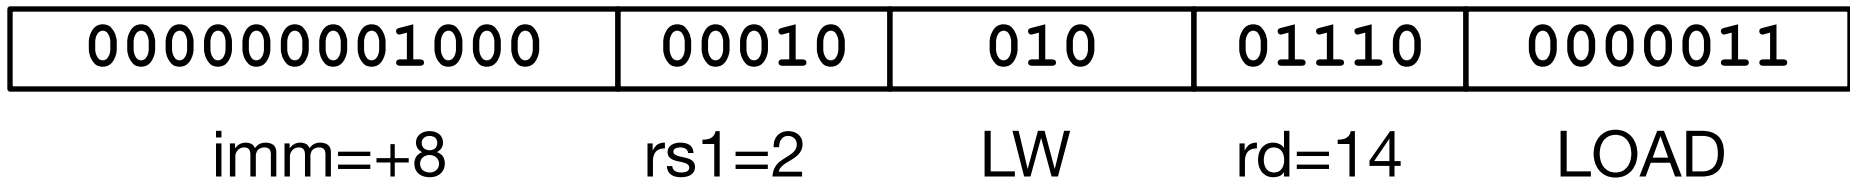
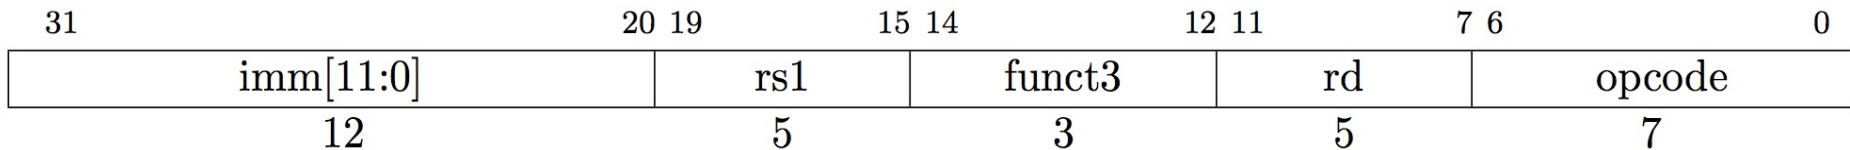


- The 12-bit signed immediate is added to the base address in register **rs1** to form the memory address
  - This is very similar to the add-immediate operation but used to create address not to create final result
- The value loaded from memory is stored in register **rd**

# I-Format Load Example

- RISC-V Assembly Instruction:

**lw x14, 8(x2)**



# All RV32 Load Instructions

Computer Science 61C Spring 2021

Kolb and Weaver

imm[11:0]	rs1	000	rd	0000011	LB
imm[11:0]	rs1	001	rd	0000011	LH
imm[11:0]	rs1	010	rd	0000011	LW
imm[11:0]	rs1	100	rd	0000011	LBU
imm[11:0]	rs1	101	rd	0000011	LHU

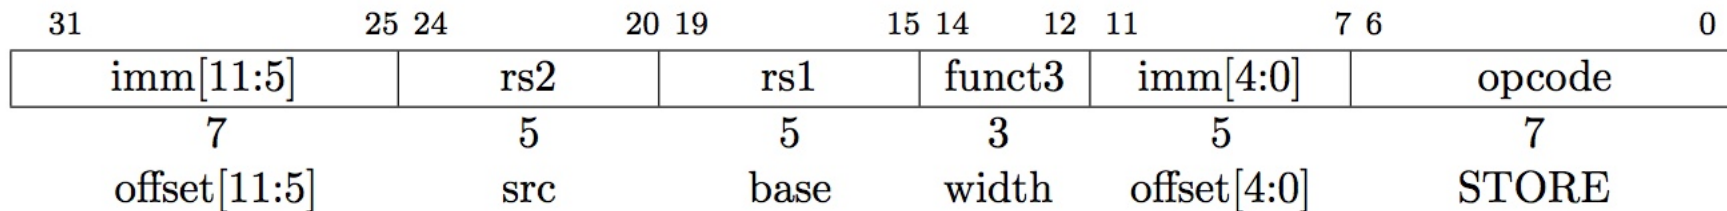
↑  
funct3 field encodes size  
and signedness of load  
data

- LBU is “load unsigned byte”
- LH is “load halfword”, which loads 16 bits (2 bytes) and sign-extends to fill destination 32-bit register
- LHU is “load unsigned halfword”, which zero-extends 16 bits to fill destination 32-bit register
- There is no LWU in RV32, because there is no sign/zero extension needed when copying 32 bits from a memory location into a 32-bit register

# S-Format Used for Stores

Compute

David Weaver



- Store needs to read two registers, **rs1** for base memory address, and **rs2** for data to be stored, as well as need immediate offset!
- Can't have both **rs2** and immediate in same place as other instructions!
- Note that stores don't write a value to the register file, **no rd!**
- RISC-V design decision is move low 5 bits of immediate to where rd field was in other instructions – keep **rs1/rs2** fields in same place.

# Keeping Registers always in the Same Place...

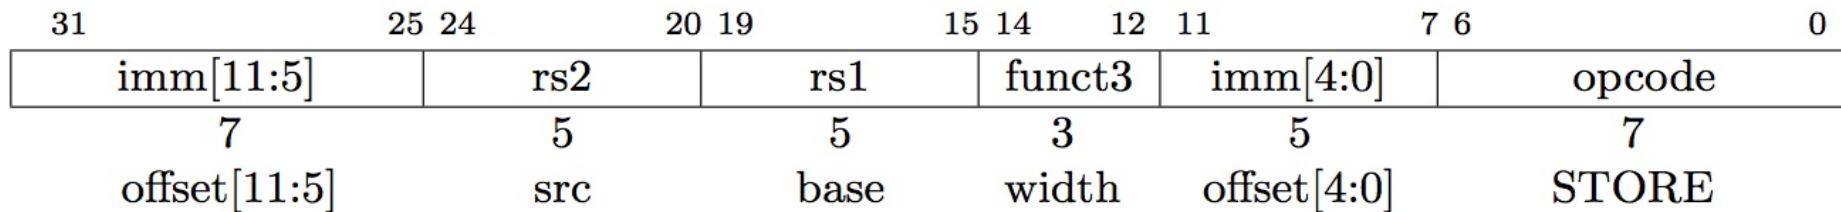
- The critical path for ***all operations*** includes fetching values from the registers
- By always placing the read sources in the same place, the register file can read without hesitation
  - If the data ends up being unnecessary (e.g. I-Type), it can be ignored
- Other RISCs have had slightly different encodings
  - Necessitating the logic to look at the instruction to determine which registers to read
- Example of one of the (many) little tweaks done in RISC-V to make things work better



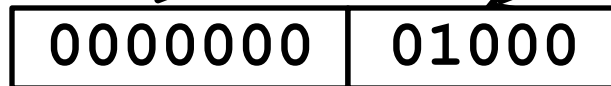
# S-Format Example

- RISC-V Assembly Instruction:

**sw x14, 8(x2)**



offset[11:5] rs2=14 rs1=2 SW offset[4:0] STORE  
=0 =8



combined 12-bit offset = 8

# All RV32 Store Instructions

imm[11:5]	rs2	rs1	000	imm[4:0]	0100011	SB
imm[11:5]	rs2	rs1	001	imm[4:0]	0100011	SH
imm[11:5]	rs2	rs1	010	imm[4:0]	0100011	SW

# RISC-V Conditional Branches

- E.g., **BEQ x1, x2, Label**
- Branches read two registers but don't write a register (similar to stores)
- How to encode label, i.e., where to branch to?

# Branching Instruction Usage

- Branches typically used for loops (`if-else`, `while`, `for`)
  - Loops are generally small ( $< 50$  instructions)
  - Function calls and unconditional jumps handled with jump instructions (J-Format)
- **Recall:** Instructions stored in a localized area of memory (Code/Text)
  - Largest branch distance limited by size of code
  - Address of current instruction stored in the program counter (PC)

# PC-Relative Addressing

- **PC-Relative Addressing:** Use the **immediate** field as a two's-complement *offset* relative to PC
  - Branches generally change the PC by a small amount
  - Could specify  $\pm 2^{11}$  addresses offset from the PC
- To improve the reach of a single branch instruction, *in principle*, could multiply the offset by four bytes before adding to PC (*instructions are 4 bytes and word aligned*).
- This would allow one branch instruction to reach  $\pm 2^{11} \times 32$ -bit instructions either side of PC
  - Four times greater reach than using byte offset
  - However ...

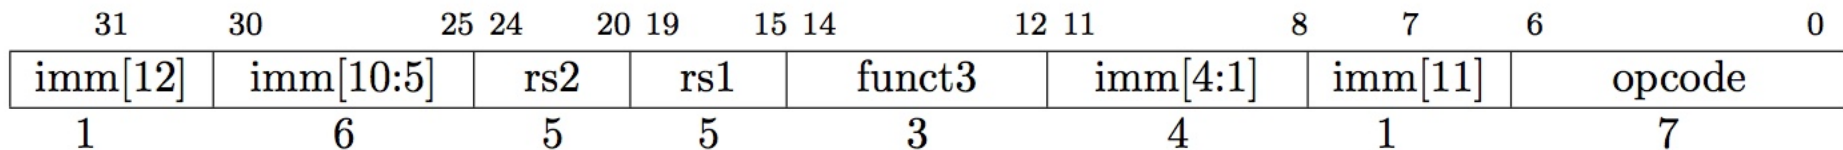
# RISC-V Feature, $n \times 16$ -bit instructions

- Extensions to RISC-V base ISA support 16-bit compressed instructions and also variable-length instructions that are multiples of 2-Bytes in length
- To enable this, RISC-V always scales the branch offset by 2 bytes - even when there are no 16-bit instructions
- (This means for us, the low bit of the stored immediate value will always be 0)
- Reduces branch reach by half:
- RISC-V conditional branches can only reach  $\pm 2^{10} \times 32$ -bit instructions either side of PC

# Branch Calculation

- If we **don't** take the branch:  
$$PC = PC + 4 \text{ (i.e., next instruction)}$$
- If we **do** take the branch:  
$$PC = PC + \text{immediate}$$
- **Observations:**
  - `immediate` is number of bytes to jump (remember, PC increments in words/half-words, so it has to be aligned) either forward (+) or backwards (−)

# RISC-V B-Format for Branches



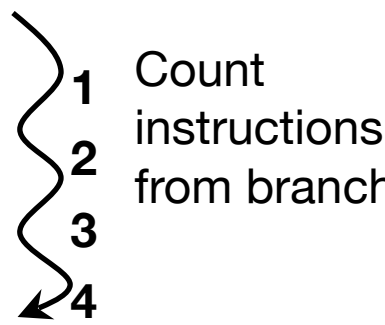
- B-format is mostly same as S-Format, with two register sources (**rs1/rs2**) and a 12-bit immediate
- But now immediate represents values -4096 to +4094 in 2-byte increments
- The 12 immediate bits encode *even* 13-bit signed byte offsets (lowest bit of offset is always zero, so no need to store it)
  - Thus the **imm[12:1]** in the total encoding, compared with **imm[11:0]** in the I-type encodings



# Branch Example, determine offset

- RISC-V Code:

```
Loop: beq x19, x10, End
      add x18, x18, x10
      addi x19, x19, -1
      j Loop
End: # target instruction
```



Count instructions from branch

- Branch offset =  $4 \times 32\text{-bit instructions} = 16 \text{ bytes}$
- (Branch with offset of 0, branches to itself)

# Branch Example, encode offset

- RISC-V Code:

```
Loop: beq    x19, x10, End
      add    x18, x18, x10
      addi   x19, x19, -1
      j      Loop
End: # target instruction
```

offset = 16 bytes = 8x2

???????	01010	10011	000	?????	1100011
imm	rs2=10	rs1=19	BEQ	imm	BRANCH

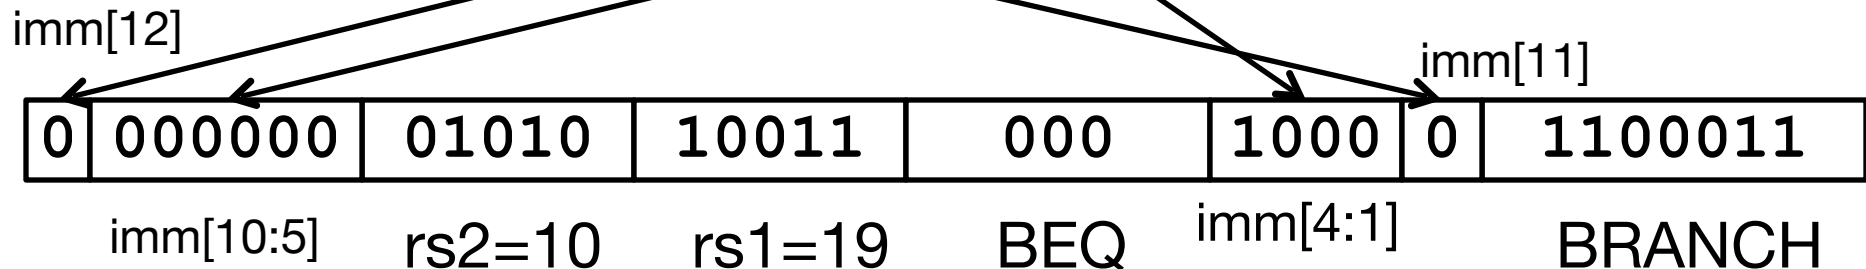
# Branch Example, complete encoding

**beq**    **x19,x10, offset = 16 bytes**

13-bit immediate, imm[12:0], with value 16

00000000010000

imm[0] discarded,  
→ always zero



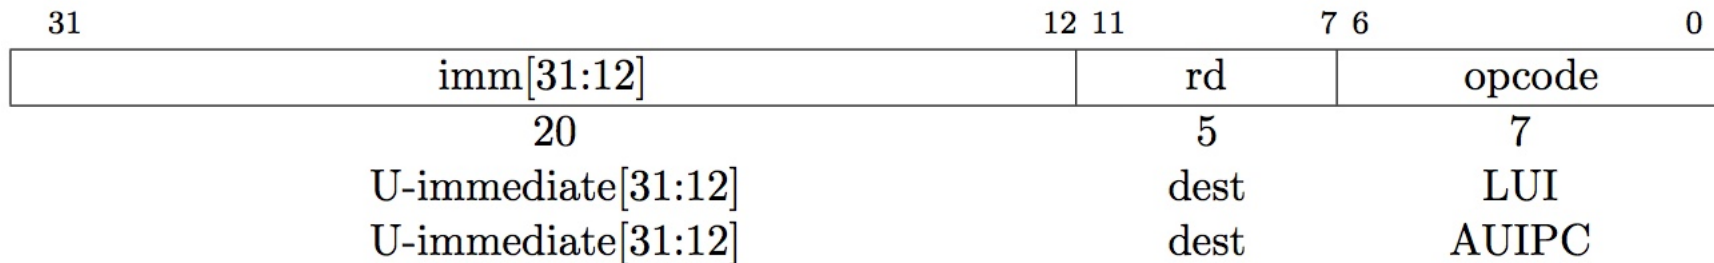
# All RISC-V Branch Instructions

imm[12 10:5]	rs2	rs1	000	imm[4:1 11]	1100011	BEQ
imm[12 10:5]	rs2	rs1	001	imm[4:1 11]	1100011	BNE
imm[12 10:5]	rs2	rs1	100	imm[4:1 11]	1100011	BLT
imm[12 10:5]	rs2	rs1	101	imm[4:1 11]	1100011	BGE
imm[12 10:5]	rs2	rs1	110	imm[4:1 11]	1100011	BLTU
imm[12 10:5]	rs2	rs1	111	imm[4:1 11]	1100011	BGEU

# Questions on PC-addressing

- Does the value in branch immediate field change if we move the code?
  - If moving individual lines of code, then yes
  - If moving all of code, then no (because PC-relative offsets)
- What do we do if destination is  $> 2^{10}$  instructions away from branch?
  - Other instructions save us
  - `beq x10,x0,far`                      `bne x10,x0,next`  
  `# next instr`                      `j far`  
   `next: # next instr`

# U-Format for “Upper Immediate” instructions



- Has 20-bit immediate in upper 20 bits of 32-bit instruction word
- One destination register, rd
- Used for two instructions
  - LUI – Load Upper Immediate
  - AUIPC – Add Upper Immediate to PC

# LUI to create long immediates

- LUI writes the upper 20 bits of the destination with the immediate value, and clears the lower 12 bits.
- Together with an ADDI to set low 12 bits, can create any 32-bit value in a register using two instructions (LUI/ADDI).

```
LUI x10, 0x87654      # x10 = 0x87654000
```

```
ADDI x10, x10, 0x321 # x10 = 0x87654321
```

# One Corner Case

How to set 0xDEADBEEF?

**LUI x10, 0xDEADB      # x10 = 0xDEADB000**

**ADDI x10, x10, 0xEEF # x10 = 0xDEAD~~A~~EEF**

ADDI 12-bit immediate is always sign-extended, if top bit is set, will subtract -1 from upper 20 bits



# Solution

How to set 0xDEADBEEF?

```
LUI x10, 0xDEADC      # x10 = 0xDEADC000
```

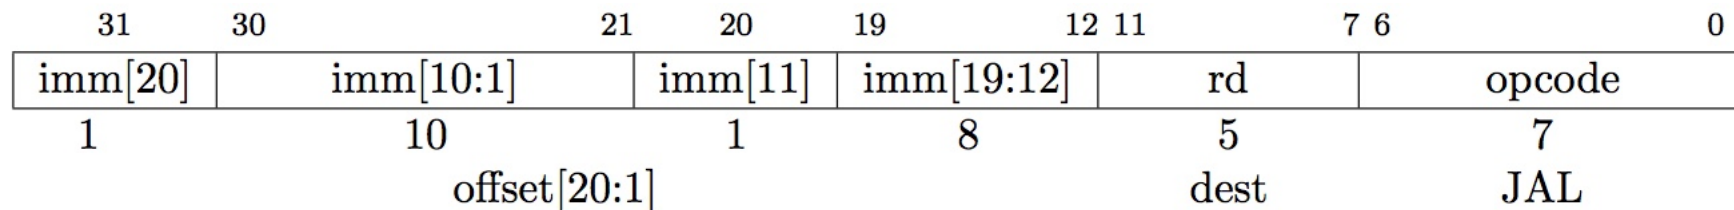
```
ADDI x10, x10, 0xEEF # x10 = 0xDEADBEEF
```

Pre-increment the value placed in upper 20 bits, if sign bit will be set on immediate in lower 12 bits.

Assembler pseudo-op handles all of this:

```
li x10, 0xDEADBEEF # Creates two instructions
```

# J-Format for Jump Instructions



- JAL saves PC+4 in register rd (the return address)
  - Assembler “j” jump is pseudo-instruction, uses JAL but sets rd=x0 to discard return address
- Set PC = PC + offset (PC-relative jump)
- Target somewhere within  $\pm 2^{19}$  locations, 2 bytes apart
  - $\pm 2^{18}$  32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost

# Uses of JAL

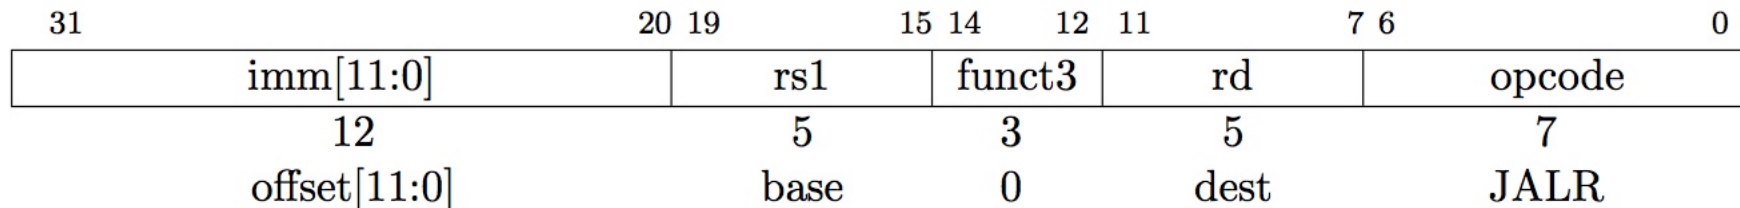
```
# j pseudo-instruction
```

```
j Label = jal x0, Label # Discard return address
```

```
# Call function within  $2^{18}$  instructions of PC
```

```
jal ra, FuncName
```

# JALR Instruction (I-Format)



- **JALR rd, rs, immediate**
  - Writes PC+4 to **rd** (return address)
  - Sets  $PC = rs + immediate$
  - Uses same immediates as arithmetic and loads
    - **no** multiplication by 2 bytes

# Uses of JALR

```
# ret and jr psuedo-instructions
```

```
ret = jr ra = jalr x0, ra, 0
```

```
# Call function at any 32-bit absolute address
```

```
lui  x1, <hi20bits>
```

```
jalr ra, x1, <lo12bits>
```

```
# Jump PC-relative with 32-bit offset
```

```
auipc x1, <hi20bits>      # Adds upper immediate value to  
                           # and places result in x1
```

```
jalr x0, x1, <lo12bits>  # Same sign extension trick needed  
                           # as LUI
```

# Summary of RISC-V Instruction Formats

31	30	25	24	21	20	19	15	14	12	11	8	7	6	0				
funct7				rs2			rs1		funct3		rd			opcode		R-type		
imm[11:0]						rs1		funct3		rd			opcode		I-type			
imm[11:5]				rs2			rs1		funct3		imm[4:0]			opcode		S-type		
imm[12]		imm[10:5]			rs2			rs1		funct3		imm[4:1]		imm[11]		opcode		B-type
imm[31:12]									rd			opcode			U-type			
imm[20]		imm[10:1]				imm[11]		imm[19:12]			rd			opcode		J-type		

# Complete RV32I ISA

Computer Science 61C Spring 2021

Kolb and Weaver

imm[31:12]				rd	0110111
imm[31:12]				rd	0010111
imm[20:10:11:19:12]				rd	1101111
imm[11:0]				rd	1100111
imm[12:10:5]	rs2	rs1	000	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	001	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	100	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	101	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	110	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	111	imm[4:1:11]	1100011
imm[11:0]				rd	0000011
imm[11:0]				rd	0000011
imm[11:0]				rd	0000011
imm[11:0]				rd	0000011
imm[11:0]				rd	0000011
imm[11:5]	rs2	rs1	000	imm[4:0]	0100011
imm[11:5]	rs2	rs1	001	imm[4:0]	0100011
imm[11:5]	rs2	rs1	010	imm[4:0]	0100011
imm[11:0]				rd	0010011
imm[11:0]				rd	0010011
imm[11:0]				rd	0010011
imm[11:0]				rd	0010011
imm[11:0]				rd	0010011
imm[11:0]				rd	0010011

LUI  
 AUIPC  
 JAL  
 JALR  
 BEQ  
 BNE  
 BLT  
 BGE  
 BLTU  
 BGEU  
 LB  
 LH  
 LW  
 LBU  
 LHU  
 SB  
 SH  
 SW  
 ADDI  
 SLTI  
 SLTIU  
 XORI  
 ORI  
 ANDI

0000000	shamt	rs1	001	rd	0010011
0000000	shamt	rs1	101	rd	0010011
0100000	shamt	rs1	101	rd	0010011
0000000	rs2	rs1	000	rd	0110011
0100000	rs2	rs1	000	rd	0110011
0000000	rs2	rs1	001	rd	0110011
0000000	rs2	rs1	010	rd	0110011
0000000	rs2	rs1	011	rd	0110011
0000000	rs2	rs1	100	rd	0110011
0000000	rs2	rs1	101	rd	0110011
0100000	rs2	rs1	101	rd	0110011
0000000	rs2	rs1	110	rd	0110011
0000000	rs2	rs1	111	rd	0110011

SLLI  
 SRLI  
 SRAI  
 ADD  
 SUB  
 SLL  
 SLT  
 SLTU  
 XOR  
 SRL  
 SRA  
 OR  
 AND  
 FENCE  
 FENCE.I  
 ECALL  
 EBREAK  
 CSRRW  
 CSRRS  
 CSRRC  
 CSRRWI  
 CSRRSI  
 CSRRCI

0000	pred	succ	00000	000	00000	0001111
0000	0000	0000	00000	001	00000	0001111
0000000000000			00000	000	00000	1110011
0000000000001			00000	000	00000	1110011
csr	rs1	001	rd	1110011		
csr	rs1	010	rd	1110011		
csr	rs1	011	rd	1110011		
csr	zimm	101	rd	1110011		
csr	zimm	110	rd	1110011		
csr	zimm	111	rd	1110011		

Not in CS61C