

# ***RISC-V Processor Datapath***

# Administrivia

- Lecture quizzes will now be due on Saturday each week
- Project 2 Deadlines
  - Part A due this Thursday (2/25)
  - Part B due next Thursday (3/4)
- Project 1 Clobbering
  - Will use the project1-practice autograder
  - Multiply your score by 0.7
  - You must submit test cases for consideration

# Great Idea #1: Abstraction

## (Levels of Representation/Interpretation)

```
lw  t0, t2, 0
lw  t1, t2, 4
sw  t1, t2, 0
sw  t0, t2, 4
```

High Level Language  
Program (e.g., C)

*Compiler*

Assembly Language  
Program (e.g., RISC-V)

*Assembler*


Machine Language  
Program (RISC-V)

```
temp = v[k];
v[k] = v[k+1];
v[k+1] = temp;
```

Anything can be represented  
as a *number*,  
i.e., data or instructions


```
0000 1001 1100 0110 1010 1111 0101 1000
1010 1111 0101 1000 0000 1001 1100 0110
1100 0110 1010 1111 0101 1000 0000 1001
0101 1000 0000 1001 1100 0110 1010 1111
```

***We are here!***



*Architecture  
Implementation*

Logic Circuit Description  
(Circuit Schematic Diagrams)



# Recap: Complete RV32I ISA

imm[31:12]				rd	0110111
imm[31:12]				rd	0010111
imm[20:10:11:19:12]				rd	1101111
imm[11:0]				rd	1100111
imm[12:10:5]	rs2	rs1	000	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	001	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	100	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	101	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	110	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	111	imm[4:1:11]	1100011
imm[11:0]				rd	0000011
imm[11:0]				rd	0000011
imm[11:0]				rd	0000011
imm[11:0]				rd	0000011
imm[11:0]				rd	0000011
imm[11:5]	rs2	rs1	000	imm[4:0]	0100011
imm[11:5]	rs2	rs1	001	imm[4:0]	0100011
imm[11:5]	rs2	rs1	010	imm[4:0]	0100011
imm[11:0]				rd	0010011
imm[11:0]				rd	0010011
imm[11:0]				rd	0010011
imm[11:0]				rd	0010011
imm[11:0]				rd	0010011
imm[11:0]				rd	0010011
imm[11:0]				rd	0010011

LUI  
AUIPC  
JAL  
JALR  
BEQ  
BNE  
BLT  
BGE  
BLTU  
BGEU  
LB  
LH  
LW  
LBU  
LHU  
SB  
SH  
SW  
ADDI  
SLTI  
SLTIU  
XORI  
ORI  
ANDI

0000000	shamt	rs1	001	rd	0010011
0000000	shamt	rs1	101	rd	0010011
0100000	shamt	rs1	101	rd	0010011
0000000	rs2	rs1	000	rd	0110011
0100000	rs2	rs1	000	rd	0110011
0000000	rs2	rs1	001	rd	0110011
0000000	rs2	rs1	010	rd	0110011
0000000	rs2	rs1	011	rd	0110011
0000000	rs2	rs1	100	rd	0110011
0000000	rs2	rs1	101	rd	0110011
0100000	rs2	rs1	101	rd	0110011
0000000	rs2	rs1	110	rd	0110011
0000000	rs2	rs1	111	rd	0110011

Not in CS61C

0000	pred	succ	00000	000	00000	0001111
0000	0000	0000	00000	001	00000	0001111
0000000000000			00000	000	00000	1110011
0000000000001			00000	000	00000	1110011
csr	rs1	001	rd	1110011		
csr	rs1	010	rd	1110011		
csr	rs1	011	rd	1110011		
csr	zimm	101	rd	1110011		
csr	zimm	110	rd	1110011		
csr	zimm	111	rd	1110011		

SLLI  
SRLI  
SRAI  
ADD  
SUB  
SLL  
SLT  
SLTU  
XOR  
SRL  
SRA  
OR  
AND  
FENCE  
FENCE.I  
ECALL  
EBREAK  
CSR  
CSR.RW  
CSR.RS  
CSR.RC  
CSR.RWI  
CSR.RSI  
CSR.RCI

# “State” Required by RV32I ISA

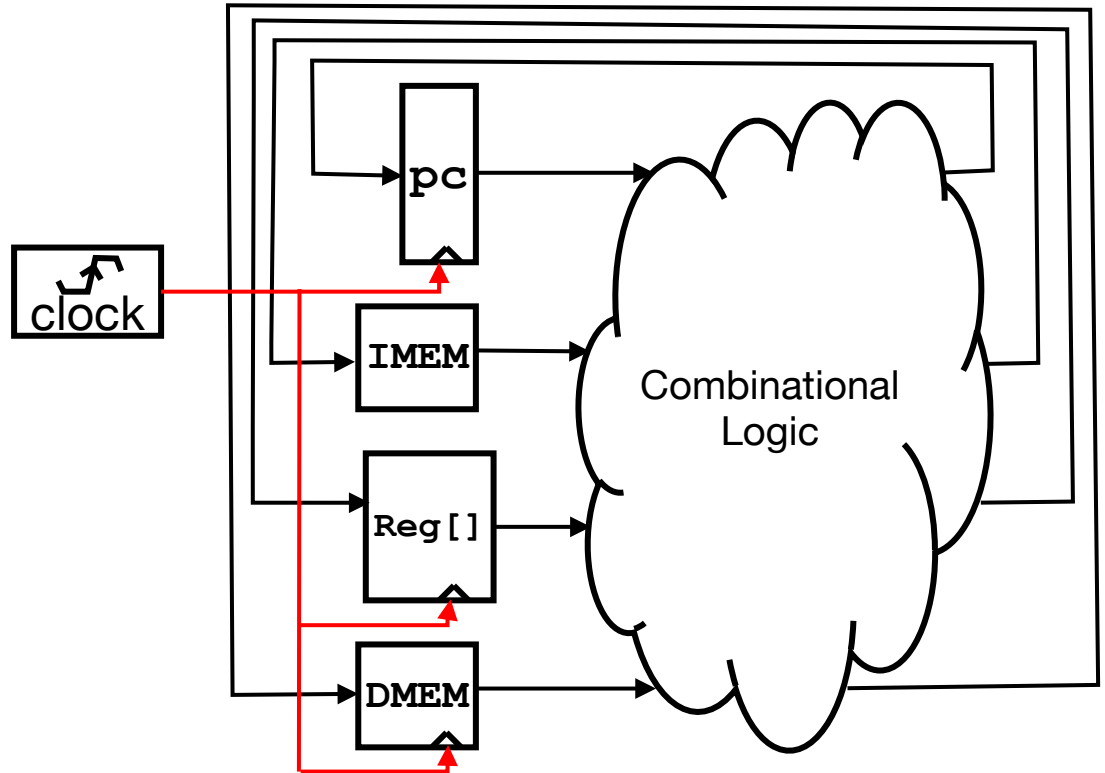
Each instruction reads and updates this state during execution:

- Registers (**x0** . . **x31**)
  - Register file (or *regfile*) **Reg** holds 32 registers x 32 bits/register: **Reg**[0] . . **Reg**[31]
  - First register read specified by *rs1* field in instruction
  - Second register read specified by *rs2* field in instruction
  - Write register (destination) specified by *rd* field in instruction
  - **x0** is always 0 (writes to **Reg**[0] are ignored)
- Program Counter (**PC**)
  - Holds address of current instruction
- Memory (**MEM**)
  - Holds both instructions & data, in one 32-bit byte-addressed memory space
  - We'll use separate memories for instructions (**IMEM**) and data (**DMEM**)
    - *Later we'll replace these with instruction and data caches*
  - Instructions are read (*fetch*) from instruction memory (assume **IMEM** read-only)
  - Load/store instructions access data memory

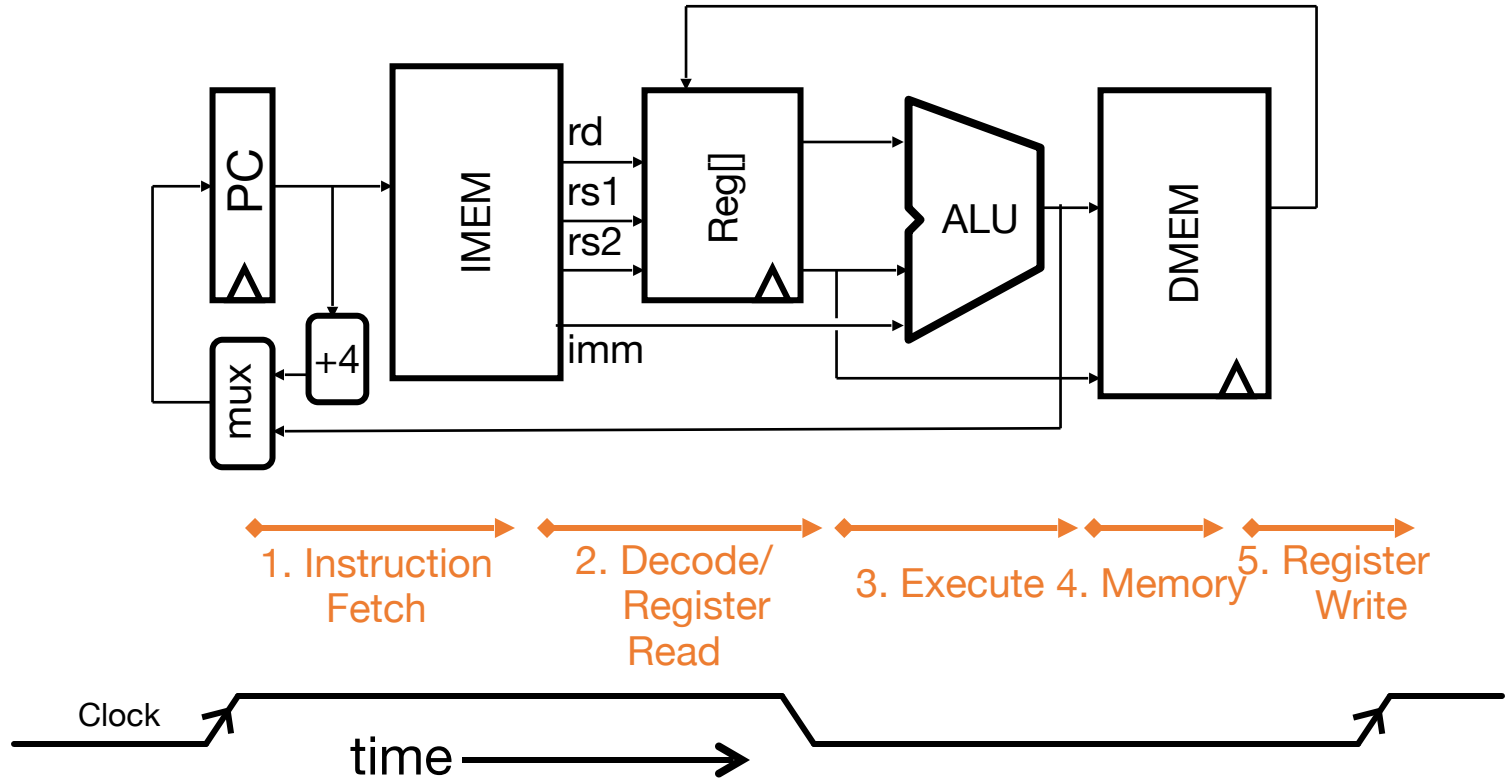
# One-Instruction-Per-Cycle RISC-V Machine

On every tick of the clock, the computer executes one instruction

1. Current state outputs drive the inputs to the combinational logic, whose outputs settles at the values of the state before the next clock edge
2. At the rising clock edge, all the state elements are updated with the combinational logic outputs, and execution moves to the next clock cycle
3. Separate instruction/data memory:  
For simplification, memory is asynchronous read (not clocked), but synchronous write (is clocked)



# Basic Phases of Instruction Execution



# Implementing the `add` instruction

0000000	rs2	rs1	000	rd	0110011	ADD
---------	-----	-----	-----	----	---------	-----

`add rd, rs1, rs2`

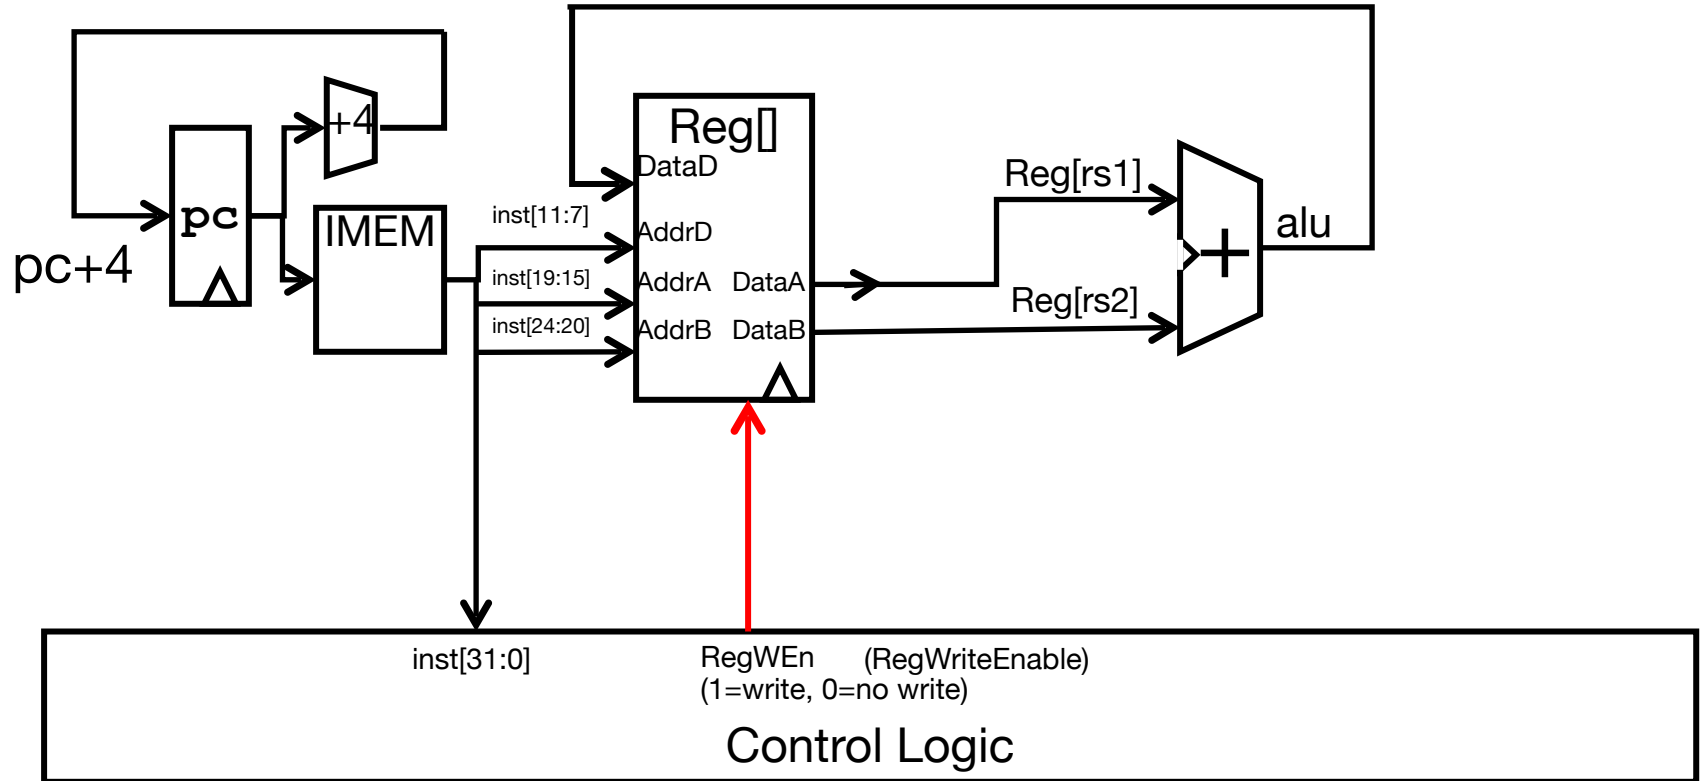
- Instruction makes two changes to machine's state:

`Reg[rd] = Reg[rs1] + Reg[rs2]`

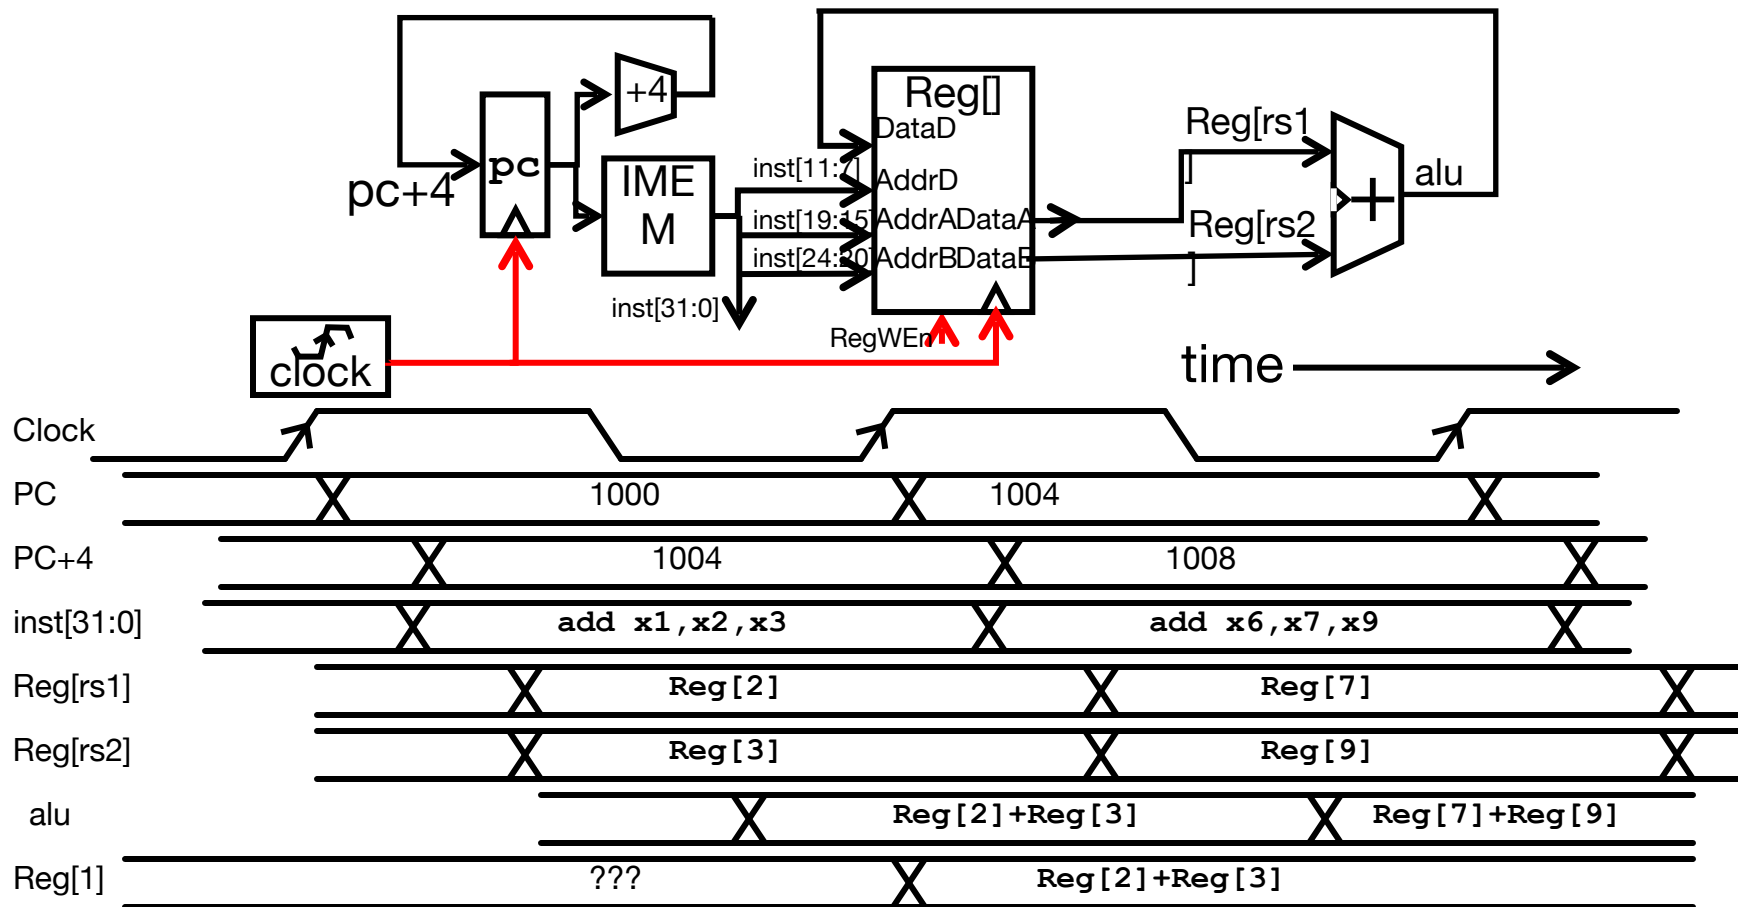
`PC = PC + 4`



# Datapath for add



# Timing Diagram for add

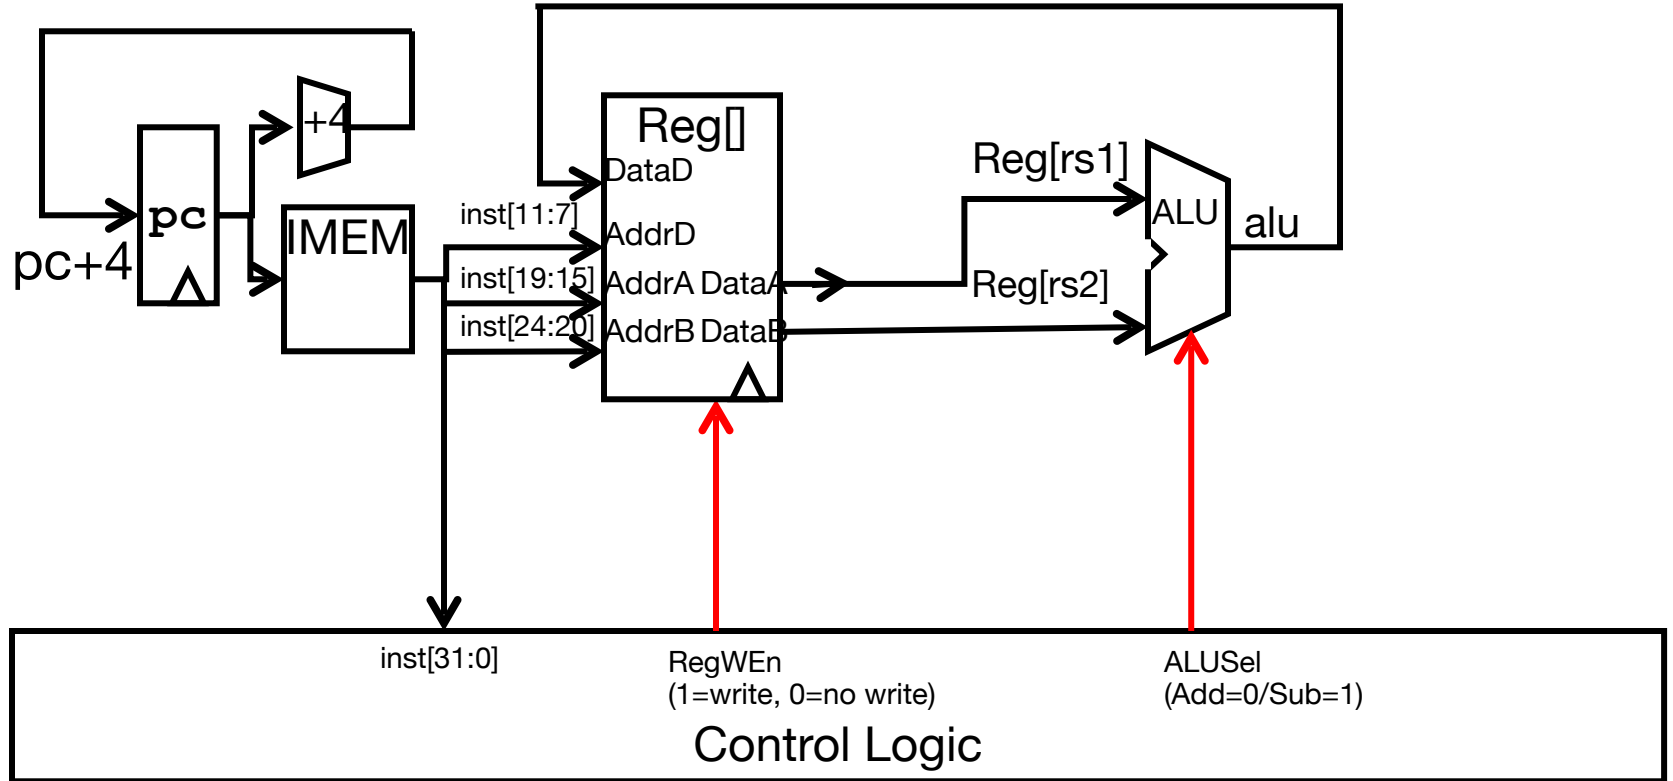


# Implementing the **sub** instruction

0000000	rs2	rs1	000	rd	0110011	ADD
0100000	rs2	rs1	000	rd	0110011	SUB

- Almost the same as add, except now have to subtract operands instead of adding them
- **inst**[30] selects between add and subtract

# Datapath for add/sub




# Implementing other R-Format instructions

0000000	rs2	rs1	000	rd	0110011	ADD
0100000	rs2	rs1	000	rd	0110011	SUB
0000000	rs2	rs1	001	rd	0110011	SLL
0000000	rs2	rs1	010	rd	0110011	SLT
0000000	rs2	rs1	011	rd	0110011	SLTU
0000000	rs2	rs1	100	rd	0110011	XOR
0000000	rs2	rs1	101	rd	0110011	SRL
0100000	rs2	rs1	101	rd	0110011	SRA
0000000	rs2	rs1	110	rd	0110011	OR
0000000	rs2	rs1	111	rd	0110011	AND

# Implementing the **addi** instruction

- RISC-V Assembly Instruction:

**addi    x15, x1, -50**



imm=-50

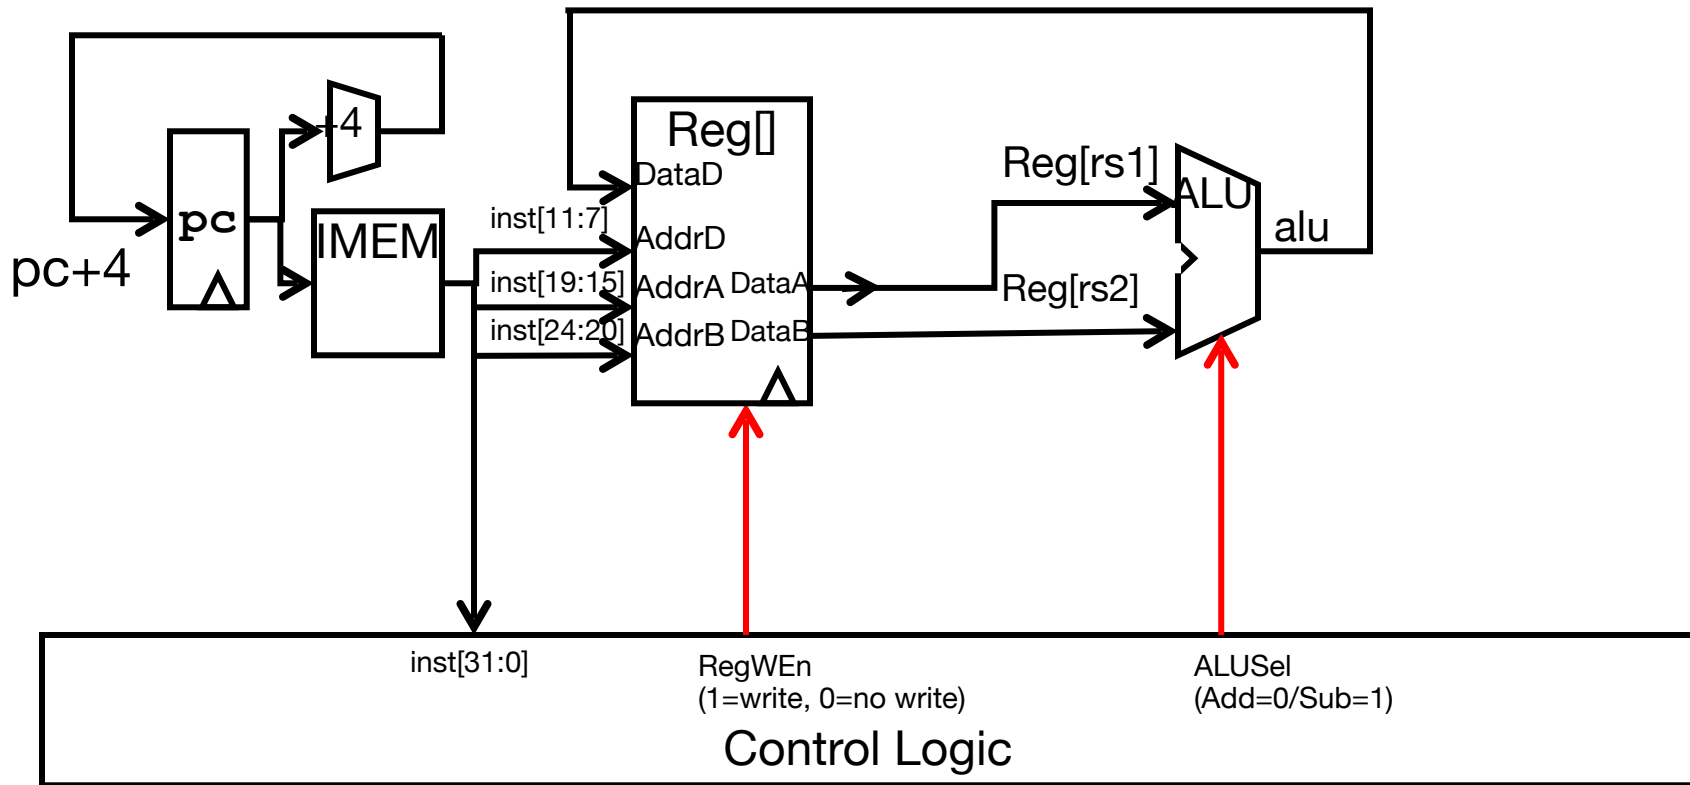
rs1=1

ADD

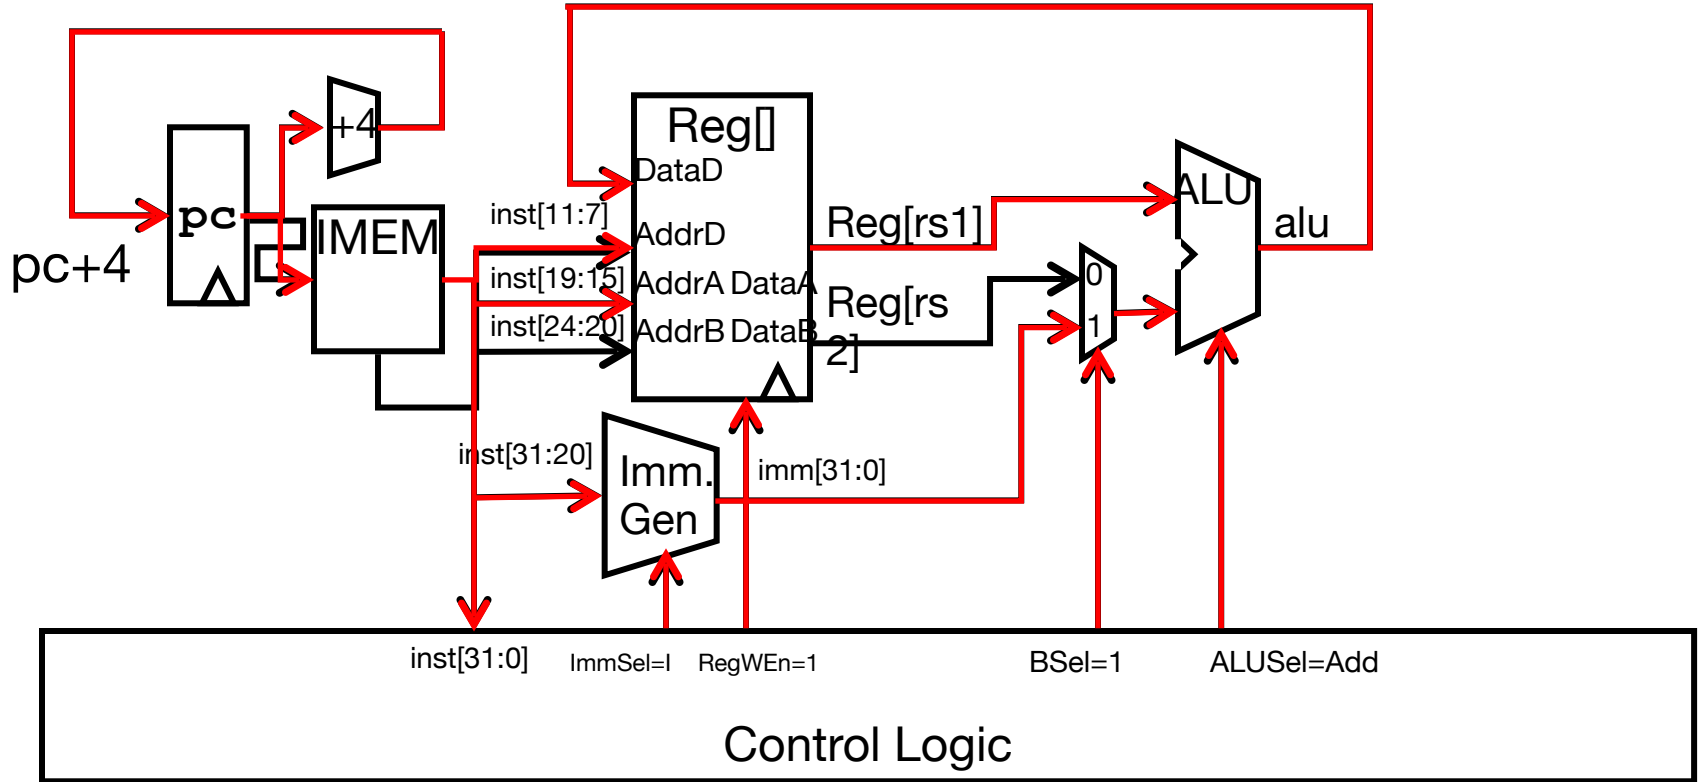
rd=15

OP-Imm

# Datapath for add/sub


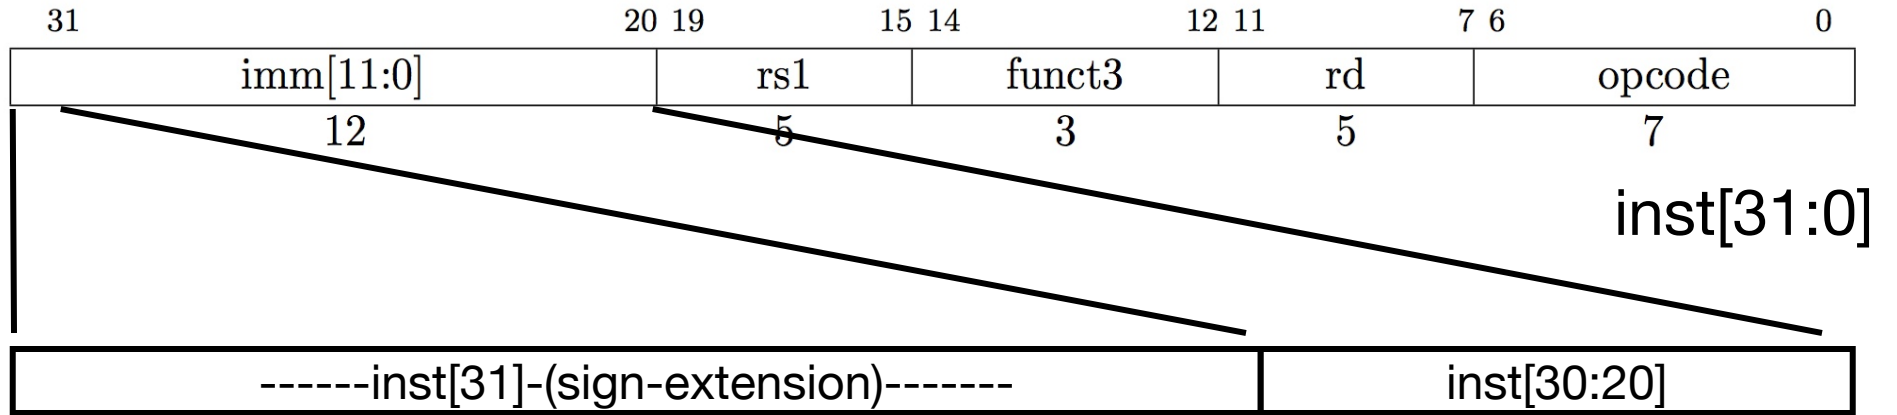


# Adding `addi` to datapath



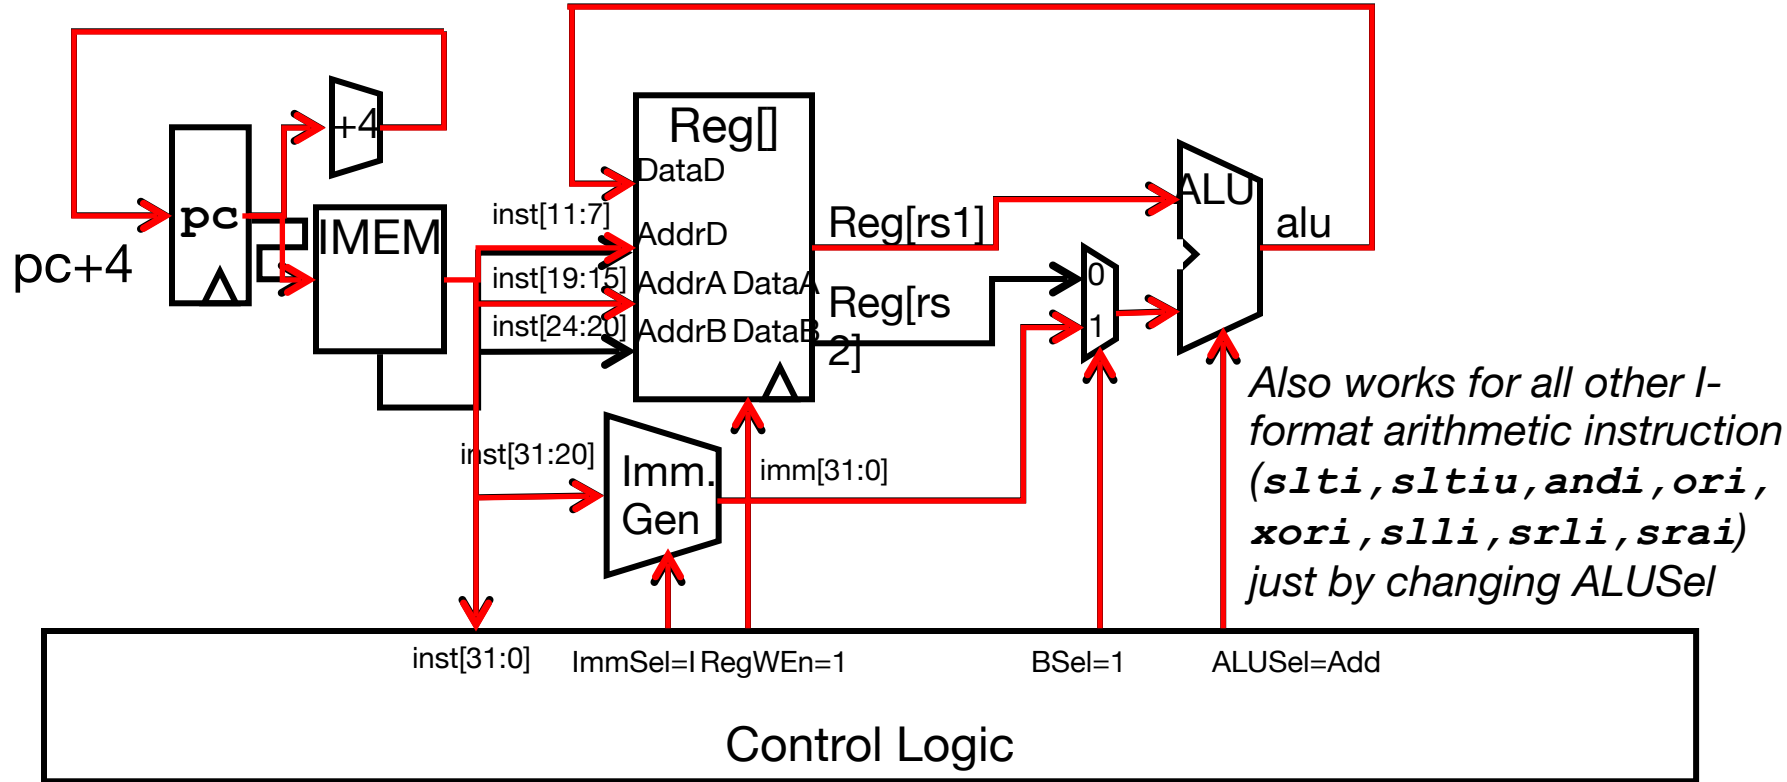


# I-Format immediates



- High 12 bits of instruction (inst[31:20]) copied to low 12 bits of immediate (imm[11:0])
- Immediate is sign-extended by copying value of inst[31] to fill the upper 20 bits of the immediate value (imm[31:12])



# Adding `addi` to datapath



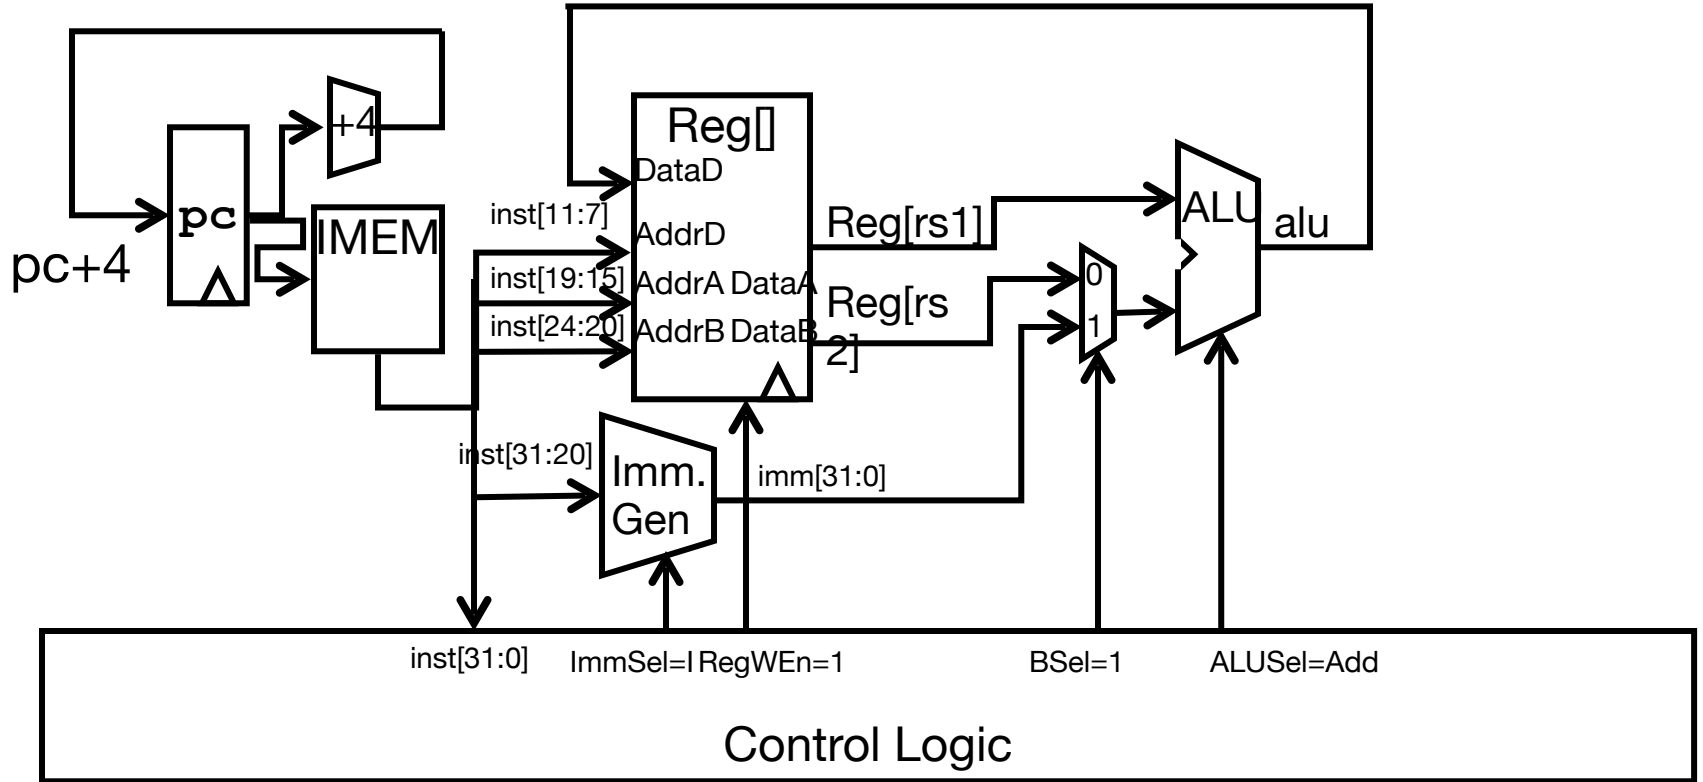
# Implementing Load Word instruction

- RISC-V Assembly Instruction:

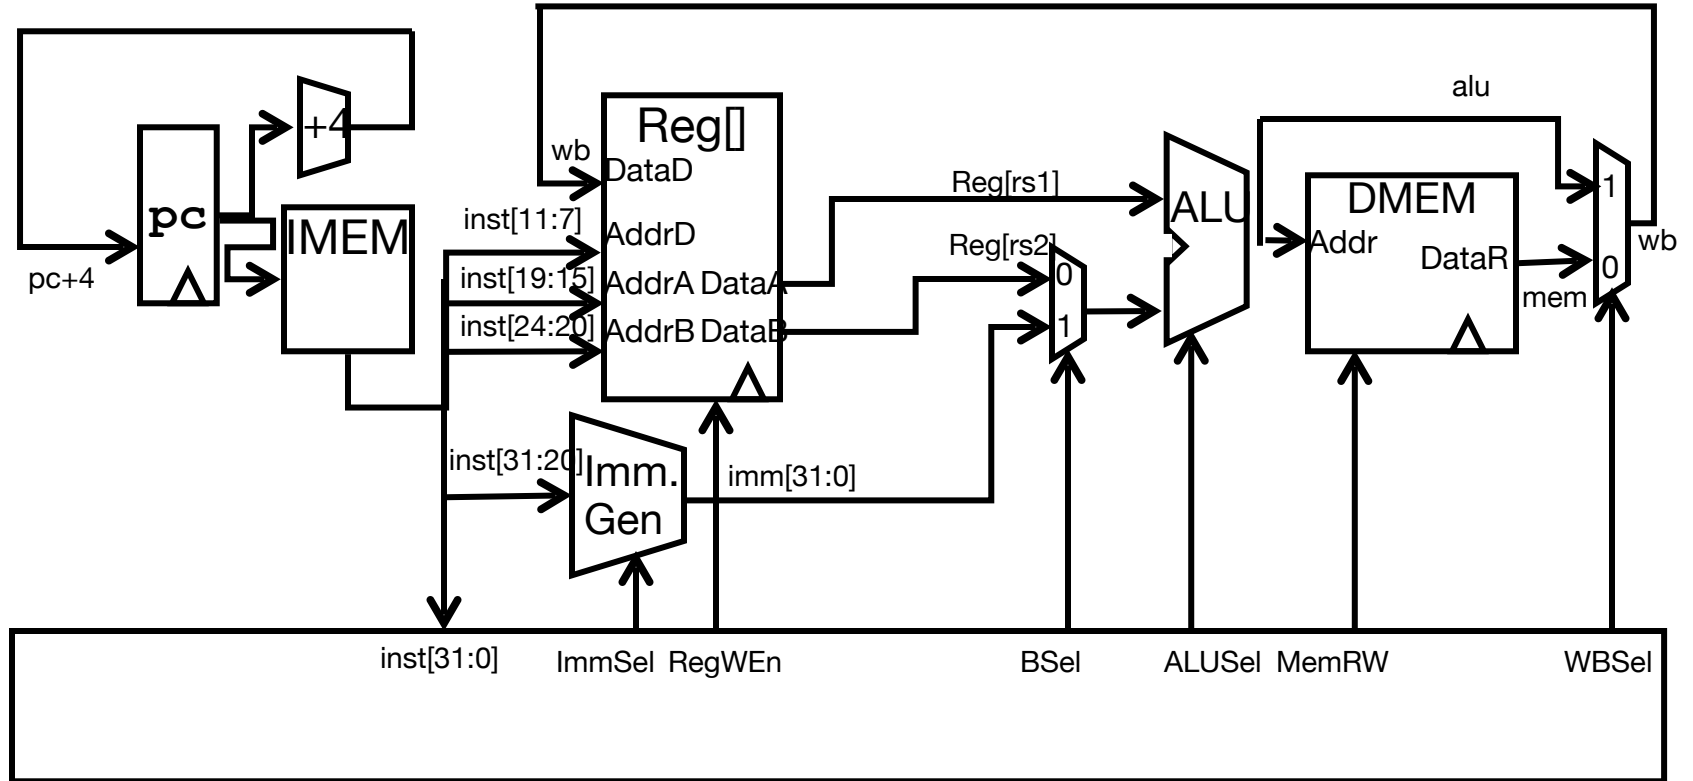
**lw x14, 8(x2)**



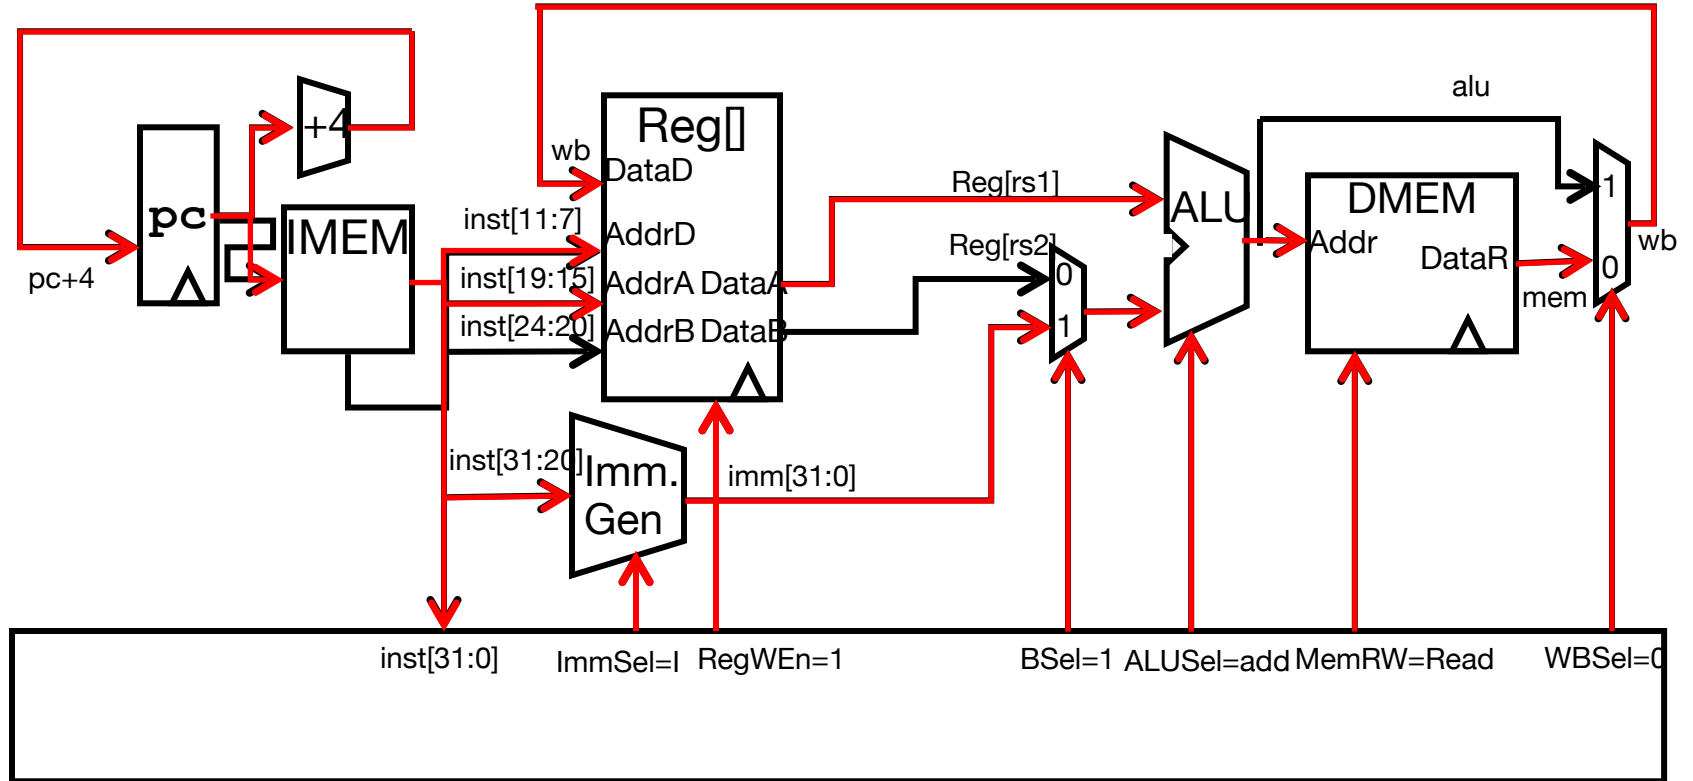
# Adding `addi` to datapath



# Adding 1w to datapath



# Adding 1w to datapath



# All RV32 Load Instructions

imm[11:0]	rs1	000	rd	0000011	LB
imm[11:0]	rs1	001	rd	0000011	LH
imm[11:0]	rs1	010	rd	0000011	LW
imm[11:0]	rs1	100	rd	0000011	LBU
imm[11:0]	rs1	101	rd	0000011	LHU


writing back to register file.

↑ funct3 field encodes size  
and signedness of load  
data


# Implementing Store Word instruction

- RISC-V Assembly Instruction:

**sw x14, 8(x2)**



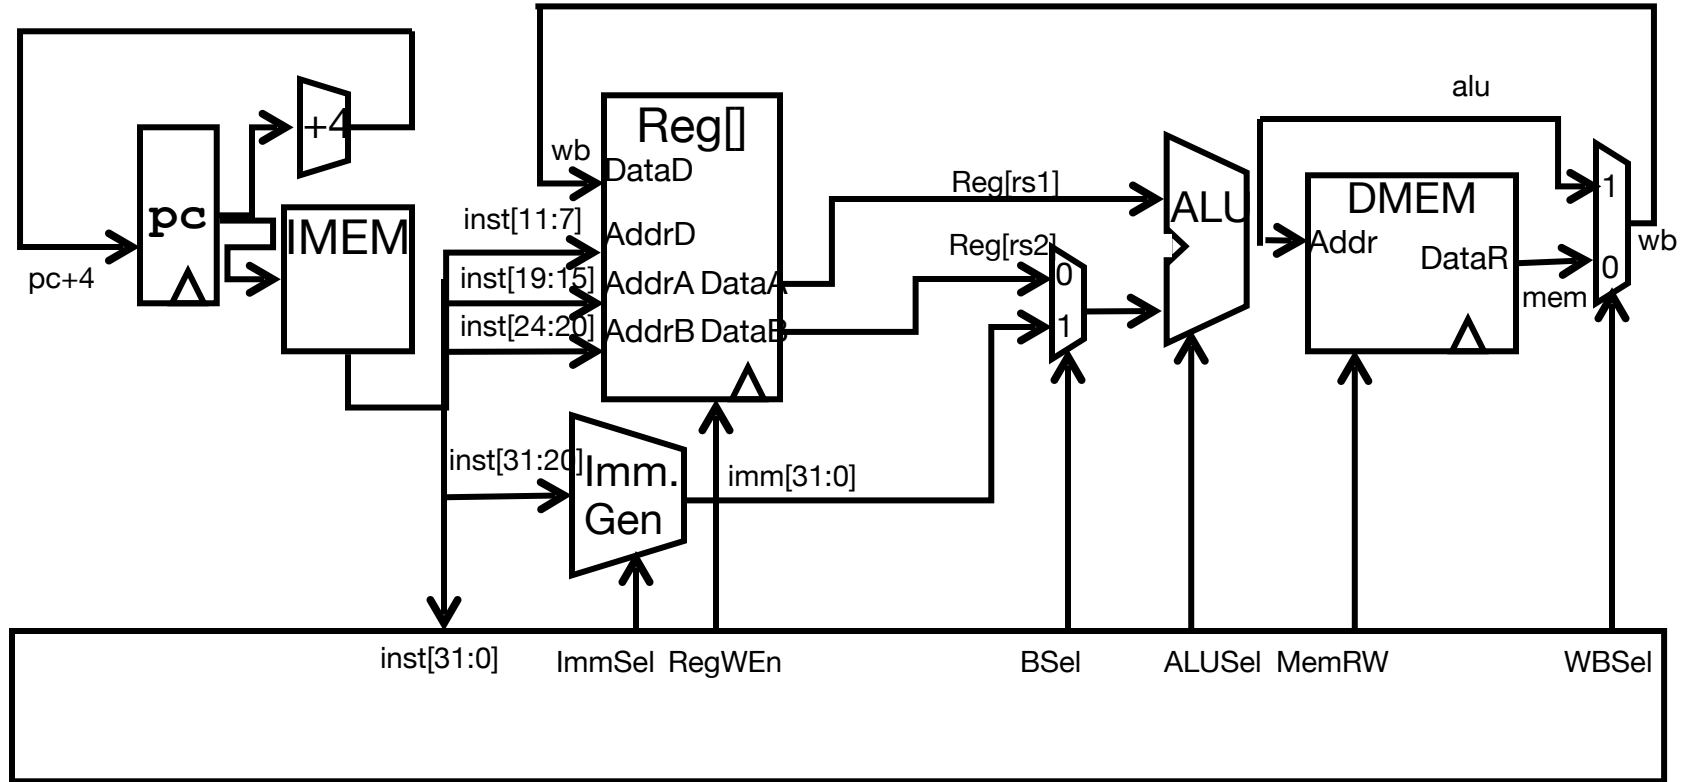
offset[11:5] rs2=14 rs1=2 SW offset[4:0] STORE  
=0 =8



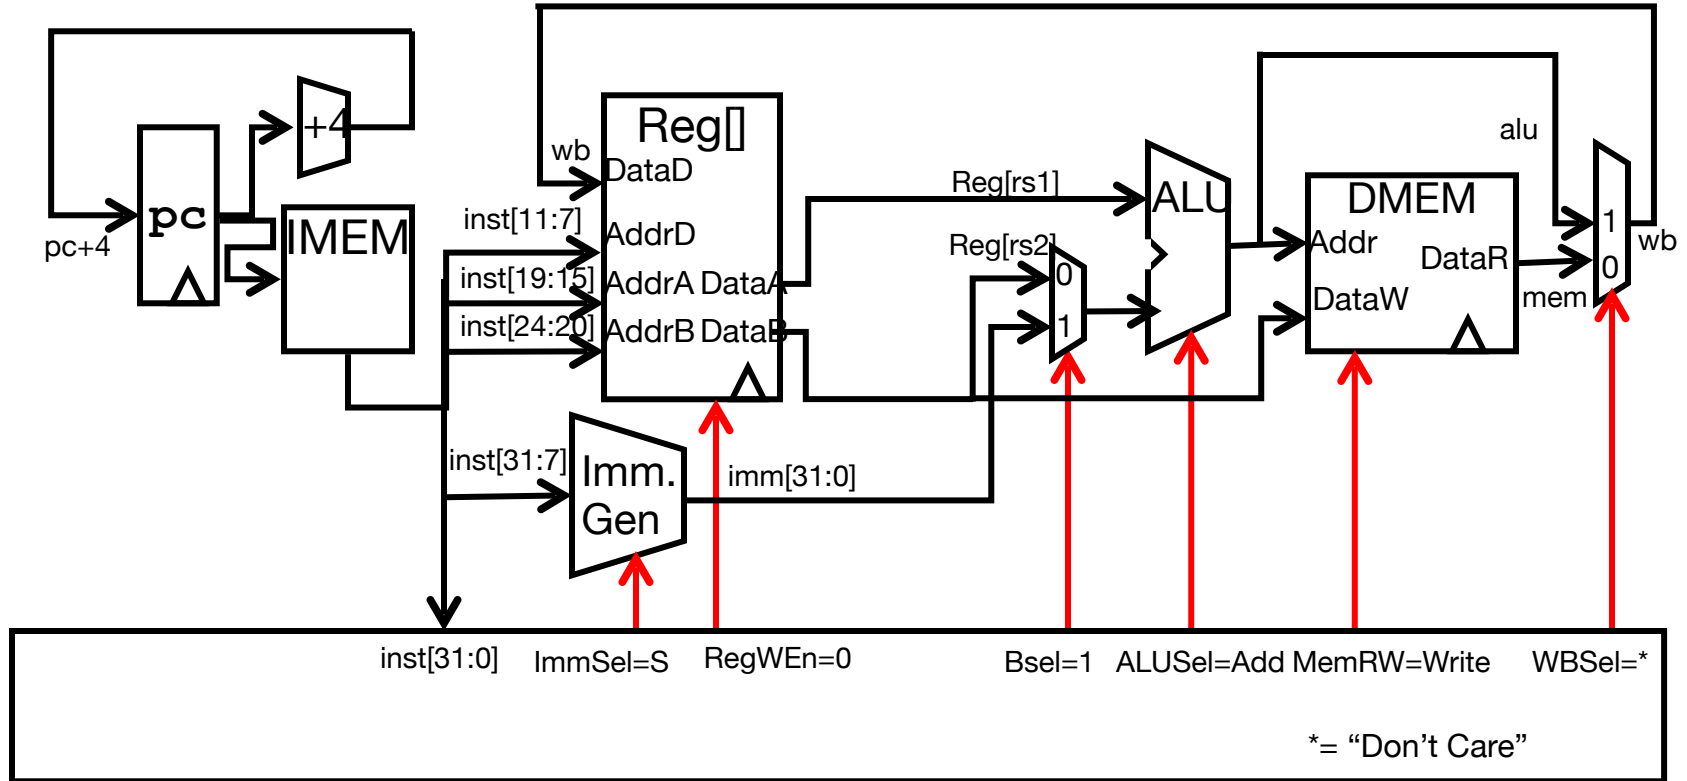
combined 12-bit offset = 8



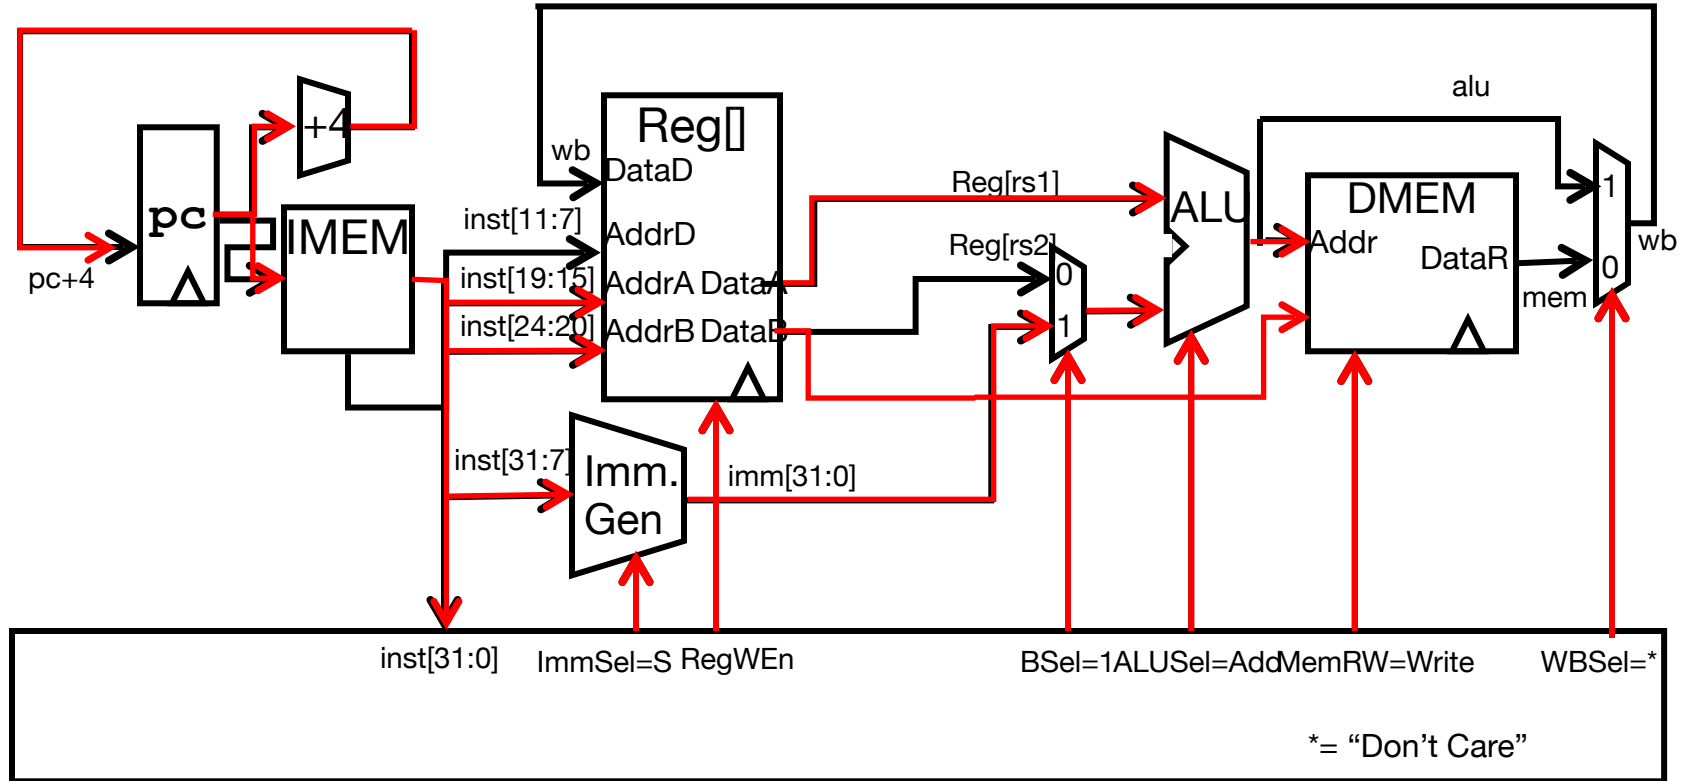
# Adding 1w to datapath




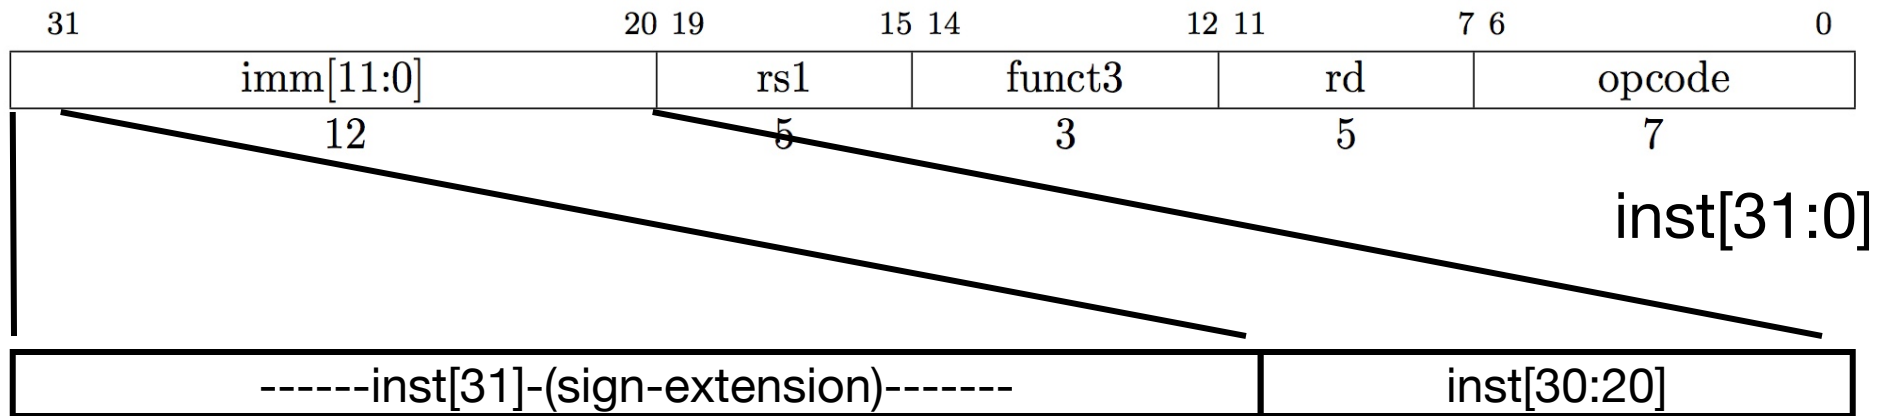
# Adding **sw** to datapath



# Adding **sw** to datapath

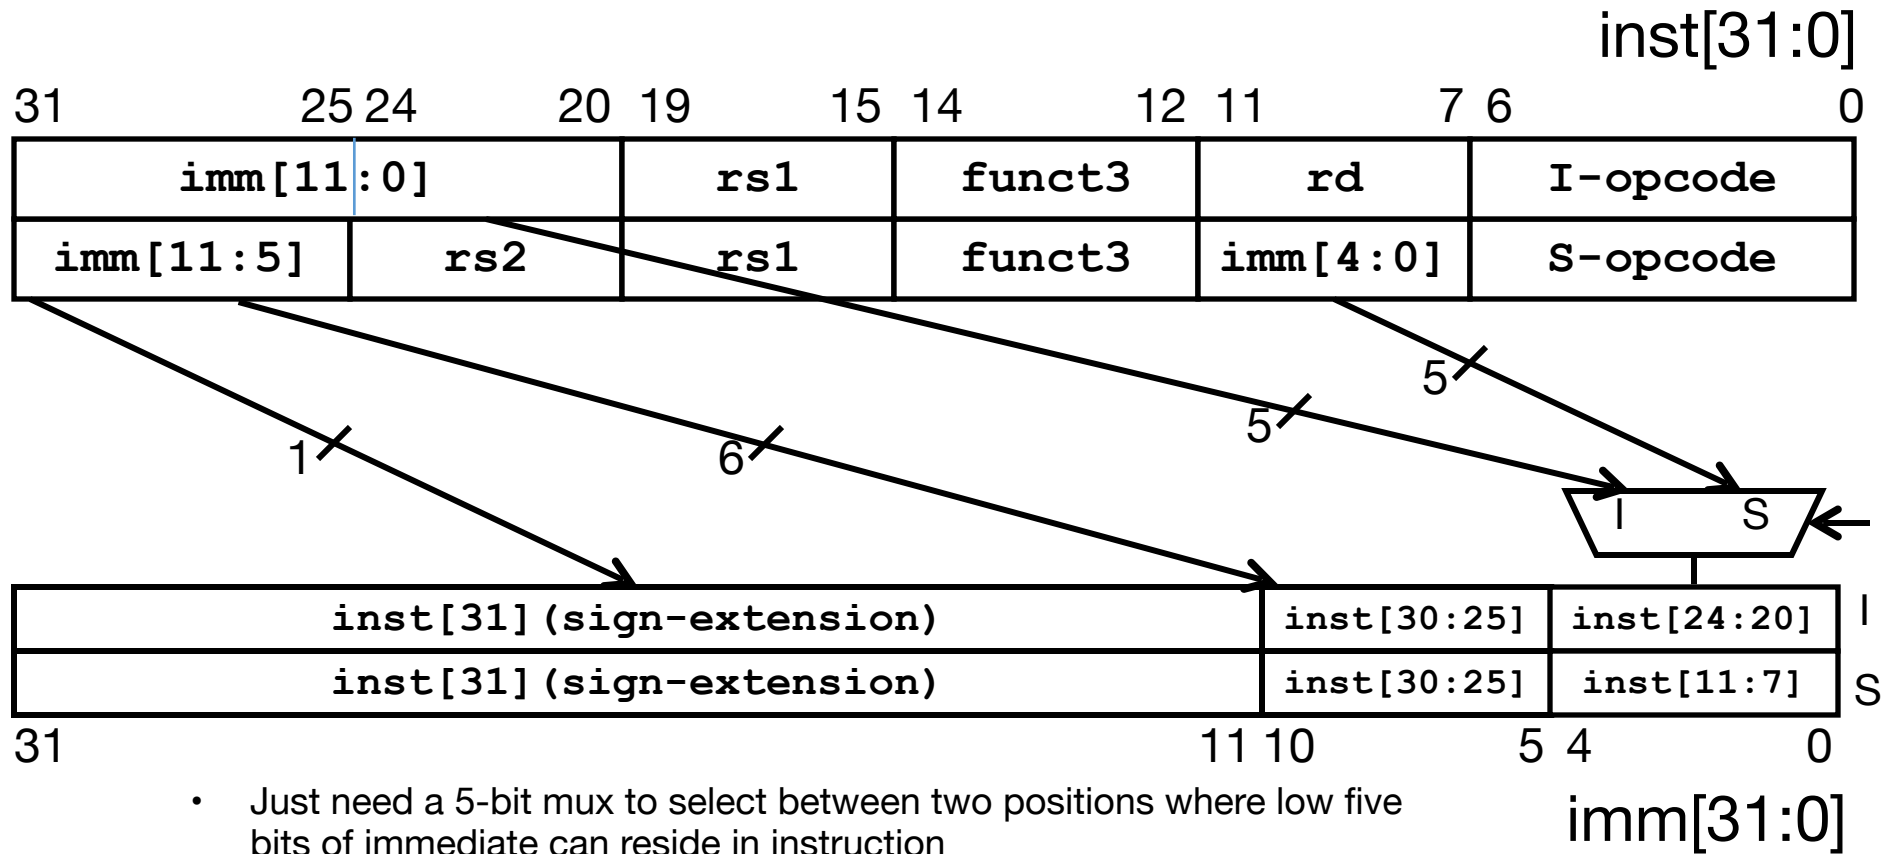


# I-Format immediates




- High 12 bits of instruction (`inst[31:20]`) copied to low 12 bits of immediate (`imm[11:0]`)
- Immediate is sign-extended by copying value of `inst[31]` to fill the upper 20 bits of the immediate value (`imm[31:12]`)

# I & S Immediate Generator



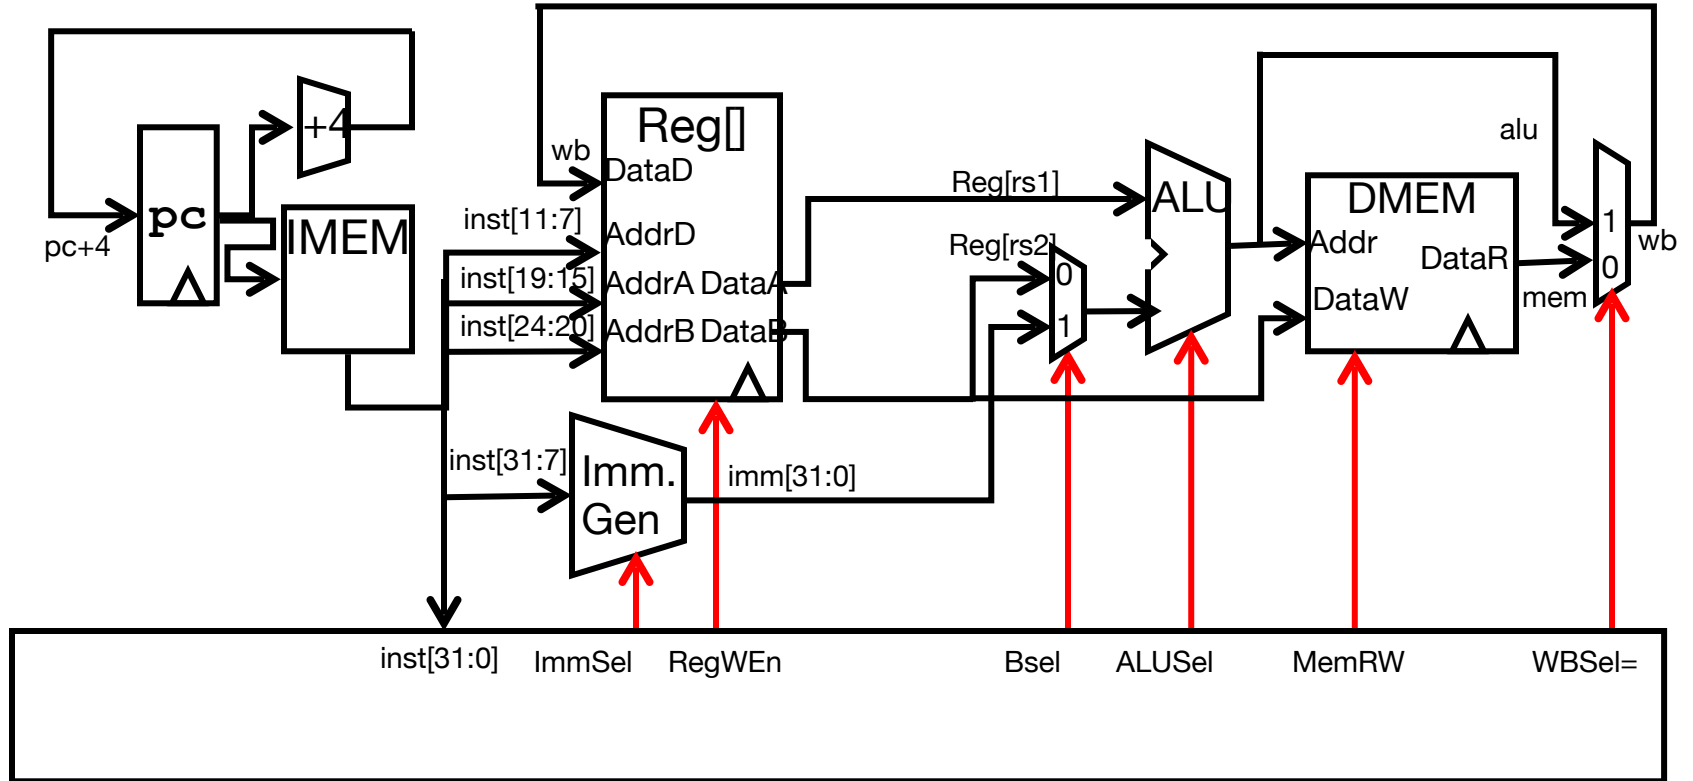
- Just need a 5-bit mux to select between two positions where low five bits of immediate can reside in instruction
- Other bits in immediate are wired to fixed positions in instruction

# Implementing Branches

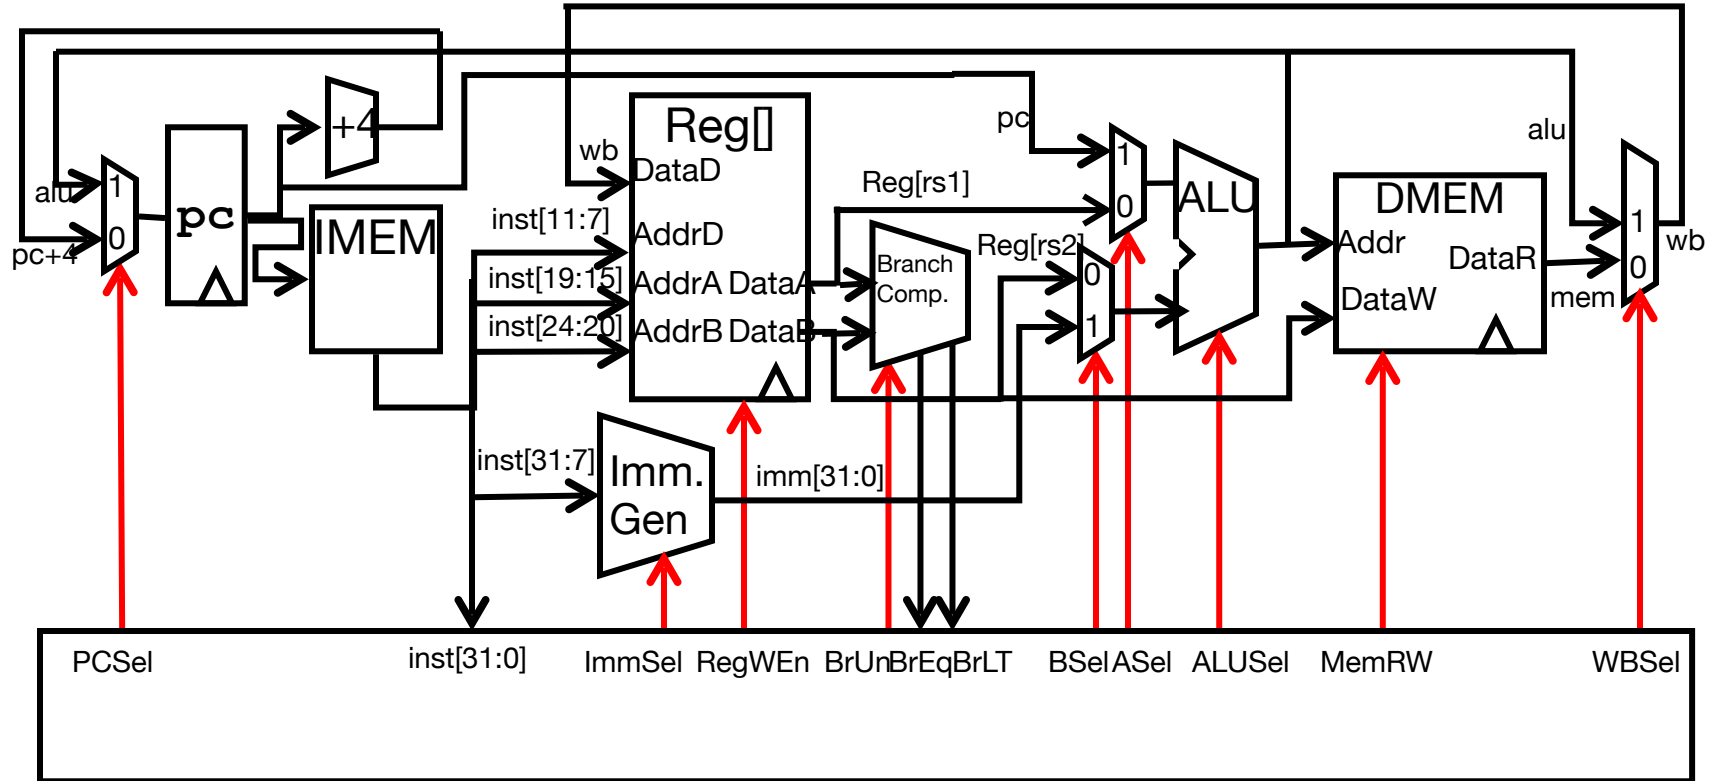


- B-format is mostly same as S-Format, with two register sources (rs1/rs2) and a 12-bit immediate
- But now immediate represents values -4096 to +4094 in 2-byte increments
- The 12 immediate bits encode *even* 13-bit signed byte offsets (lowest bit of offset is always zero, so no need to store it)

# Adding **sw** to datapath

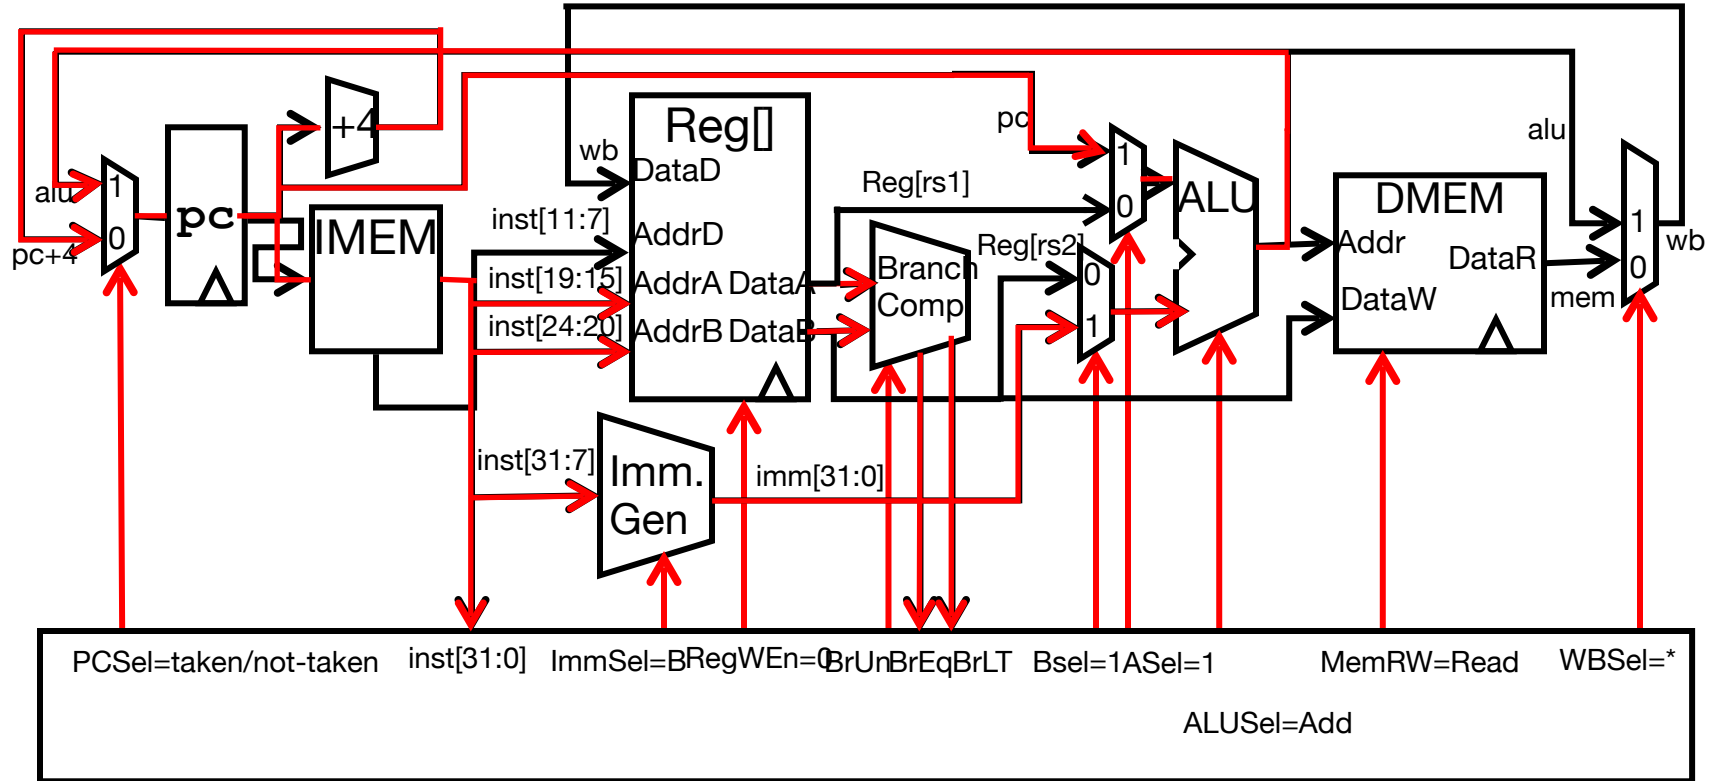


# Adding branches to datapath






# Adding branches to datapath




# Branch Comparator



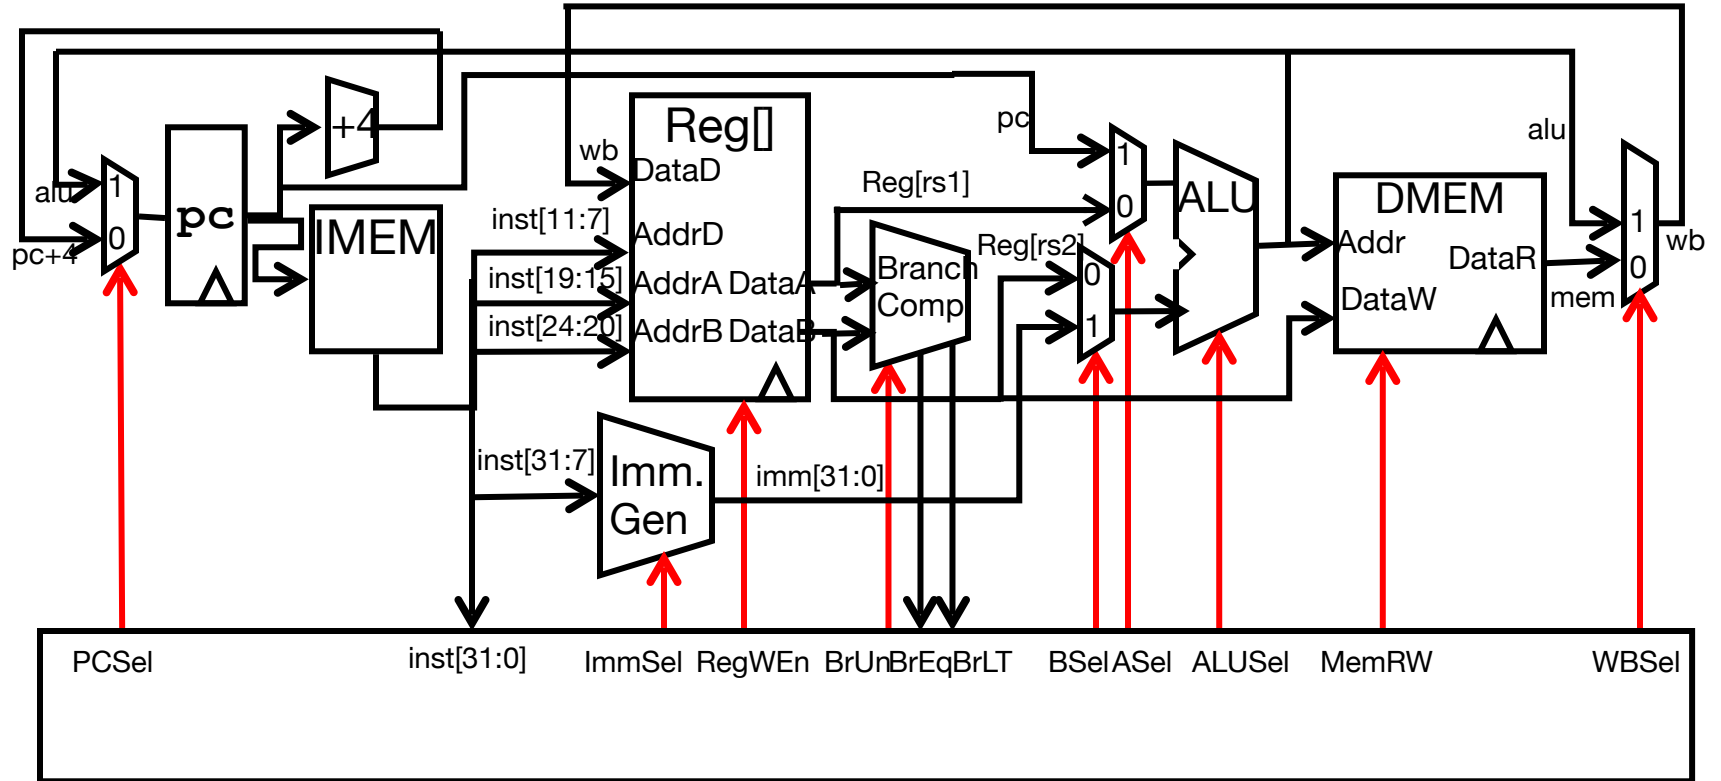
- $\text{BrEq} = 1$ , if  $A=B$
- $\text{BrLT} = 1$ , if  $A < B$
- $\text{BrUn} = 1$  selects unsigned comparison for  $\text{BrLT}$ , 0=signed
- BGE branch:  $A \geq B$ , if  $\neg(A < B)$

# Implementing **JALR** Instruction (I-Format)

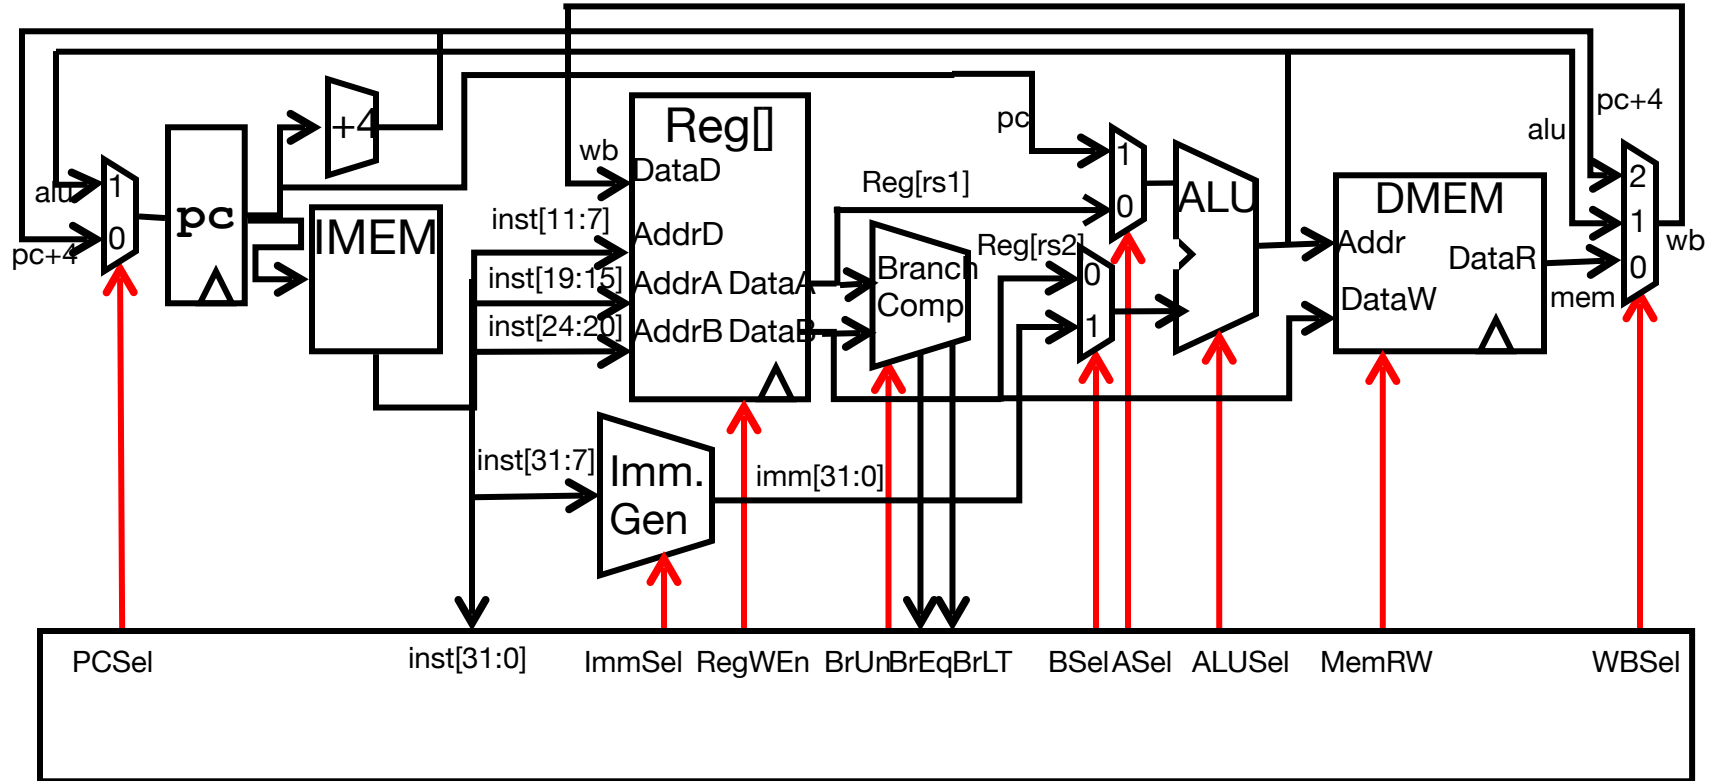


- Sets  $PC = \text{Reg}[\text{rs1}] + \text{immediate}$
- Uses same immediates as arithmetic and loads
  - **no** multiplication by 2 bytes

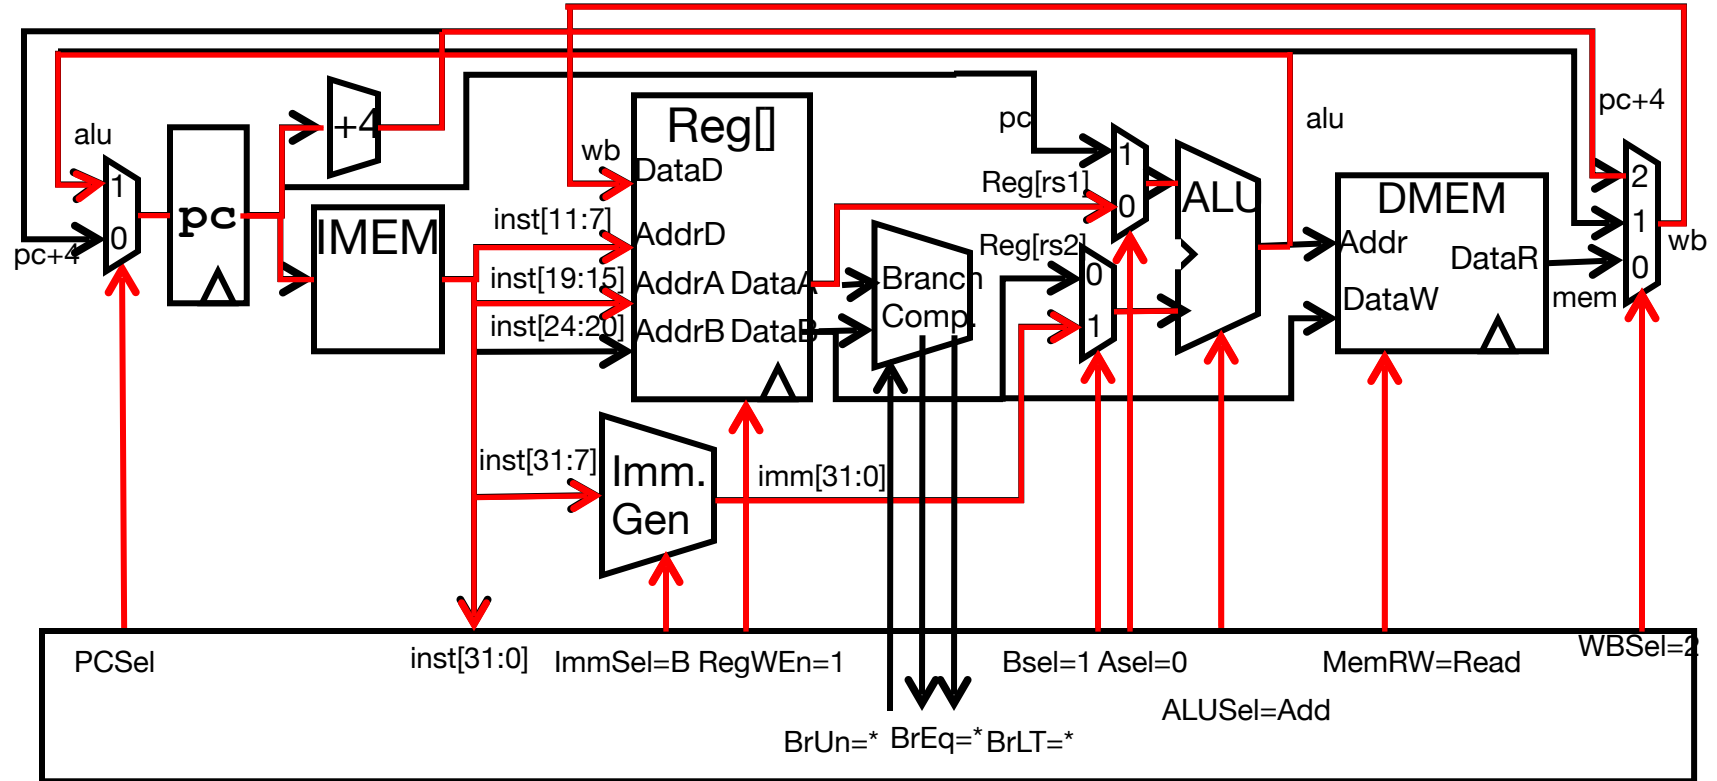
# Adding branches to datapath




# Adding jalr to datapath



# Adding jalr to datapath

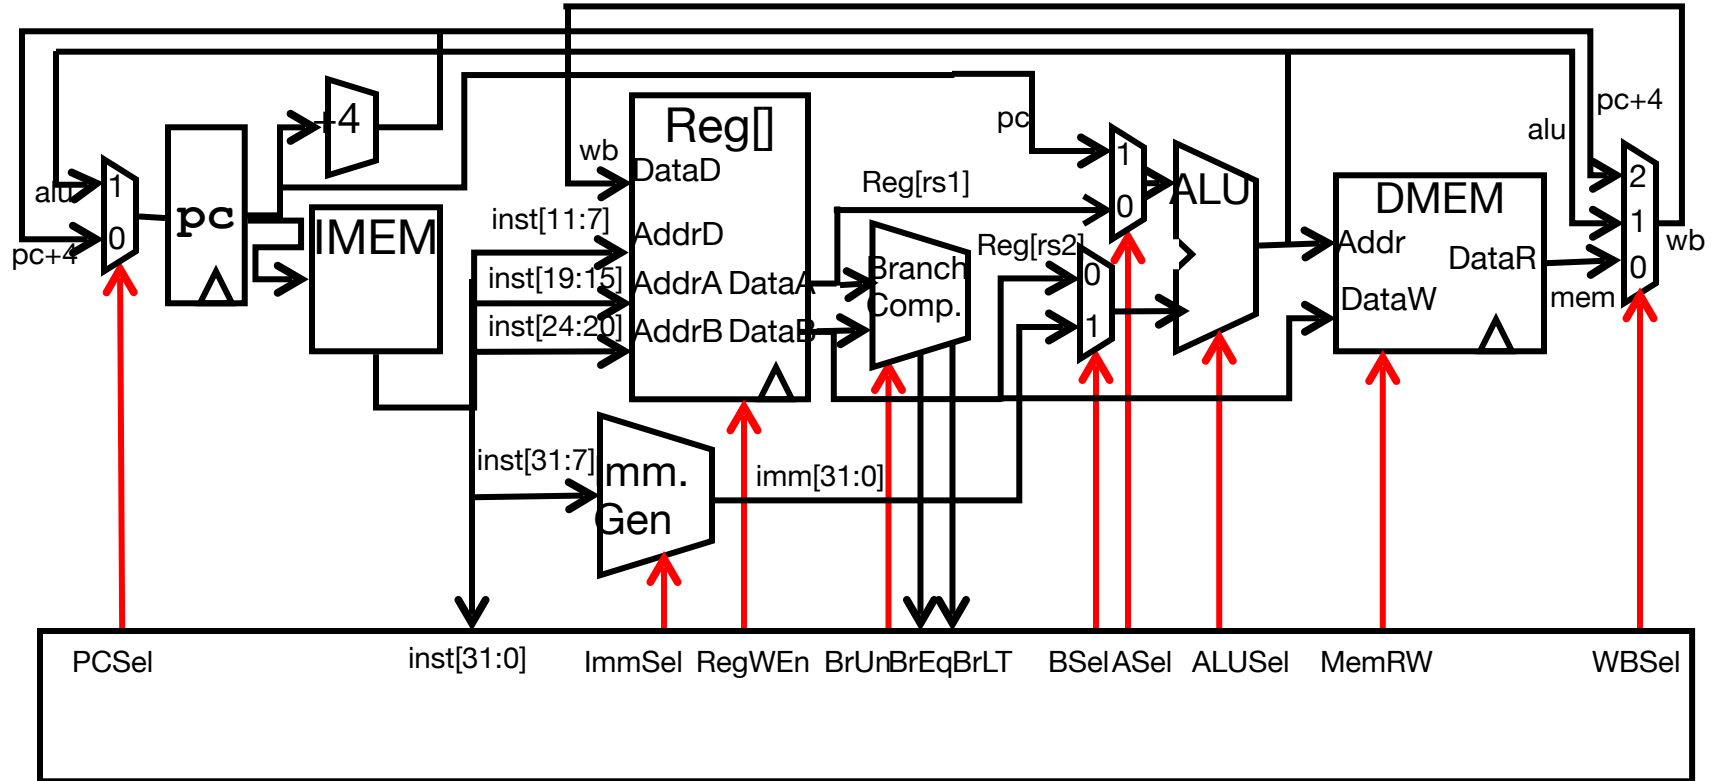


# Implementing jal Instruction



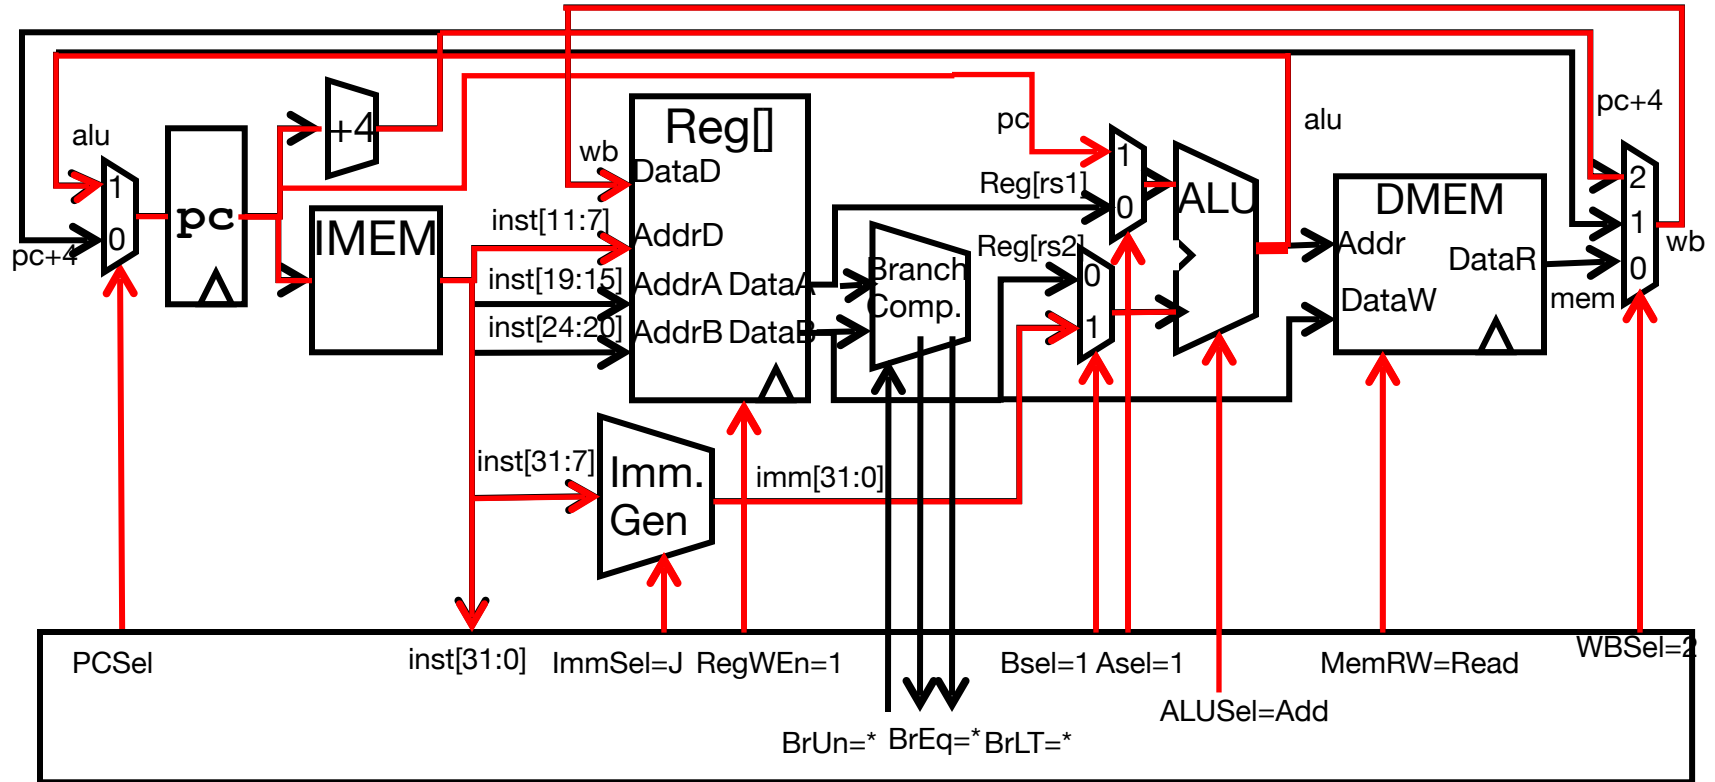
- JAL saves PC+4 in Reg[rd] (the return address)
- Set PC = PC + offset (PC-relative jump)
- Target somewhere within  $\pm 2^{19}$  locations, 2 bytes apart
  - $\pm 2^{18}$  32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost

# Adding jal to datapath

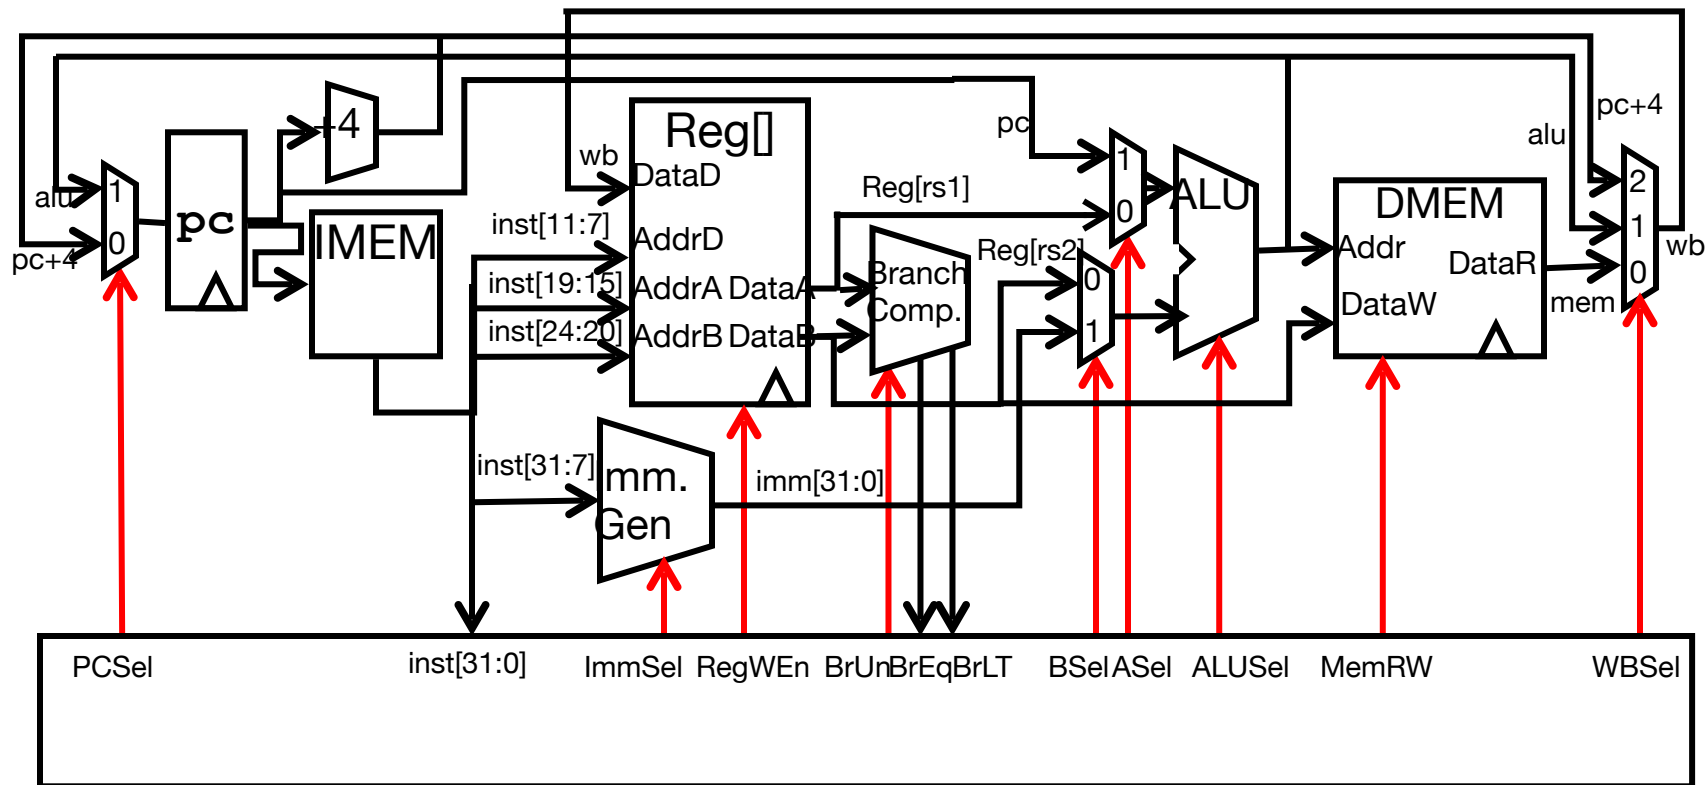




# Adding jal to datapath



# Single-Cycle RISC-V RV32I Datapath



# And in Conclusion, ...

- Universal datapath
  - Capable of executing all RISC-V instructions in one cycle each
  - Not all units (hardware) used by all instructions
- 5 Phases of execution
  - IF, ID, EX, MEM, WB
  - Not all instructions are active in all phases
- Controller specifies how to execute instructions