

**CS 61C:**

**Great Ideas in Computer Architecture**

**Lecture 13: *Control & Operating Speed***

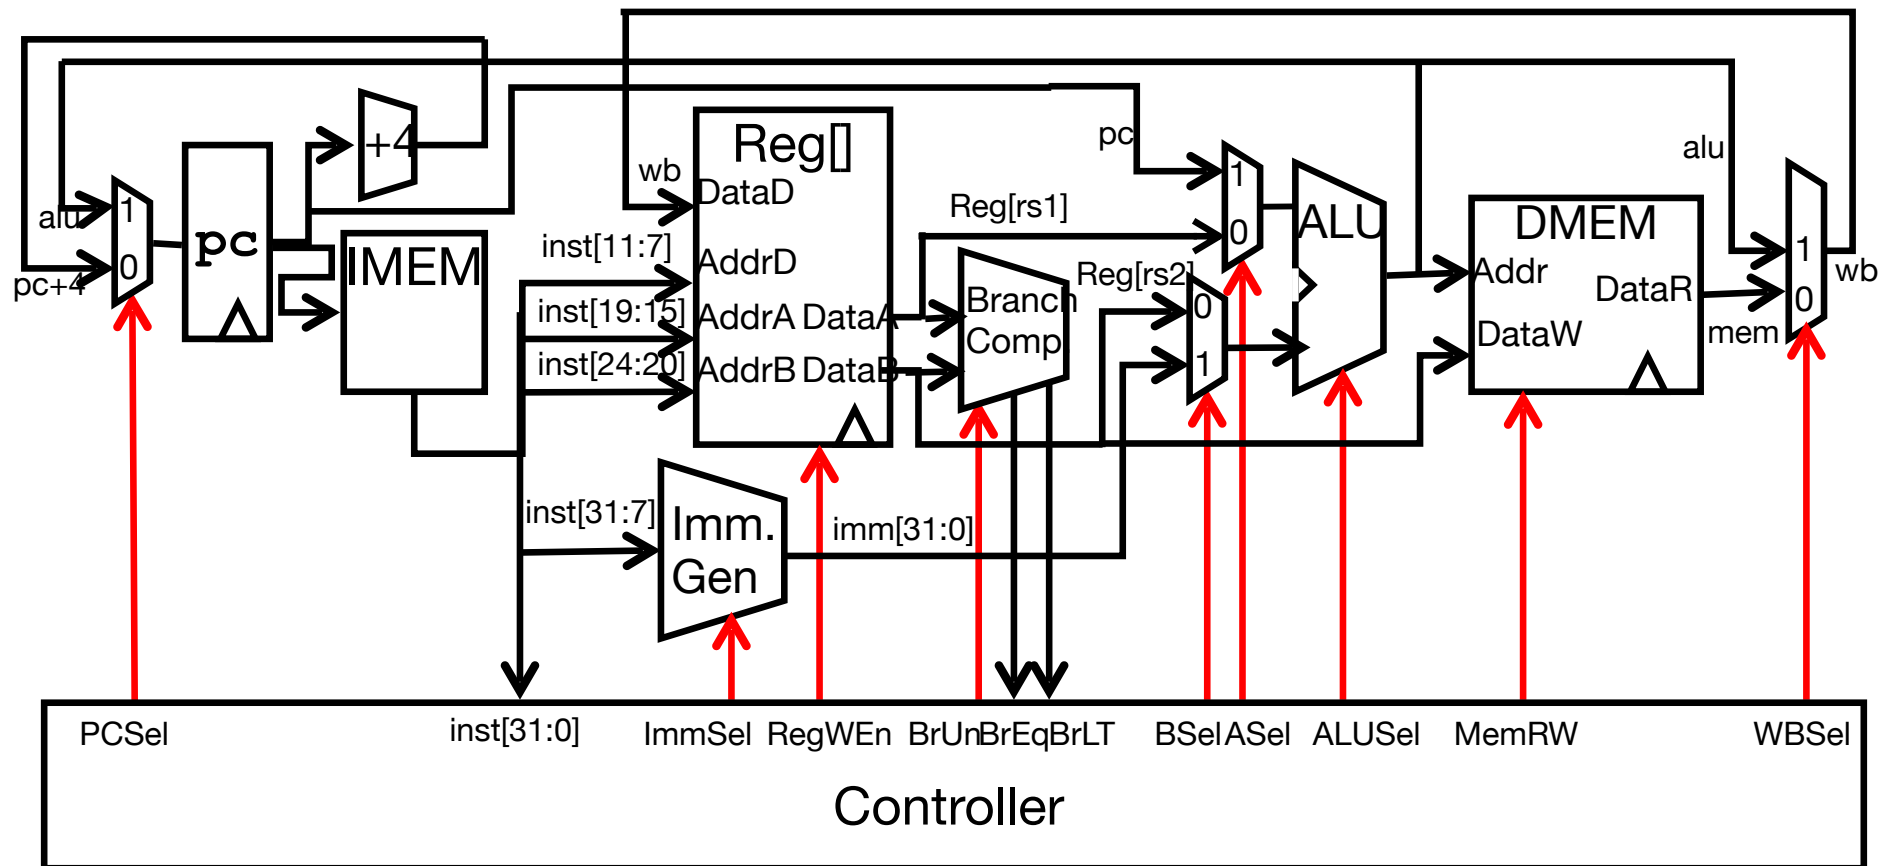
# Administrivia

- Project 2B is due this Thursday (3/4)!
- We will be releasing exam study resources this week for the midterm.
- We will also release information about the exam later this week.
- More information about the Project 1 Clobber will be released shortly.
- Reminder: Office hours have different priorities of ticket types:
  - Normal OH: Conceptual/Homework Questions > Project Questions > Lab Questions
  - Project OH: Project Questions > Conceptual/Homework Questions > Lab Questions
    - Note that different projects do NOT have a difference in priority.
  - Lab OH: Lab Questions > Project/Conceptual/Homework Questions


# Agenda

- **Review Single-Cycle RISC-V Datapath**
- Controller
- Instruction Timing
- Performance Measures
- Introduction to Pipelining
- Pipelined RISC-V Datapath
- And in Conclusion, ...

# Recap: Adding branches to datapath

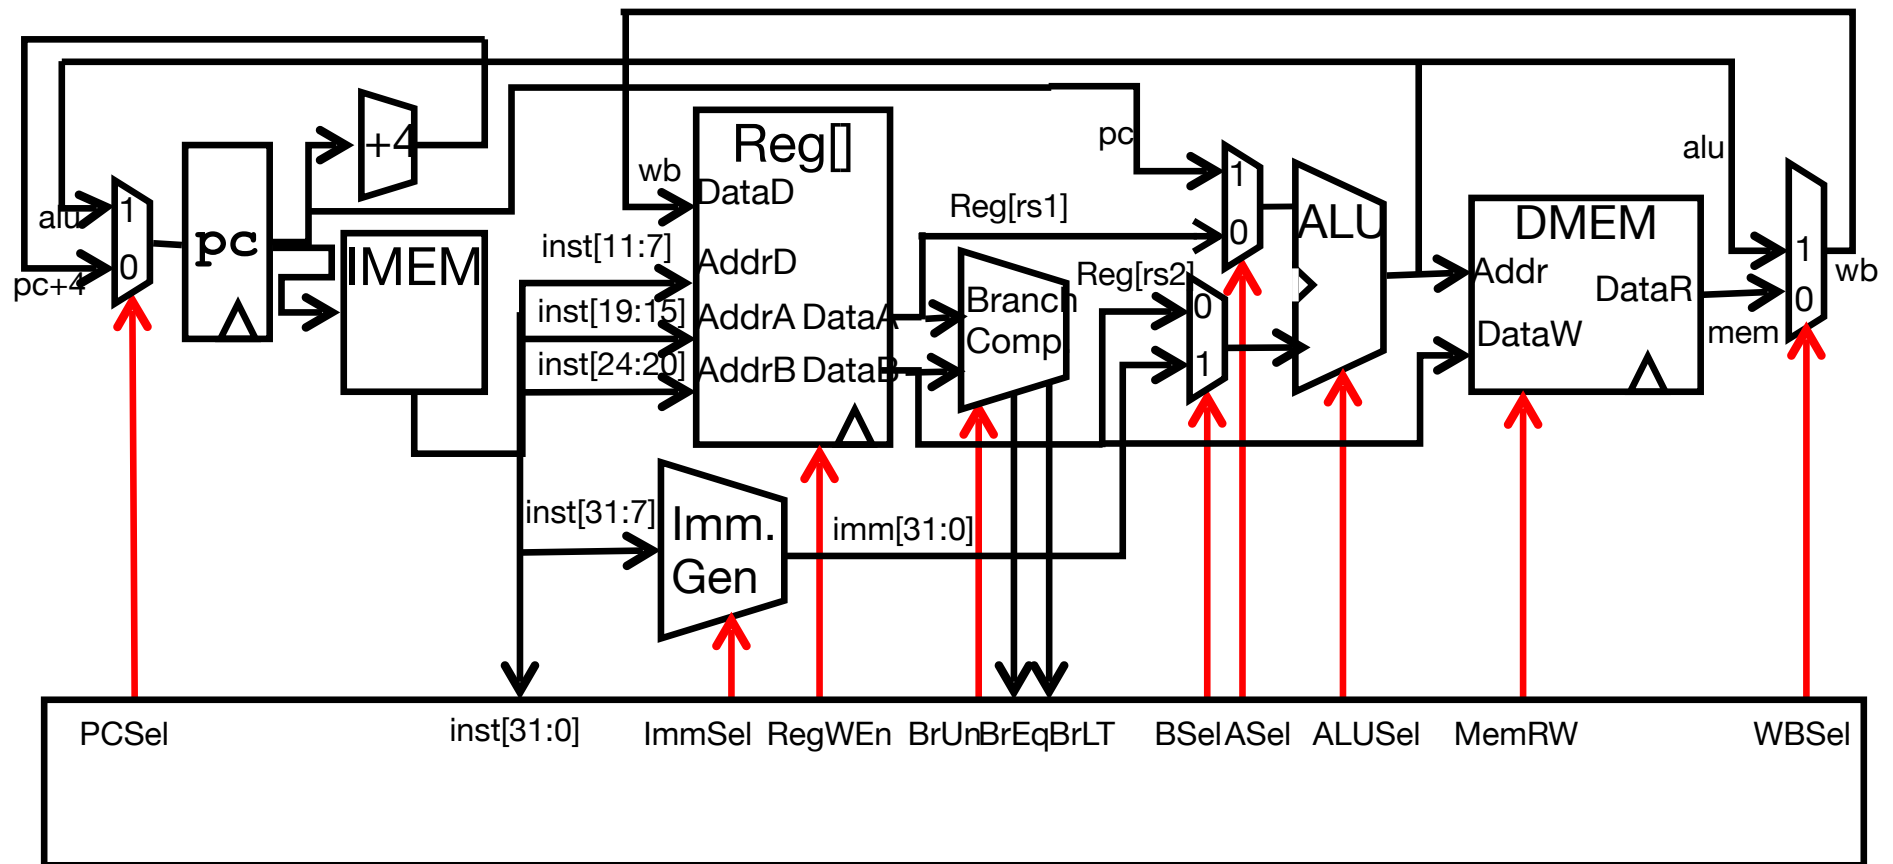


# Implementing **JALR** Instruction (I-Format)

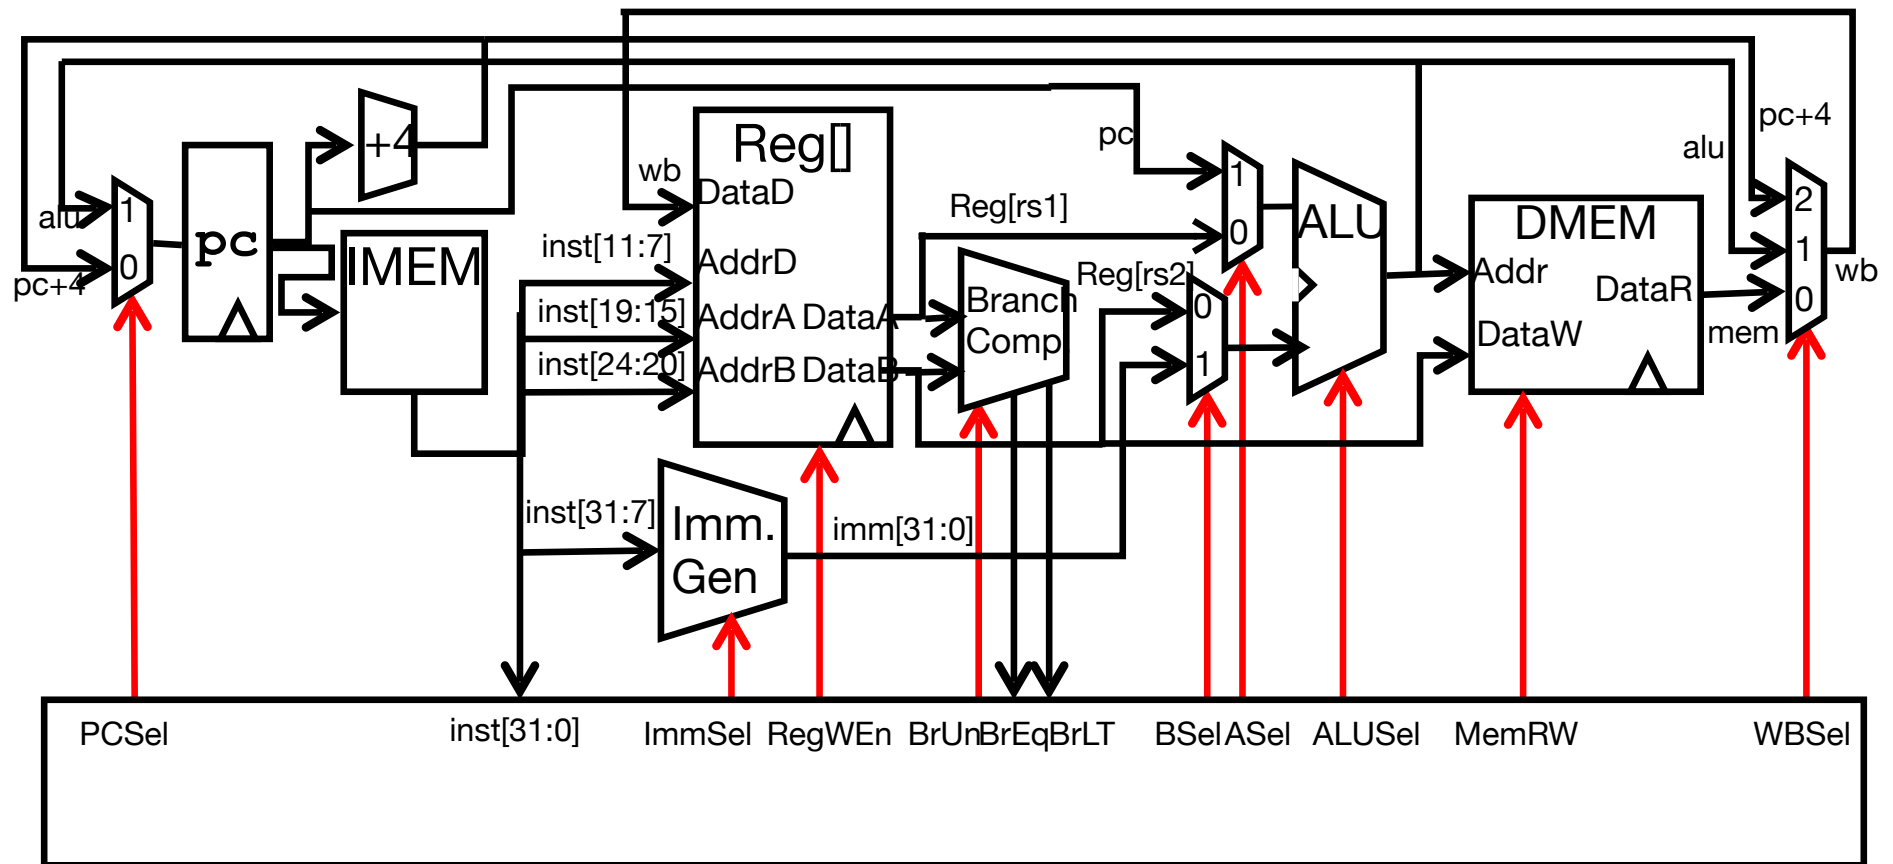


- JALR rd, rs, immediate
  - Writes PC+4 to Reg[rd] (return address)
  - Sets PC = Reg[rs1] + immediate
  - Uses same immediates as arithmetic and loads
    - **no** multiplication by 2 bytes

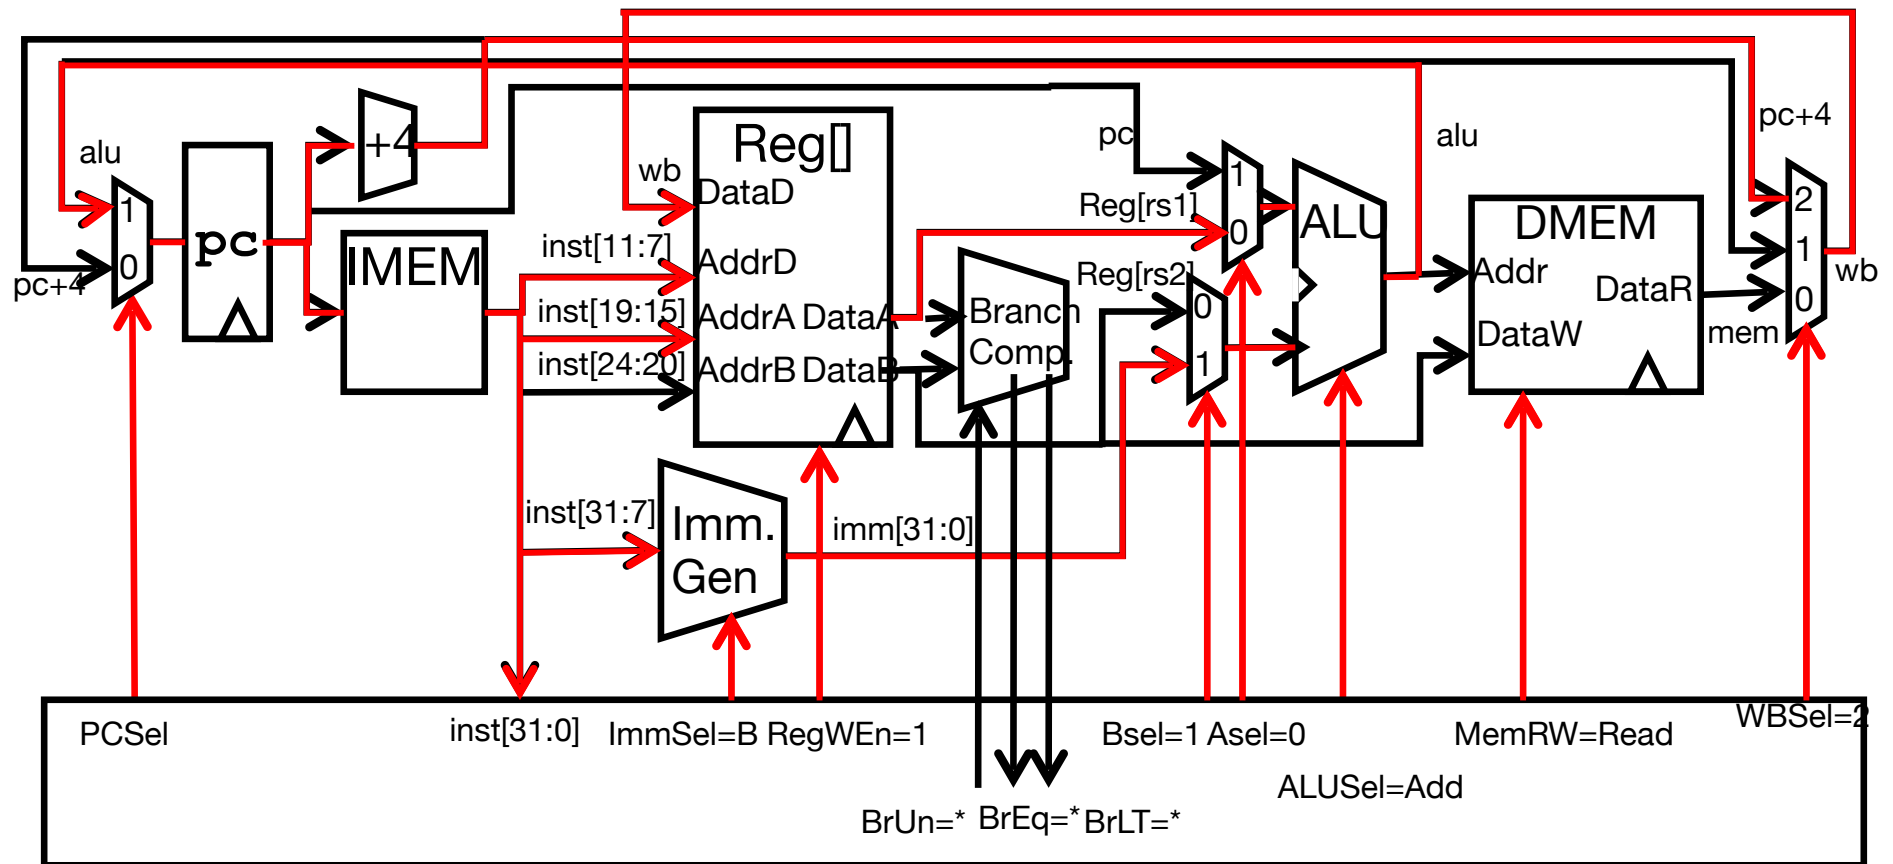
# Datapath with Branches



# Adding `jalr` to datapath




# Adding jalr to datapath



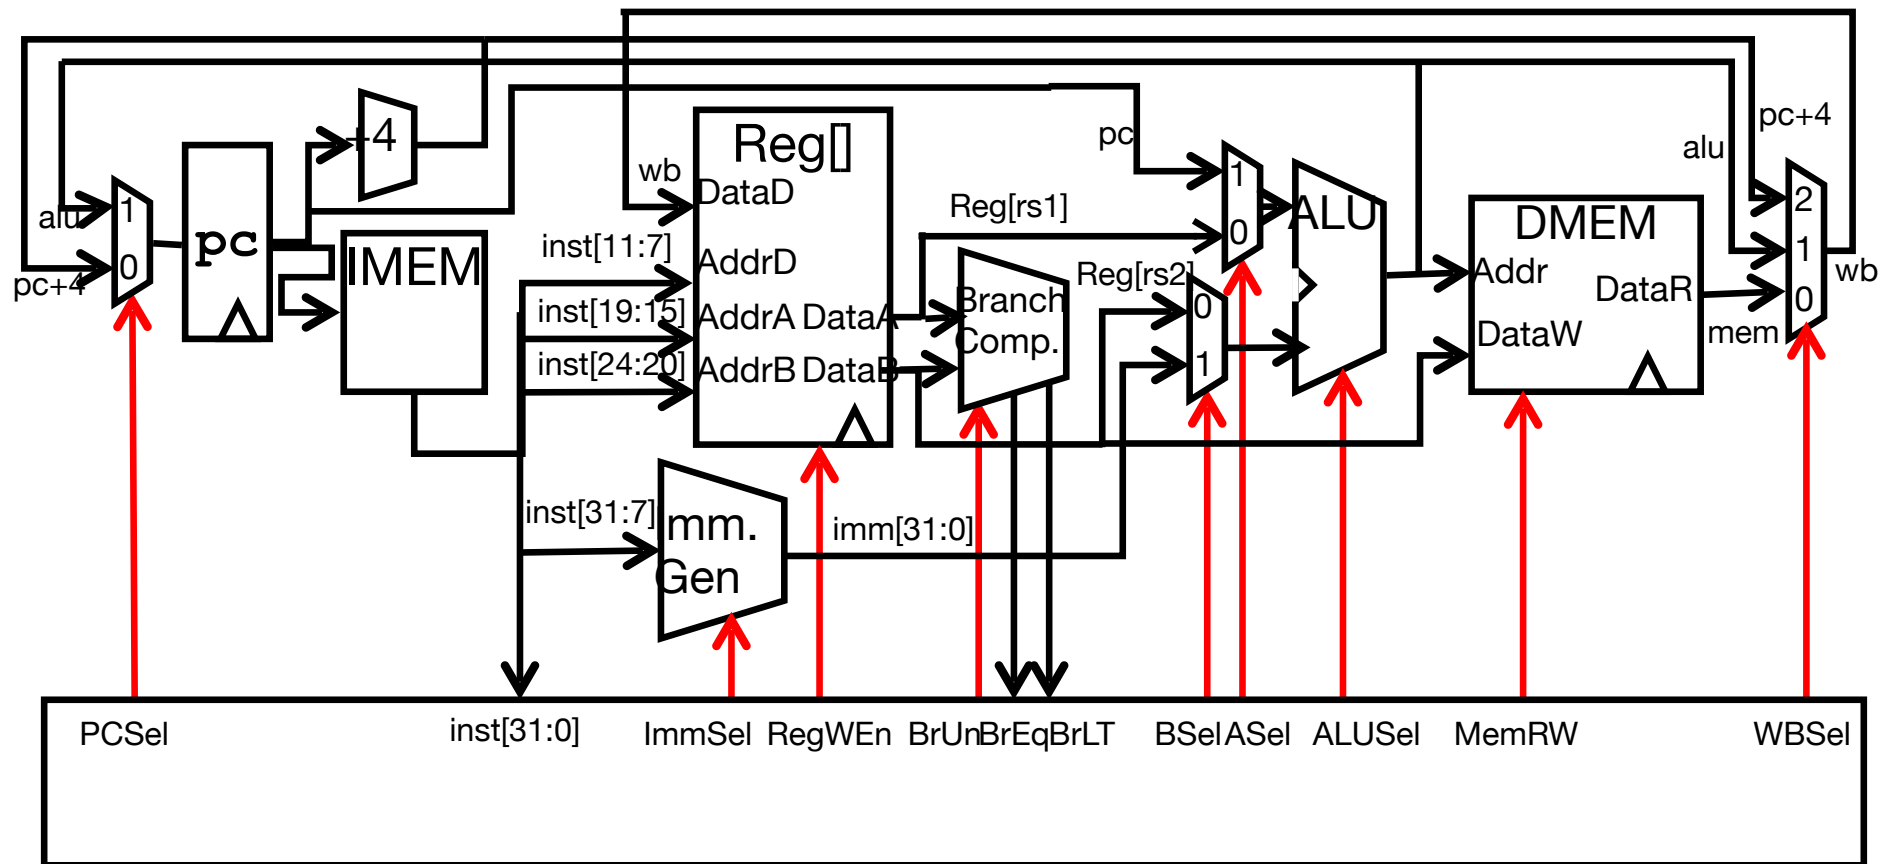


# Implementing `jal` Instruction

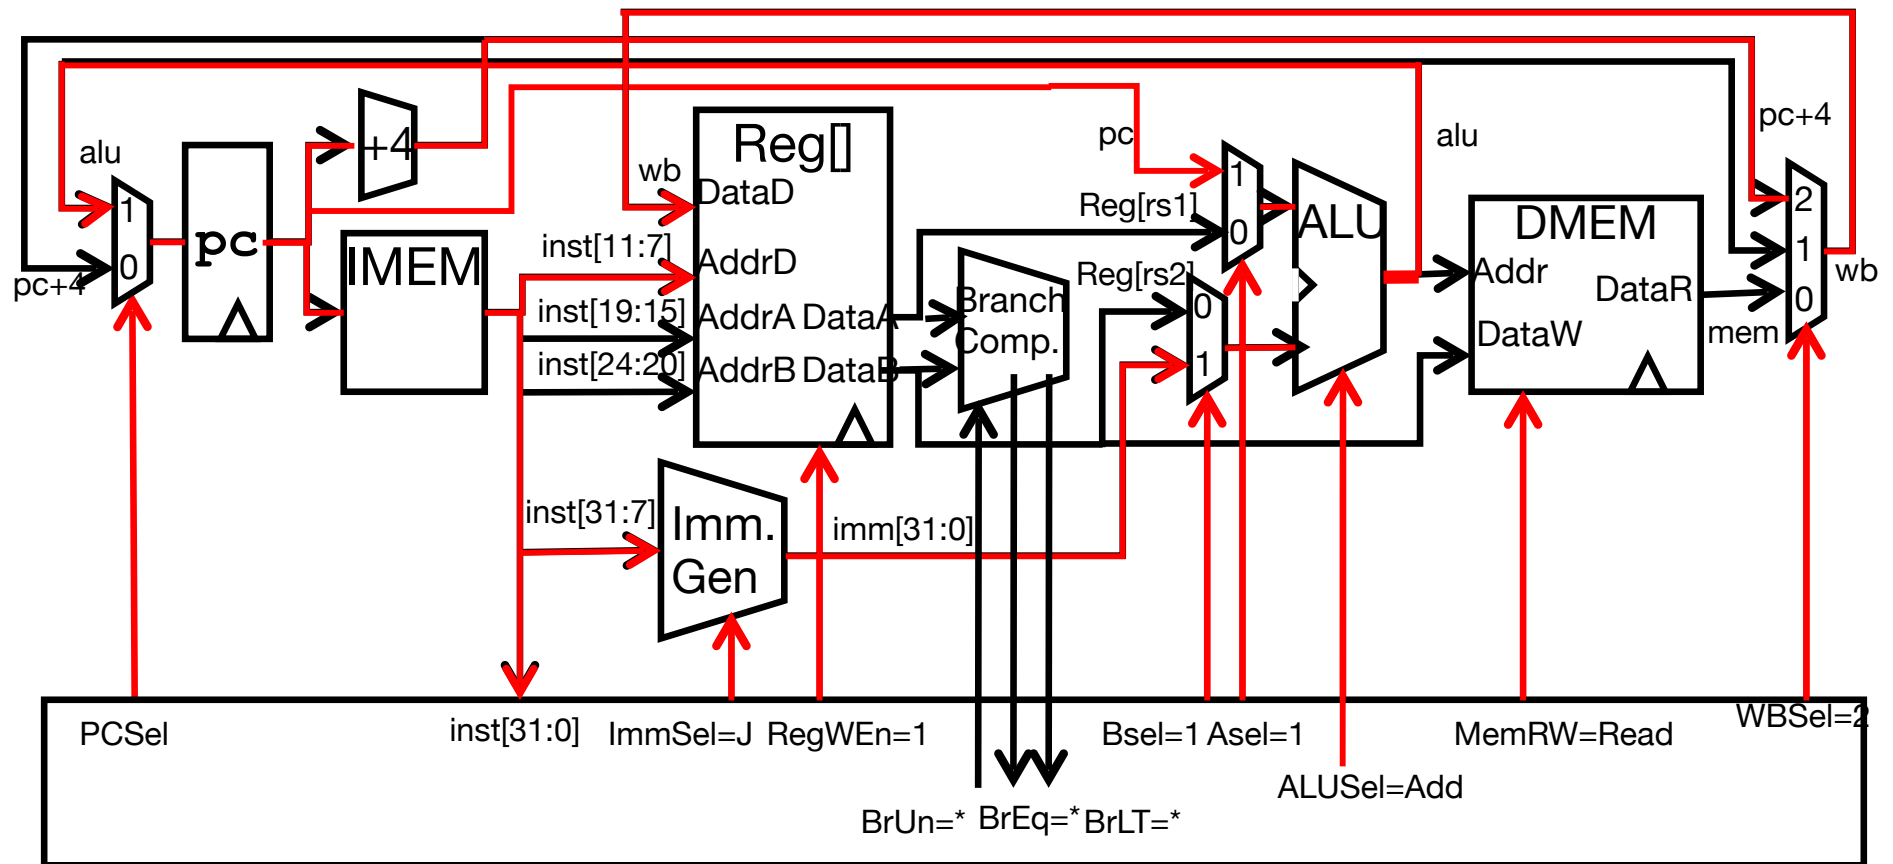


- JAL saves PC+4 in Reg[rd] (the return address)
- Set PC = PC + offset (PC-relative jump)
- Target somewhere within  $\pm 2^{19}$  locations, 2 bytes apart
  - $\pm 2^{18}$  32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost

# Adding jal to datapath



# Adding jal to datapath



# Recap: Complete RV32I ISA

imm[31:12]				rd	0110111
imm[31:12]				rd	0010111
imm[20:10:1 11 19:12]				rd	1101111
imm[11:0]		rs1	000	rd	1100111
imm[12:10:5]		rs2	rs1	000	imm[4:1 11]
imm[12:10:5]		rs2	rs1	001	imm[4:1 11]
imm[12:10:5]		rs2	rs1	100	imm[4:1 11]
imm[12:10:5]		rs2	rs1	101	imm[4:1 11]
imm[12:10:5]		rs2	rs1	110	imm[4:1 11]
imm[12:10:5]		rs2	rs1	111	imm[4:1 11]
imm[11:0]		rs1	000	rd	0000011
imm[11:0]		rs1	001	rd	0000011
imm[11:0]		rs1	010	rd	0000011
imm[11:0]		rs1	100	rd	0000011
imm[11:0]		rs1	101	rd	0000011
imm[11:5]		rs2	rs1	000	imm[4:0]
imm[11:5]		rs2	rs1	001	imm[4:0]
imm[11:5]		rs2	rs1	010	imm[4:0]
imm[11:0]		rs1	000	rd	0010011
imm[11:0]		rs1	010	rd	0010011
imm[11:0]		rs1	011	rd	0010011
imm[11:0]		rs1	100	rd	0010011
imm[11:0]		rs1	110	rd	0010011
imm[11:0]		rs1	111	rd	0010011

LUI  
 AUIPC  
 JAL  
 JALR  
 BEQ  
 BNE  
 BLT  
 BGE  
 BLTU  
 BGEU  
 LB  
 LH  
 LW  
 LBU  
 LHU  
 SB  
 SH  
 SW  
 ADDI  
 SLTI

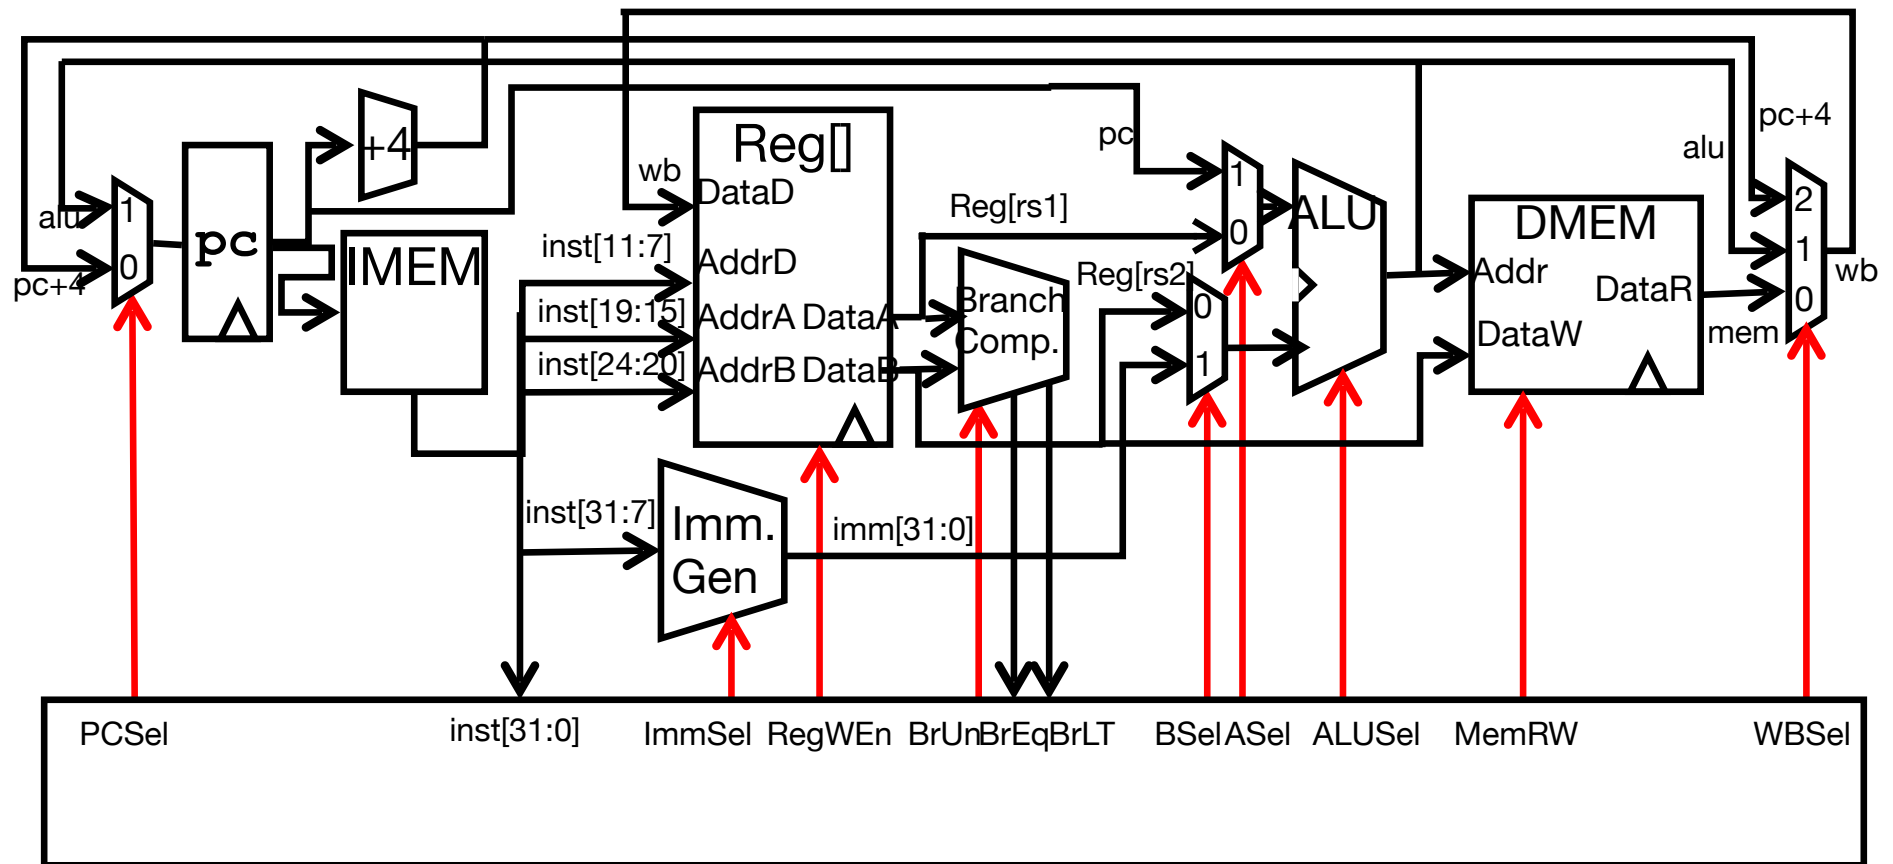
0000000	shamt	rs1	001	rd	0010011	SLLI	
0000000	shamt	rs1	101	rd	0010011	SRLI	
0100000	shamt	rs1	101	rd	0010011	SRAI	
0000000	rs2	rs1	000	rd	0110011	ADD	
0100000	rs2	rs1	000	rd	0110011	SUB	
0000000	rs2	rs1	001	rd	0110011	SLL	
0000000	rs2	rs1	010	rd	0110011	SLT	
0000000	rs2	rs1	011	rd	0110011	SLTU	
0000000	rs2	rs1	100	rd	0110011	XOR	
0000000	rs2	rs1	101	rd	0110011	SRL	
0100000	rs2	rs1	101	rd	0110011	SRA	
0000000	rs2	rs1	110	rd	0110011	OR	
0000000	rs2	rs1	111	rd	0110011	AND	
0000	pred	succ	00000	000	00000	0001111	FENCE
0000	0000	0000	00000	001	00000	0001111	FENCE.I
0000000000000			00000	000	00000	1110011	ECALL
0000000000001			00000	000	00000	1110011	EBREAK
csr		rs1	001	rd	1110011	CSR.W	
csr		rs1	011	rd	1110011	CSR.RS	

Not in CS61C

RV32I has 47 instructions total  
 37 instructions covered in CS61C

Remaining instructions (ex: lui, auipc) can be implemented with no significant additions to the datapath: adding a “pass B” option to the ALU and another immediate decoding option. Rest is all control logic

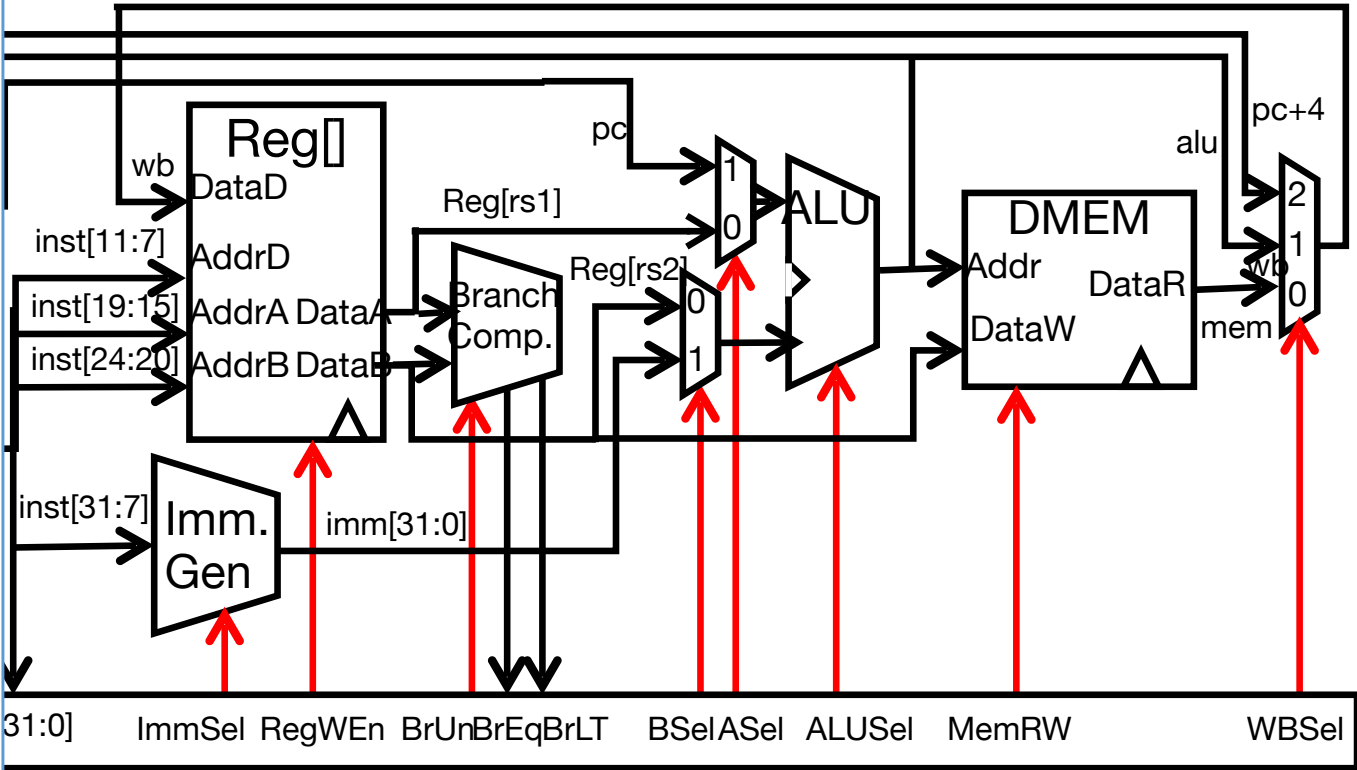
# Single-Cycle RISC-V RV32I Datapath



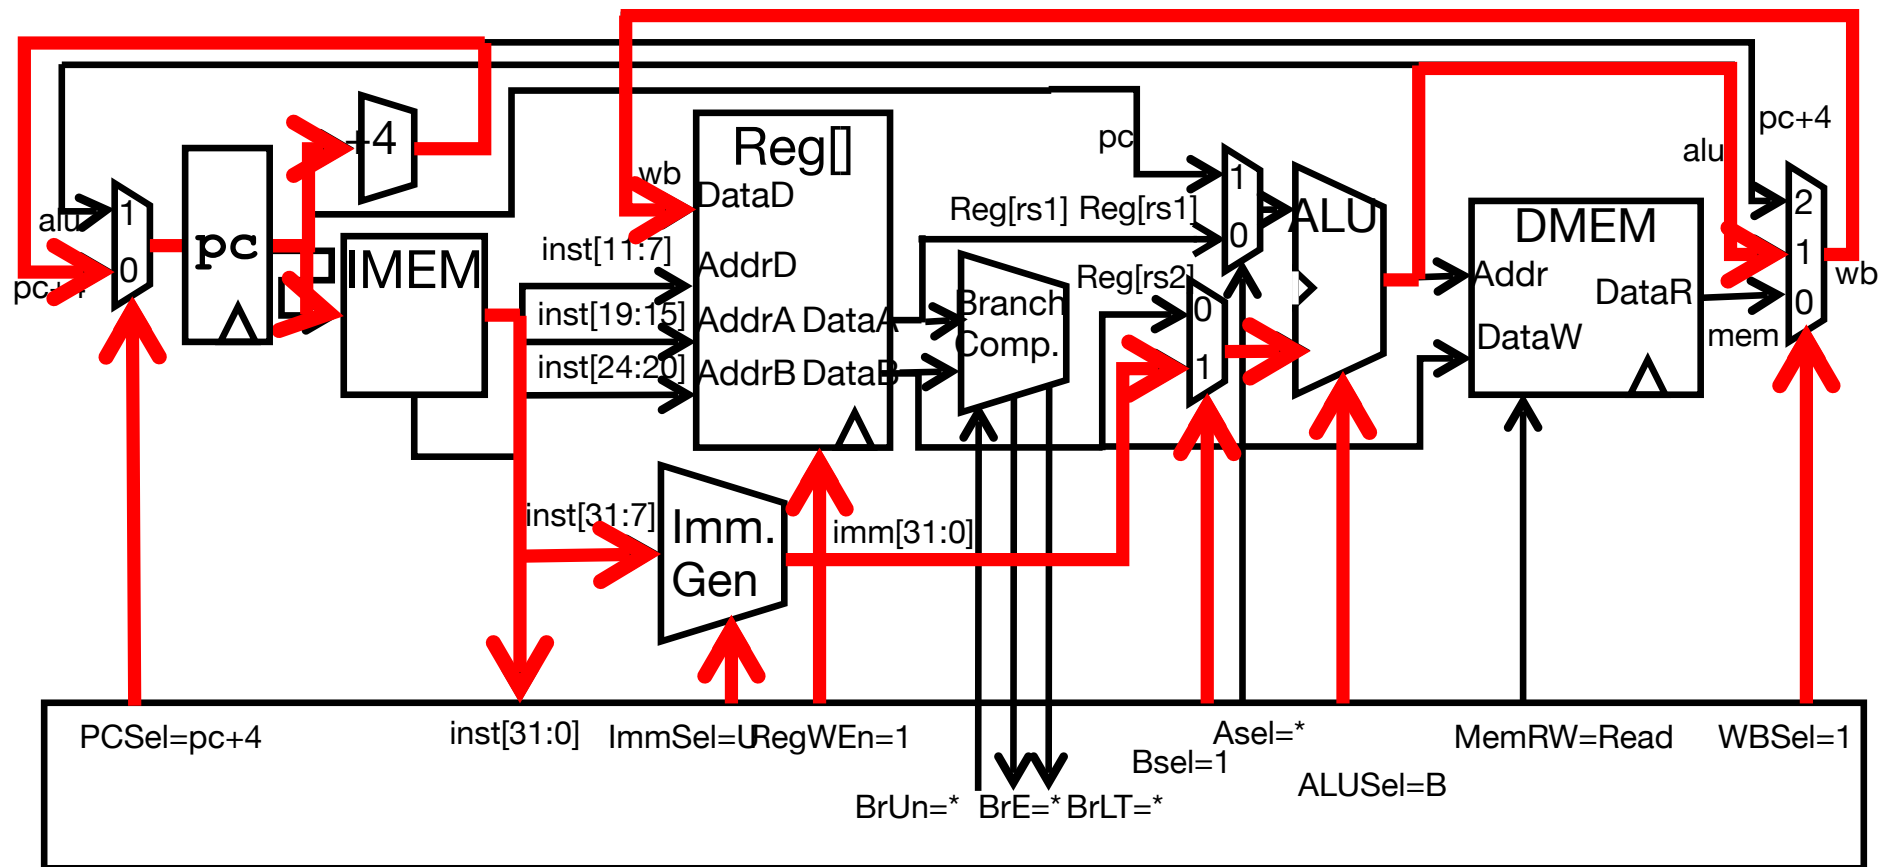
# Clicker Question

What are proper control signals for **lui** instruction?

- A:** BSel=0, ASel=0, WBSel=0
- B:** BSel=0, ASel=0, WBSel=1
- C:** BSel=0, ASel=1, WBSel=1
- D:** BSel=1, ASel=1, WBSel=1



# Implementing lui

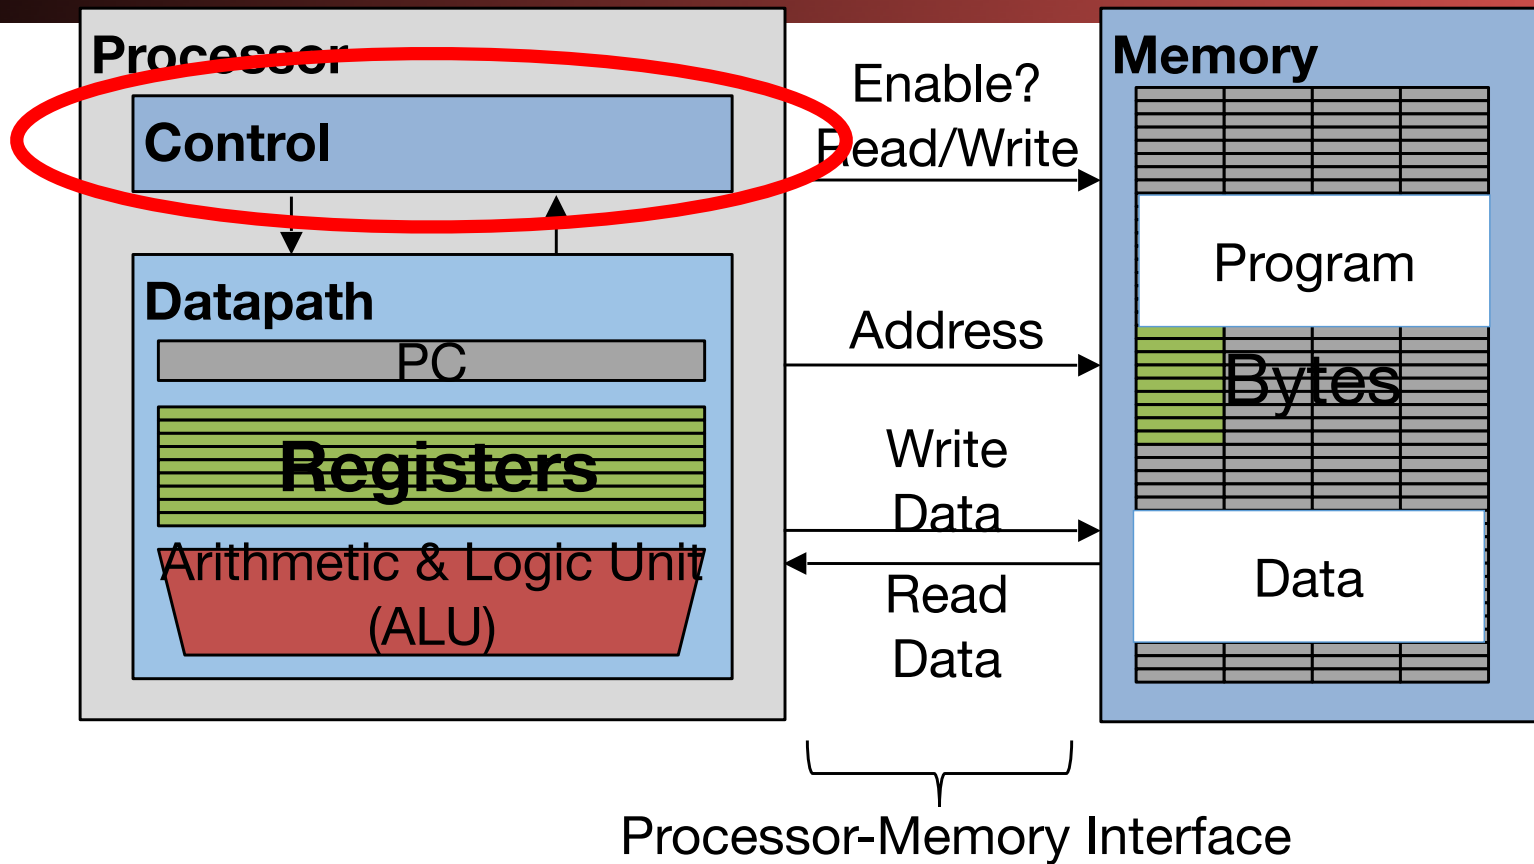


# Agenda

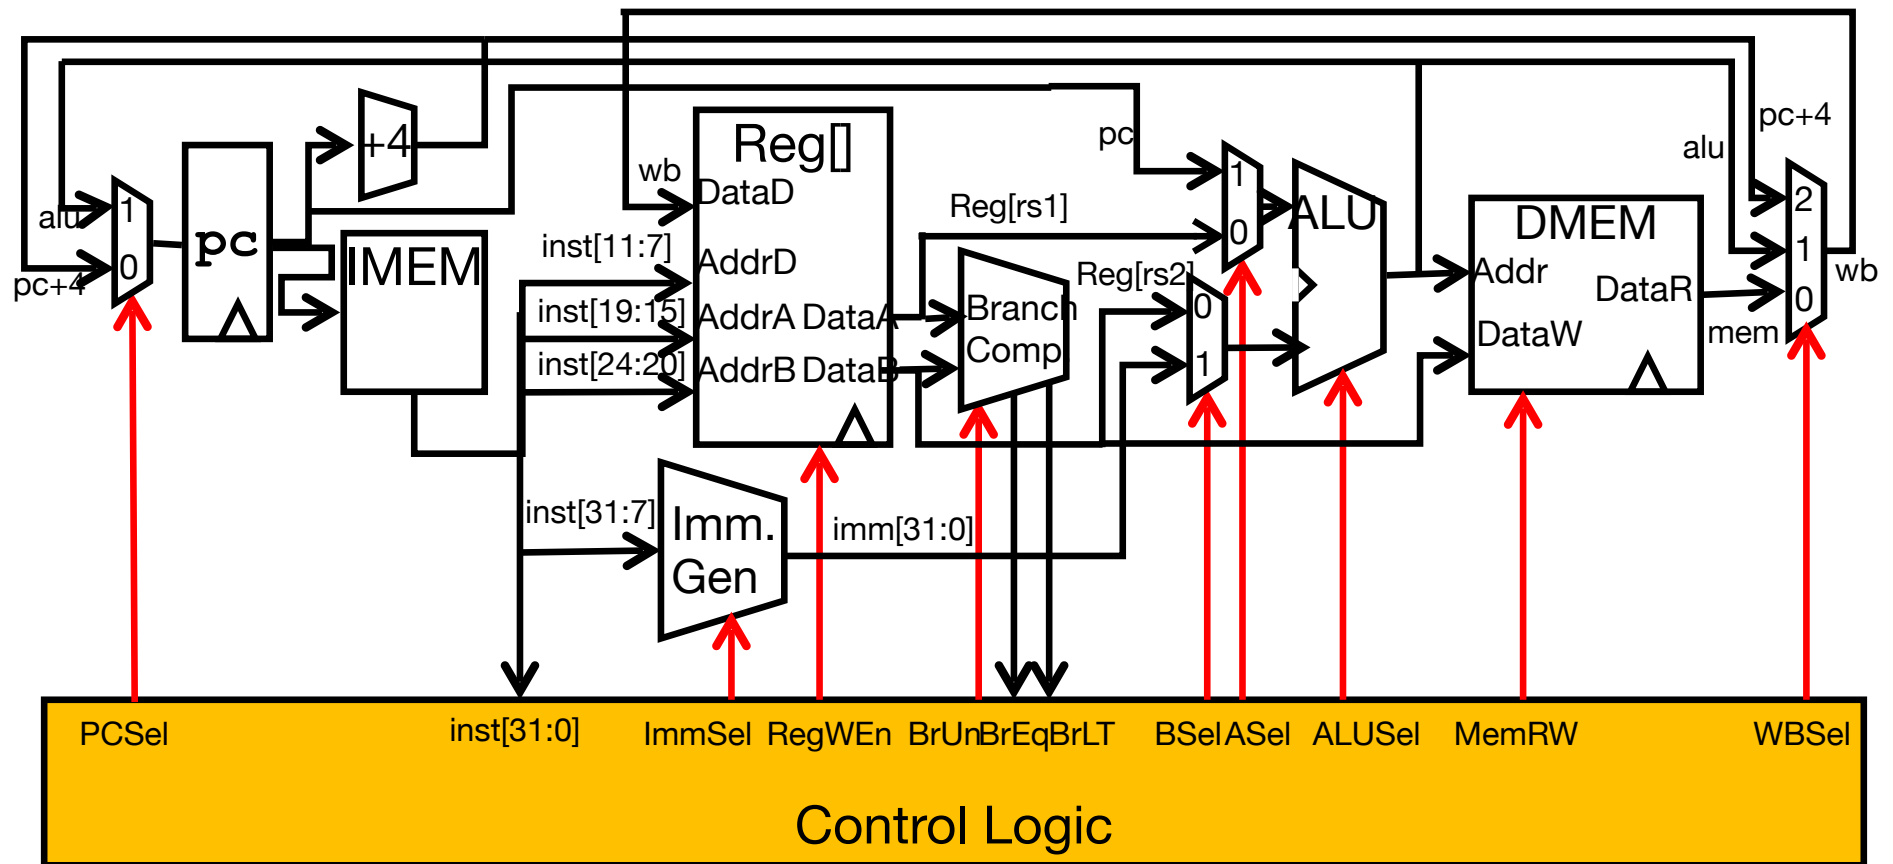
- Finish Single-Cycle RISC-V Datapath
- **Controller**
- Instruction Timing
- Performance Measures
- Introduction to Pipelining
- Pipelined RISC-V Datapath
- And in Conclusion, ...



# Processor



# Single-Cycle RISC-V RV32I Datapath



# Control Logic Truth Table (incomplete)

Inst[31:0]	BrEq	BrLT	PCSel	ImmSel	BrUn	ASel	BSel	ALUSel	MemRW	RegWEn	WBSel
add	*	*	+4	*	*	Reg	Reg	Add	Read	1	ALU
sub	*	*	+4	*	*	Reg	Reg	Sub	Read	1	ALU
(R-R Op)	*	*	+4	*	*	Reg	Reg	(Op)	Read	1	ALU
addi	*	*	+4	I	*	Reg	Imm	Add	Read	1	ALU
lw	*	*	+4	I	*	Reg	Imm	Add	Read	1	Mem
sw	*	*	+4	S	*	Reg	Imm	Add	Write	0	*
beq	0	*	+4	B	*	PC	Imm	Add	Read	0	*
beq	1	*	ALU	B	*	PC	Imm	Add	Read	0	*
bne	0	*	ALU	B	*	PC	Imm	Add	Read	0	*
bne	1	*	+4	B	*	PC	Imm	Add	Read	0	*
blt	*	1	ALU	B	0	PC	Imm	Add	Read	0	*
bltu	*	1	ALU	B	1	PC	Imm	Add	Read	0	*
jalr	*	*	ALU	I	*	Reg	Imm	Add	Read	1	PC+4
jal	*	*	ALU	J	*	PC	Imm	Add	Read	1	PC+4
auipc	*	*	+4	U	*	PC	Imm	Add	Read	1	ALU

Computer S

Kolb & Weaver

# Instruction type encoded using only 9 bits inst[30], inst[14:12], inst[6:2]

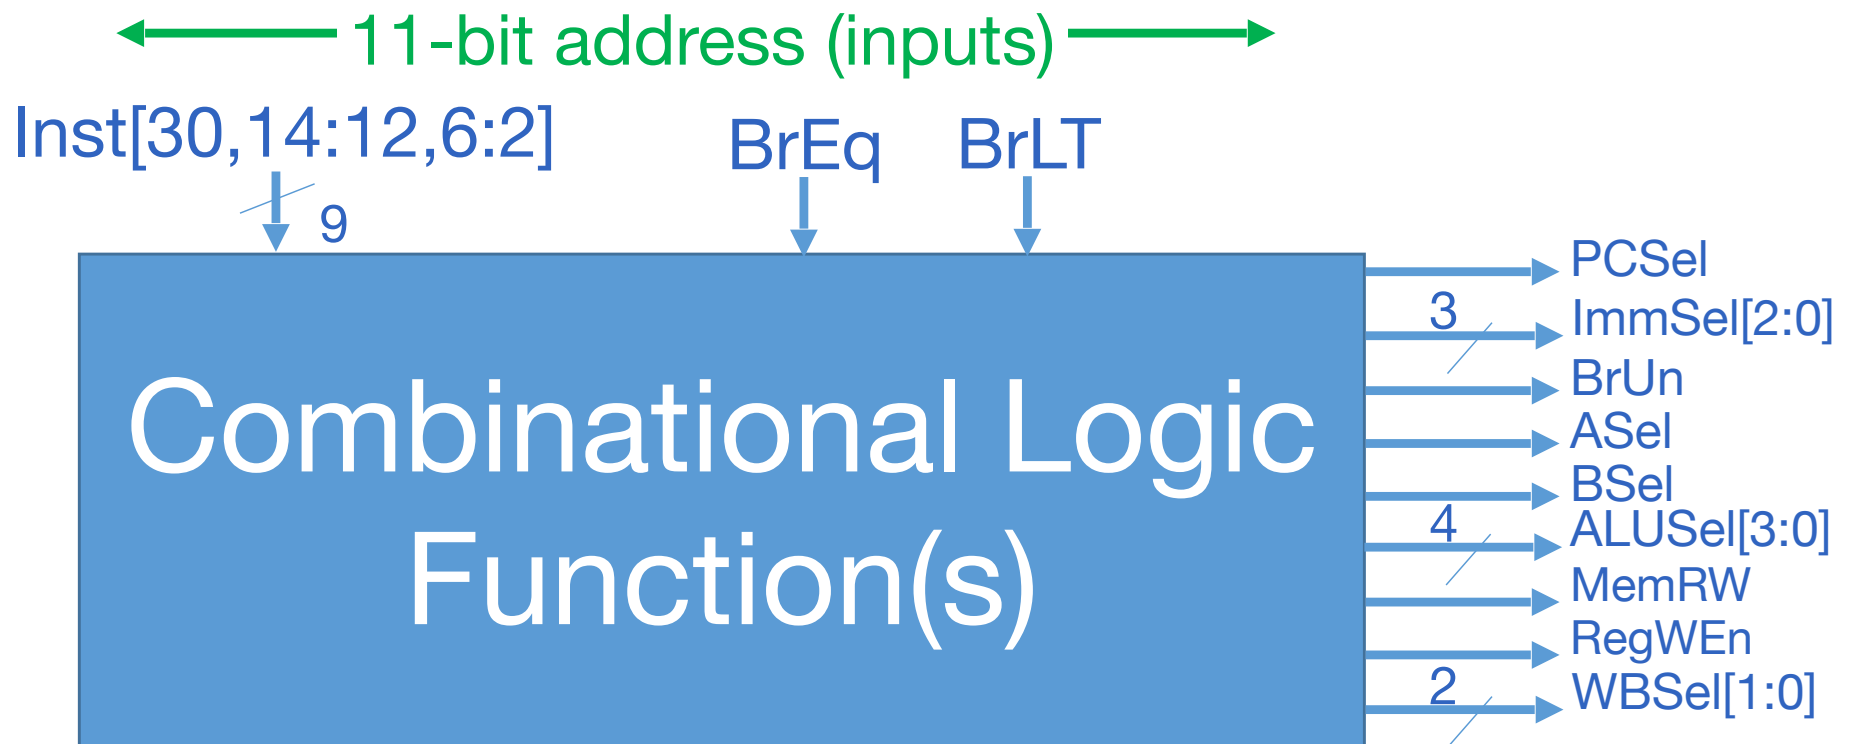
imm[31:12]				rd	011011	LUI
imm[31:12]				rd	001011	AUIPC
imm[20:10:1 11 19:12]				rd	110111	JAL
imm[11:0]	rs1	000		rd	110011	JALR
imm[12:10:5]	rs2	rs1	000	imm[4:1 11]	110001	BEQ
imm[12:10:5]	rs2	rs1	001	imm[4:1 11]	110001	BNE
imm[12:10:5]	rs2	rs1	100	imm[4:1 11]	110001	BLT
imm[12:10:5]	rs2	rs1	101	imm[4:1 11]	110001	BGE
imm[12:10:5]	rs2	rs1	110	imm[4:1 11]	110001	BLTU
imm[12:10:5]	rs2	rs1	111	imm[4:1 11]	110001	BGEU
imm[11:0]	rs1	000		rd	000001	LB
imm[11:0]	rs1	001		rd	000001	LH
imm[11:0]	rs1	010		rd	000001	LW
imm[11:0]	rs1	100		rd	000001	LBU
imm[11:0]	rs1	101		rd	000001	LHU
imm[11:5]	rs2	rs1	000	imm[4:0]	010001	SB
imm[11:5]	rs2	rs1	001	imm[4:0]	010001	SH
imm[11:5]	rs2	rs1	010	imm[4:0]	010001	SW
imm[11:0]	rs1	000		rd	001001	ADDI
imm[11:0]	rs1	010		rd	001001	SLTI
imm[11:0]	rs1	011		rd	001001	SLTIU
imm[11:0]	rs1	100		rd	001001	XORI
imm[11:0]	rs1	110		rd	001001	ORI
imm[11:0]	rs1	111		rd	001001	ANDI

inst[30]      inst[14:12]      inst[6:2]

0000000	shamt	rs1	00	rd	0010011	SLLI	
0000000	shamt	rs1	101	rd	0010011	SRLI	
0100000	shamt	rs1	101	rd	0010011	SRAI	
0000000	rs2	rs1	000	rd	0110011	ADD	
0100000	rs2	rs1	000	rd	0110011	SUB	
0000000	rs2	rs1	001	rd	0110011	SLL	
0000000	rs2	rs1	010	rd	0110011	SLT	
0000000	rs2	rs1	011	rd	0110011	SLTU	
0000000	rs2	rs1	100	rd	0110011	XOR	
0000000	rs2	rs1	101	rd	0110011	SRL	
0100000	rs2	rs1	101	rd	0110011	SRA	
0000000	rs2	rs1	110	rd	0110011	OR	
0000000	rs2	rs1	111	rd	0110011	AND	
0000	pred	succ	00000	000	00000	0001111	FENCE
0000	0000	0000	00000	001	00000	0001111	FENCE.I
000000000000			00000	000	00000	1110011	ECALL
000000000001			00000	000	00000	1110011	EBREAK
csr		rs1	001	rd	1110011	CSRRLW	
csr		rs1	010	rd	1110011	CSRRLS	
csr		rs1	011	rd	1110011	CSRRLC	
csr		zimm	101	rd	1110011	CSRRLWI	
csr		zimm	110	rd	1110011	CSRRLSI	
csr		zimm	111	rd	1110011	CSRRLCI	

Not in CS61C

# Control Block Design

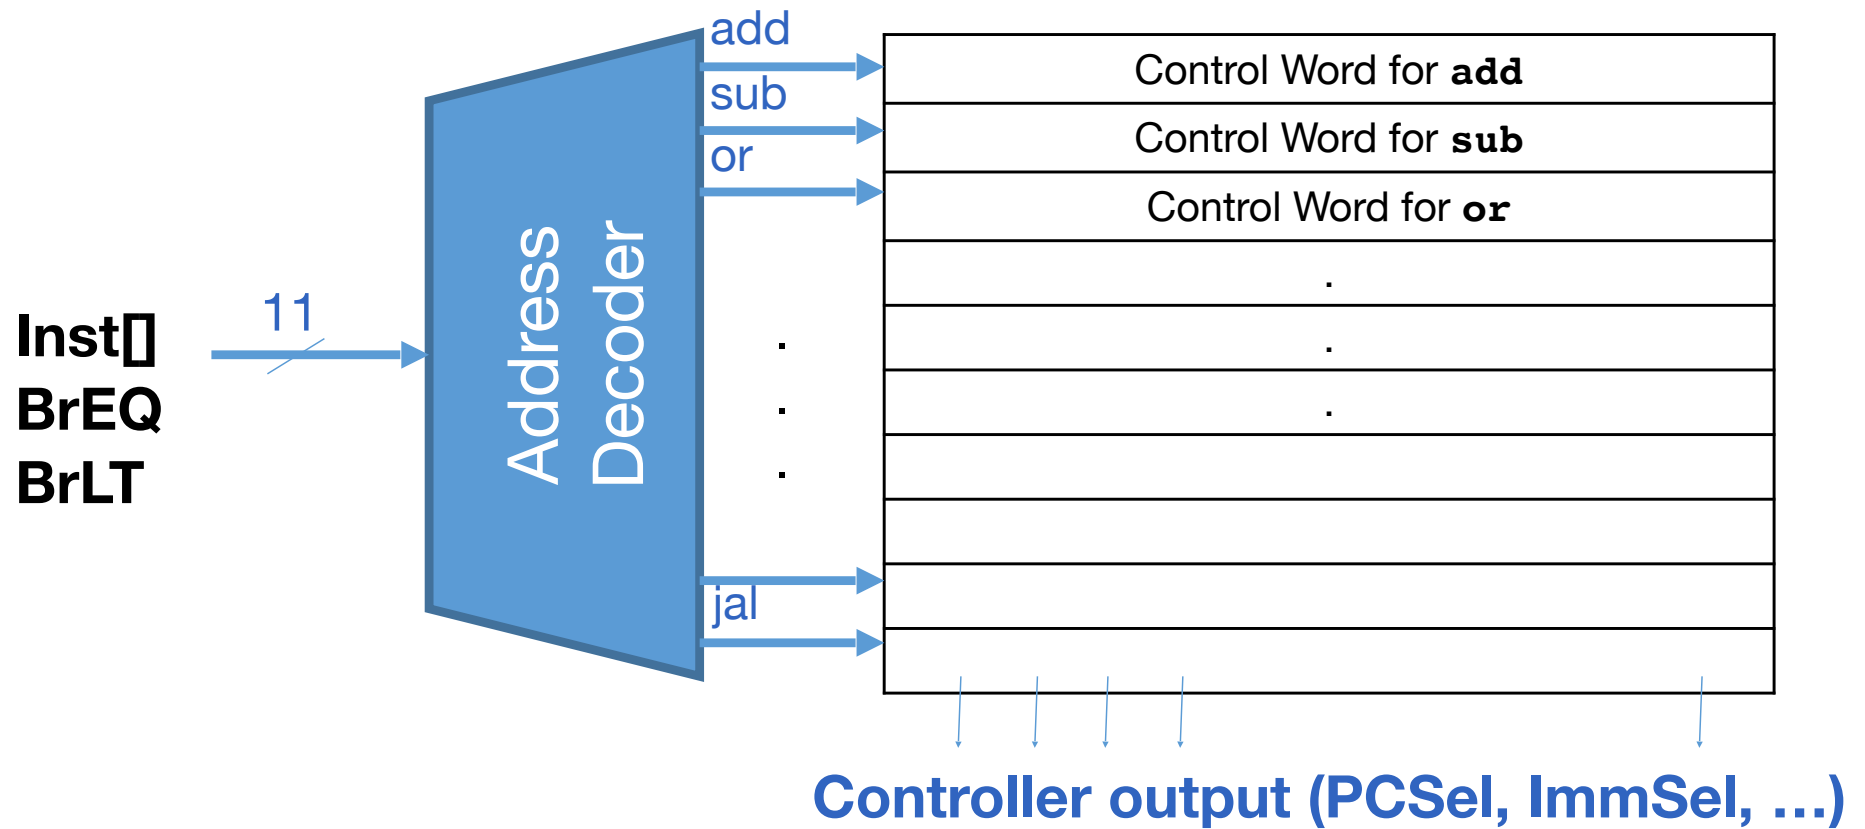


15 data bits (outputs)

# Control Realization Options

- ROM
  - “Read-Only Memory”
  - Regular structure
  - Can be easily reprogrammed
    - fix errors
    - add instructions
  - Popular when designing control logic manually
- Combinatorial Logic
  - Today, chip designers use logic synthesis tools to convert truth tables to networks of gates

# ROM Controller Implementation

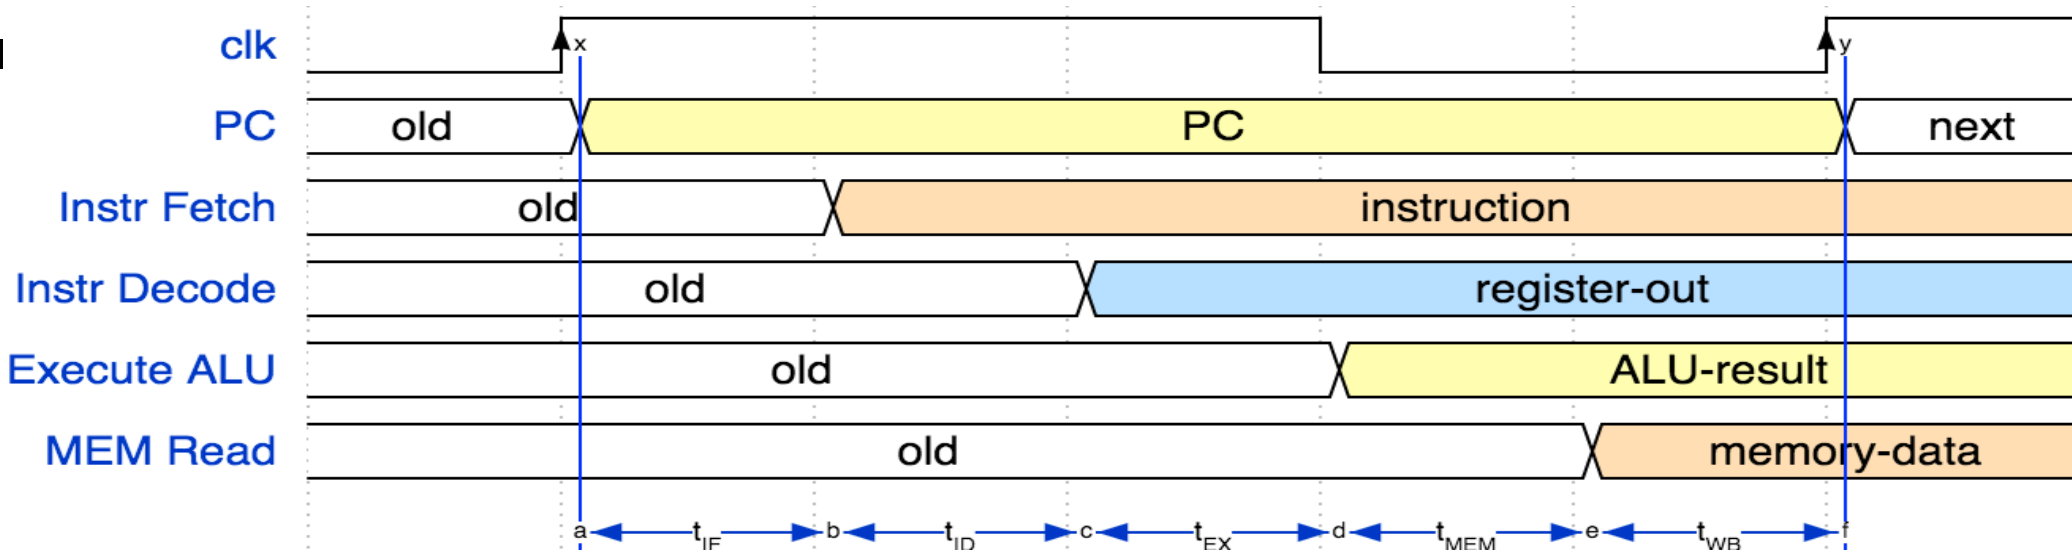


# Agenda

- Finish Single-Cycle RISC-V Datapath
- Controller
- **Instruction Timing**
- Performance Measures
- Introduction to Pipelining
- Pipelined RISC-V Datapath
- And in Conclusion, ...



# Approximate Instruction Timing



IF	ID	EX	MEM	WB	Total
I-MEM	Reg Read	ALU	D-MEM	Reg W	
200 ps	100 ps	200 ps	200 ps	100 ps	800 ps

# Instruction Timing

Instr	IF = 200ps	ID = 100ps	ALU = 200ps	MEM=200ps	WB = 100ps	Total
add	X	X	X		X	600ps
beq	X	X	X			500ps
jal	X	X	X			500ps
lw	X	X	X	X	X	800ps
sw	X	X	X	X		700ps

- Maximum clock frequency
  - $f_{\max} = 1/800\text{ps} = 1.25 \text{ GHz}$
- Most blocks idle most of the time
  - E.g.  $f_{\max, \text{ALU}} = 1/200\text{ps} = 5 \text{ GHz!}$
  - How can we keep ALU busy all the time?
  - 5 billion adds/sec, rather than just 1.25 billion?
  - Idea: Factories use three employee shifts - equipment is always busy!

# Agenda

- Finish Single-Cycle RISC-V Datapath
- Controller
- Instruction Timing
- **Performance Measures**
- Introduction to Pipelining
- Pipelined RISC-V Datapath
- And in Conclusion, ...

# Performance Measures

- “Our” RISC-V executes instructions at 1.25 GHz
  - 1 instruction every 800 ps
- Can we improve its performance?
  - What do we mean with this statement?
  - Not so obvious:
    - Quicker response time, so one job finishes faster?
    - More jobs per unit time (e.g. web server returning pages)?
    - Longer battery life?

# Transportation Analogy



Computer Science 61C

Kolb & Weaver

	Race Car	Bus
Passenger Capacity	1	50
Travel Speed	200 mph	50 mph
Gas Mileage	5 mpg	2 mpg

## 50 Mile trip:

	Race Car	Bus
Travel Time	15 min	60 min
Time for 100 passengers	1500 min	120 min
Gallons per passenger	10 gallons	0.5 gallons

# Computer Analogy

Transportation	Computer
Trip Time	Program execution time ( <b>latency</b> ): e.g. time to update display
Time for 100 passengers	<b>Throughput</b> : e.g. number of server requests handled per hour
Gallons per passenger	Energy per task*: e.g. how many movies you can watch per battery charge or energy bill for datacenter

- \* Note: power is not a good measure, since low-power CPU might run for a long time to complete one task consuming more energy than faster computer running at higher power for a shorter time

# “Iron Law” of Processor Performance

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} * \frac{\text{Cycles}}{\text{Instruction}} * \frac{\text{Time}}{\text{Cycle}}$$

# Instructions per Program

- Determined by
  - Task
  - Algorithm, e.g.  $O(N^2)$  vs  $O(N)$
  - Programming language
  - Compiler
  - Instruction Set Architecture (ISA)
- Input
  - It is, of course, the halting problem to know how many instructions a program will take in advance



# (Average) Clock cycles per Instruction

- Determined by
  - ISA (CISC versus RISC)
  - Processor implementation (or *microarchitecture*)
    - E.g. for “our” single-cycle RISC-V design,  $\text{CPI} = 1$
  - Superscalar processors,  $\text{CPI} < 1$  (next lecture)

# Time per Cycle ( $1/\text{Frequency}$ )

- Determined by
  - Processor microarchitecture (determines critical path through logic gates)
  - Technology (e.g. 14nm versus 28nm)
  - Power budget (lower voltages reduce transistor speed)

# Speed Tradeoff Example

- For some task (e.g. image compression) ...

	Processor A	Processor B
# Instructions	1 Million	1.5 Million
Average CPI	2.5	1
Clock rate $f$	2.5 GHz	2 GHz
Execution time	1 ms	0.75 ms

Processor B is faster for this task, despite executing more instructions and having a lower clock rate!

# Energy per Task

$$\frac{\text{Energy}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} * \frac{\text{Energy}}{\text{Instruction}}$$

$$\frac{\text{Energy}}{\text{Program}} \propto \frac{\text{Instructions}}{\text{Program}} * C V^2$$

“Capacitance” depends on technology, microarchitecture, circuit details

Supply voltage, e.g. 1V

Want to reduce capacitance and voltage to reduce energy/task

# Energy Tradeoff Example

- “Next-generation” processor (Moore’s law)
  - Capacitance,  $C$ : reduced by 15 %
  - Supply voltage,  $V_{\text{sup}}$ : reduced by 15 %
  - Energy consumption:  $(.85C)(.85V)^2 = .63E \Rightarrow$  **-39 % reduction**
- Significantly improved energy efficiency thanks to
  - Moore’s Law **AND**
  - Reduced supply voltage

# Energy “Iron Law”

- Energy efficiency (e.g., instructions/Joule) is key metric in all computing devices
- For power-constrained systems (e.g., 20MW datacenter), need better energy efficiency to get more performance at same power
- For energy-constrained systems (e.g., 1W phone), need better energy efficiency to prolong battery life


$$\text{Performance} = \frac{\text{Power}}{\text{Energy Efficiency}}$$

*(Tasks/Second)      (Joules/Second)      (Tasks/Joule)*

# End of Scaling

- In recent years, industry has not been able to reduce supply voltage much, as reducing it further would mean increasing “leakage power” where transistor switches don’t fully turn off (more like dimmer switch than on-off switch)
- Also, size of transistors and hence capacitance, not shrinking as much as before between transistor generations
  - Rather than horizontal modern CMOS uses vertically-aligned transistors to pack them closer together... But that doesn't reduce capacitance just allows for higher density
- Power becomes a growing concern – the “power wall”
- Cost-effective air-cooled chip limit around ~150W

# Processor Trends





# Agenda

- Finish Single-Cycle RISC-V Datapath
- Controller
- Instruction Timing
- Performance Measures
- **Introduction to Pipelining**
- Pipelined RISC-V Datapath
- And in Conclusion, ...

# Pipelining

- A familiar example:
  - Getting a university degree



Year 1



Year 2



Year 3

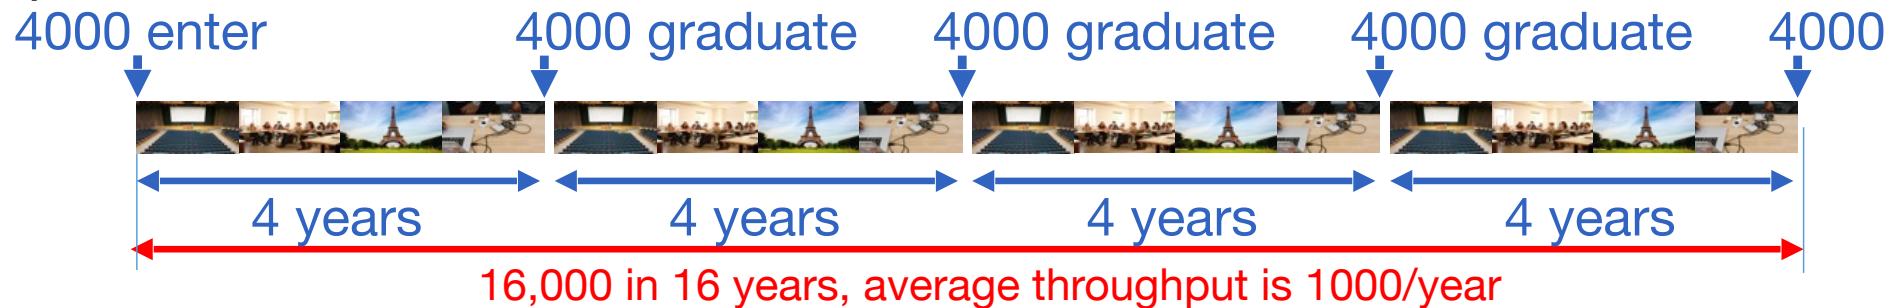


Year 4


- Shortage of Computer scientists (your startup is growing):
  - How long does it take to educate 16,000 students?

# Computer Scientist Education

- Option 1: **serial**



- Option 2: **pipelining**



- 16,000 in 7 years
- Steady state throughput is 4000/year
- Resources used efficiently
- **4-fold improvement over serial education**

# Latency versus Throughput

- **Latency**

- Time from entering college to graduation
- Serial 4 years
- Pipelining 4 years

- **Throughput**


- Average number of students graduating each year
- Serial 1000
- Pipelining 4000

- **Pipelining**

- Increases throughput (4x in this example)
- But can ***never improve*** latency
  - sometimes worse (additional overhead e.g. for shift transition)
  - AKA: "Why Nick's Ph.D. dissertation was a f@#@)(@#\* stupid idea!"

# Simultaneous versus Sequential







- What happens *sequentially*?
- What happens *simultaneously*?



# Agenda

- Finish Single-Cycle RISC-V Datapath
- Controller
- Instruction Timing
- Performance Measures
- Introduction to Pipelining
- **Pipelined RISC-V Datapath**
- And in Conclusion, ...

# Pipelining with RISC-V

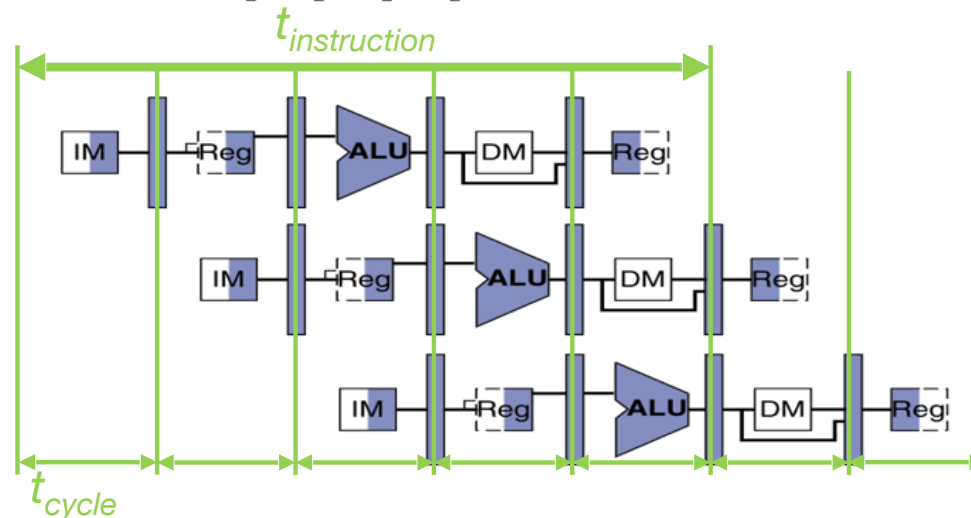
Phase	Pictogram	$t_{step}$ Serial	$t_{cycle}$ Pipelined
Instruction Fetch		200 ps	200 ps
Reg Read		100 ps	200 ps
ALU		200 ps	200 ps
Memory		200 ps	200 ps
Register Write		100 ps	200 ps
$t_{instruction}$		<b>800 ps</b>	<b>1000 ps</b>

instruction sequence

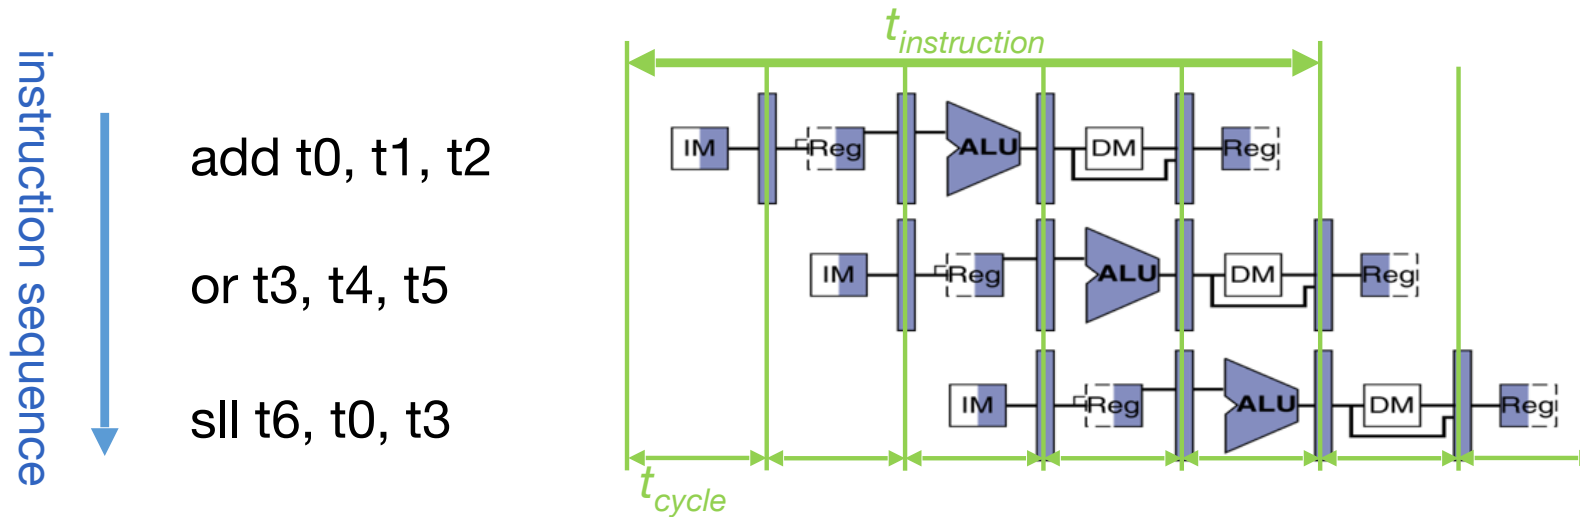
add t0, t1, t2

or t3, t4, t5

sll t6, t0, t3



# Pipelining with RISC-V

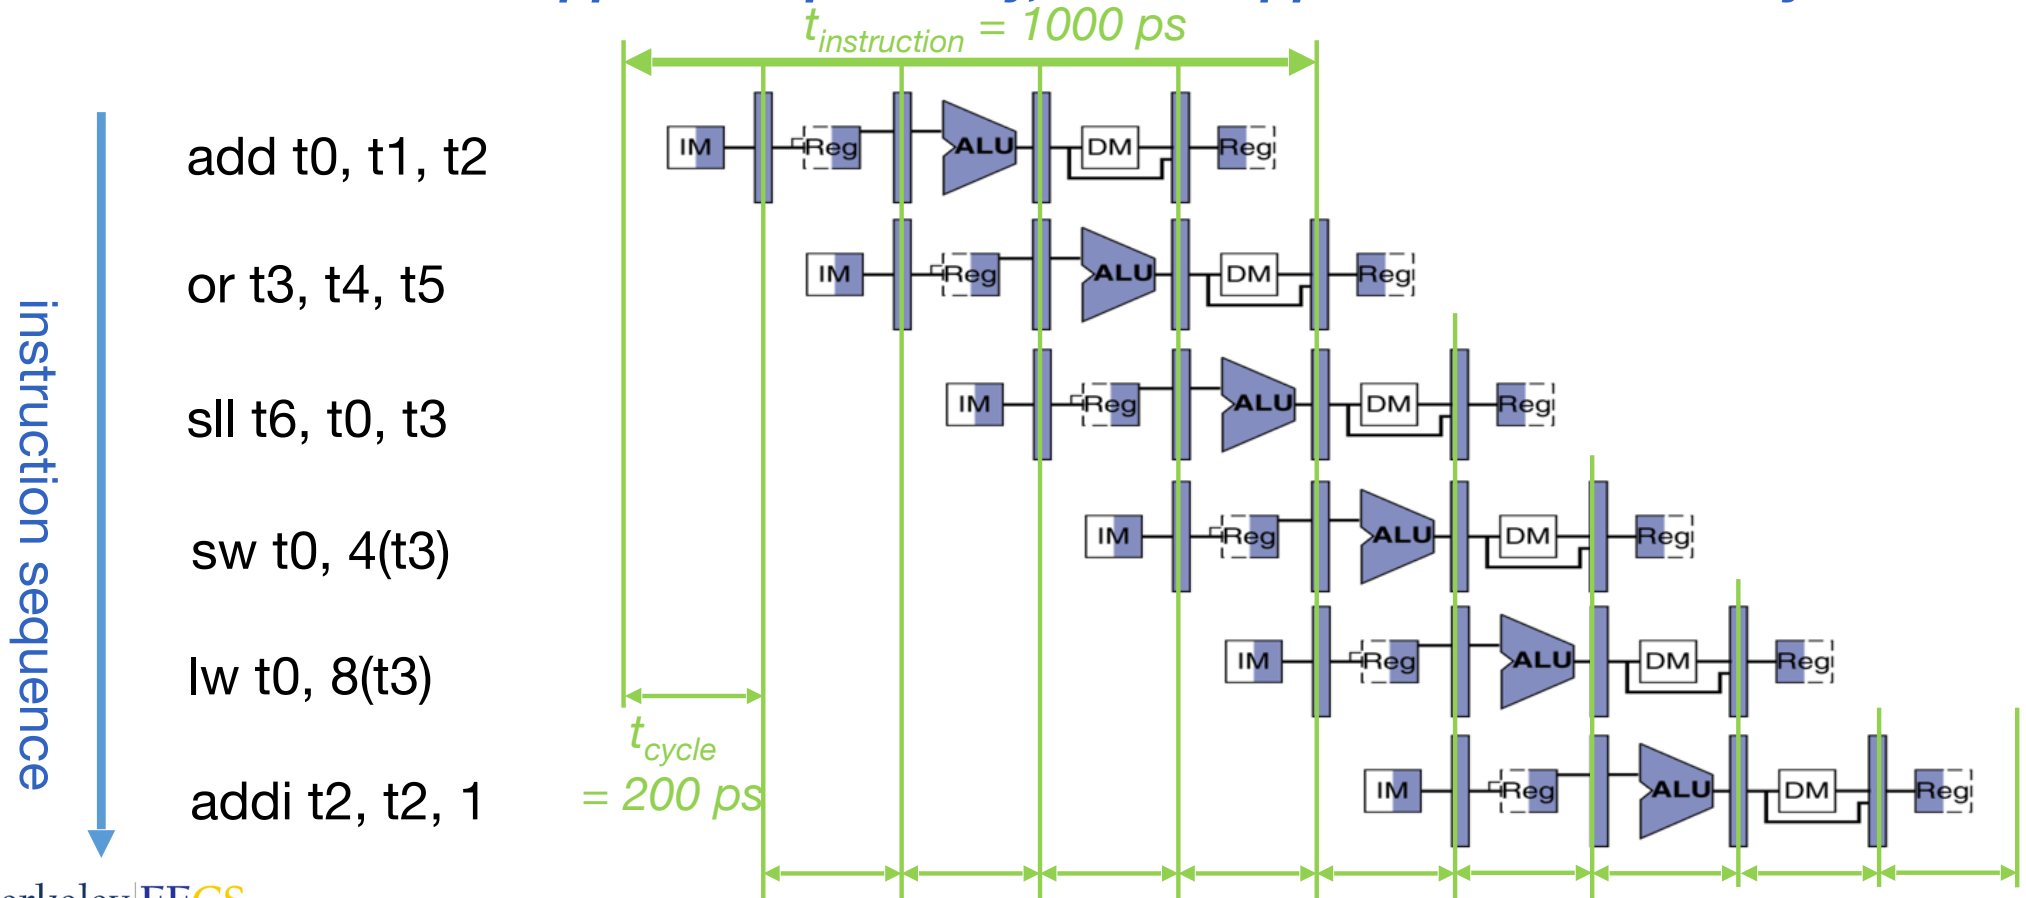


	Single Cycle	Pipelining
Timing	$t_{step} = 100 \dots 200 \text{ ps}$	$t_{cycle} = 200 \text{ ps}$
	Register access only 100 ps	All cycles same length
Instruction time, $t_{instruction}$	$= t_{cycle} = 800 \text{ ps}$	1000 ps
Clock rate, $f_s$	$1/800 \text{ ps} = 1.25 \text{ GHz}$	$1/200 \text{ ps} = 5 \text{ GHz}$
Relative speed	1 x	4 x



# Sequential vs Simultaneous

*What happens sequentially, what happens simultaneously?*



# Agenda

- Finish Single-Cycle RISC-V Datapath
- Controller
- Instruction Timing
- Performance Measures
- Introduction to Pipelining
- Pipelined RISC-V Datapath
- **And in Conclusion, ...**

# And in Conclusion, ...

- Controller
  - Tells universal datapath how to execute each instruction
- Instruction timing
  - Set by instruction complexity, architecture, technology
  - Pipelining increases clock frequency, “instructions per second”
    - But does not reduce time to complete instruction
- Performance measures
  - Different measures depending on objective
    - Response time
    - Jobs / second
    - Energy per task